

Machine Learning for Image Analysis

Lecture 1: Neural Networks - Theory and Image Classification

Dagmar Kainmueller, 16.04.2024



Course Outline

- **Introduction to Image Analysis**
- **Basics: Neural Networks**
- Convolutional Neural Networks
- Transformers
- Model Interpretability
- Self-supervised Learning
- Generative Models (GANs, Diffusion)

Requirements

- Linear Regression, Logistic Regression
 - Hypothesis function
 - Loss function (cross-entropy, mean squared error)
 - Regularization: Weight decay
- Gradient Descent
- Basic linear algebra and multi-variate calculus

Introduction to Image Analysis

Image Classification



Assign one class label to an image
{dog, cat, truck, plane, ...}

→ CAT

- No spatial extent of the output
- What if localized information is sought in an image? → Pixel-wise tasks

Image Segmentation: Semantic Segmentation

Assign one class label to each pixel in an image

{dog, cat, truck, plane, ...}



{cat, grass, tree, sky}

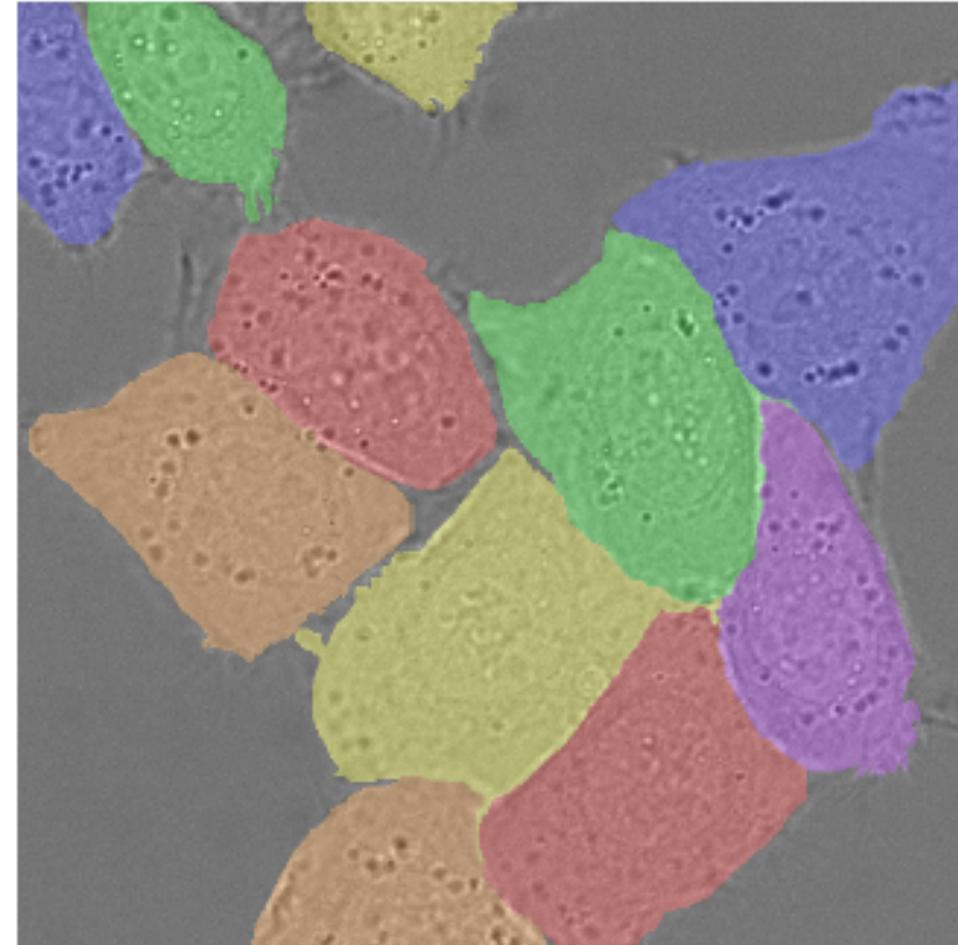
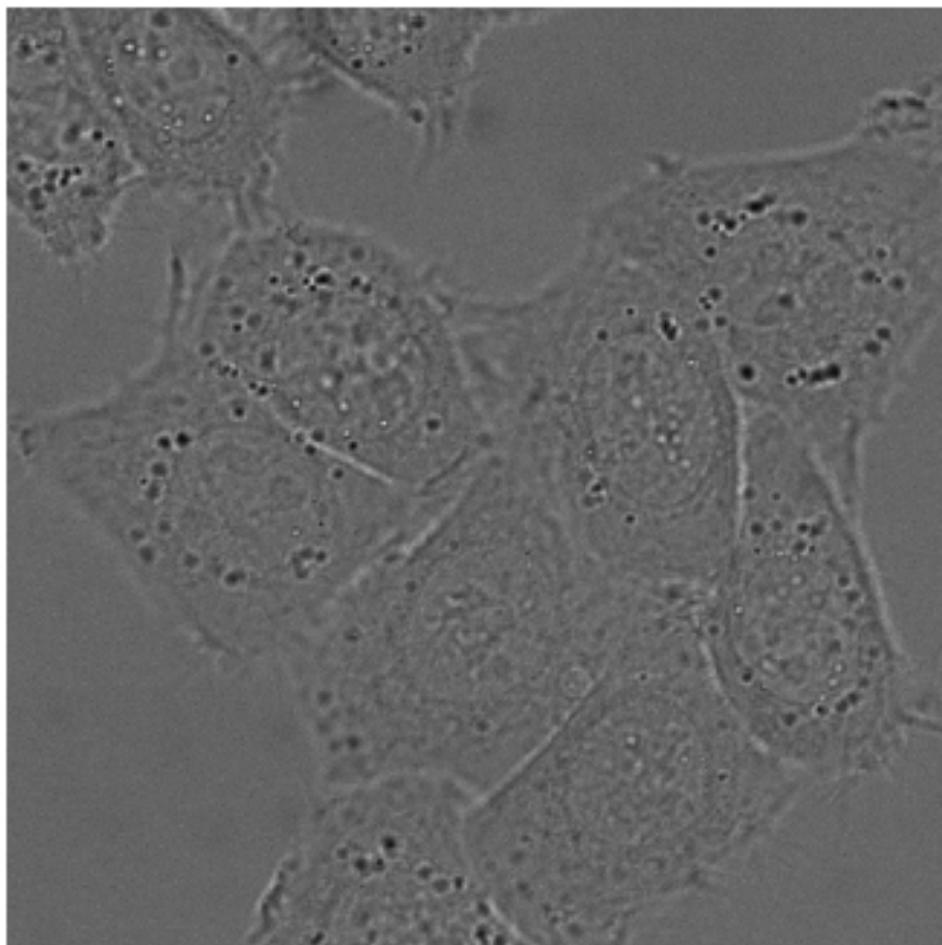


{cow, grass, tree, sky}

- Does not distinguish individual object instances

Instance Segmentation

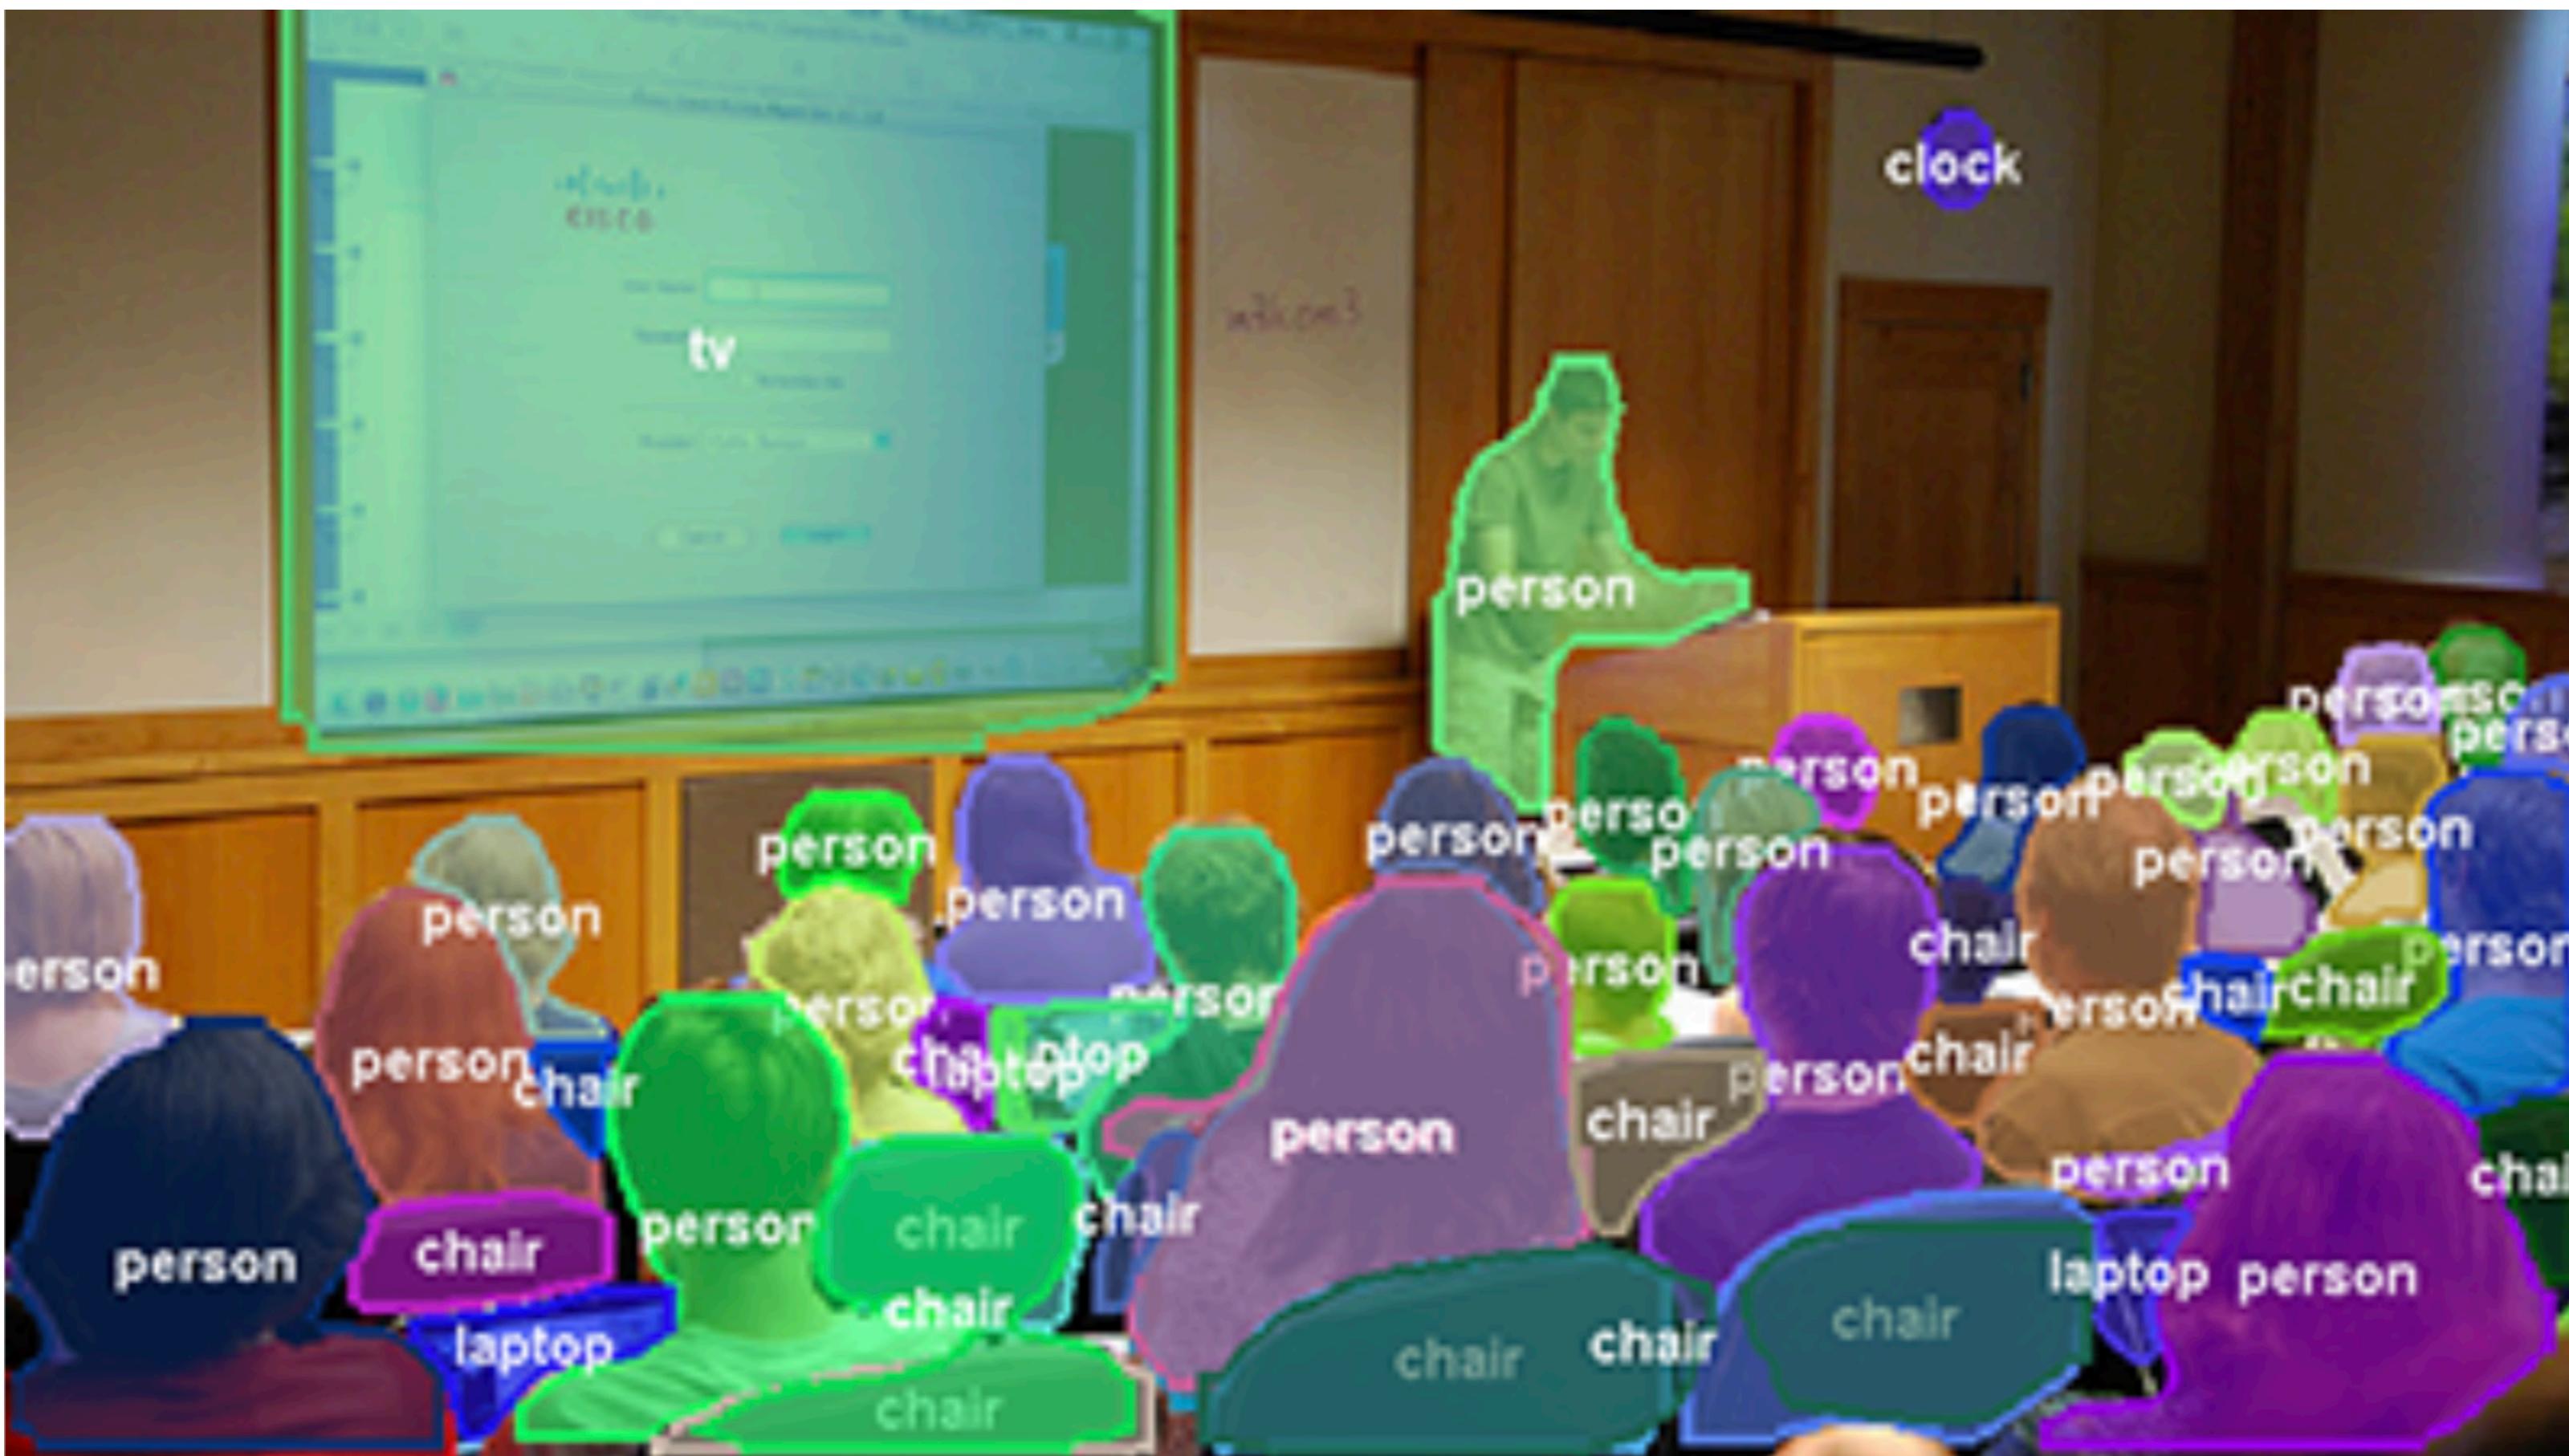
Assign one instance label to each pixel in an image



<https://segment-anything.com/demo#>

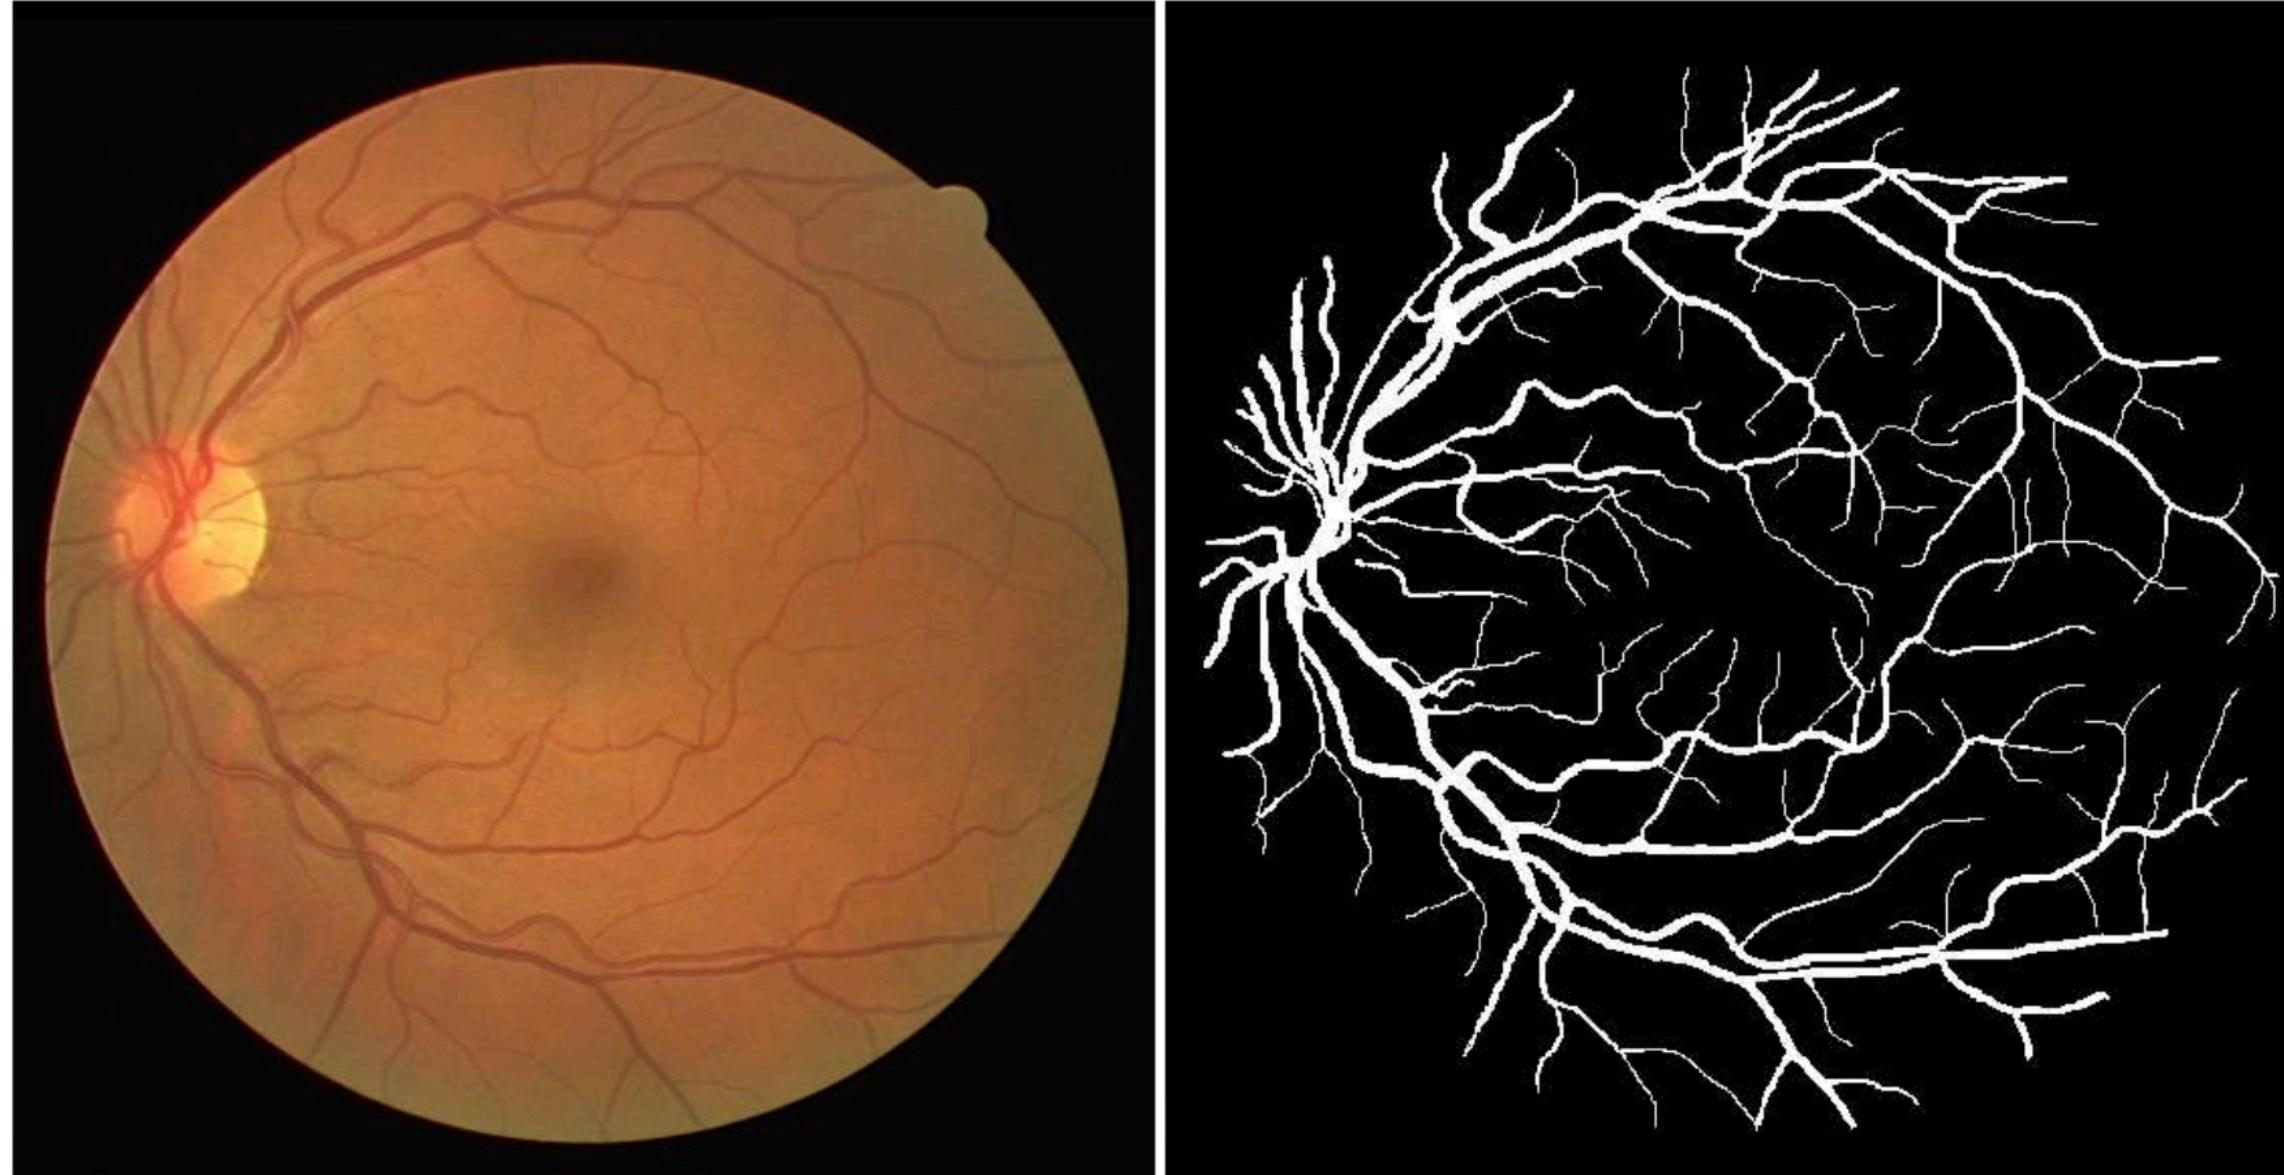
Semantic Instance Segmentation (aka Panoptic Segmentation)

Assign one class and instance label to each pixel in an image



Some more segmentation examples ...

Foreground/Background Segmentation



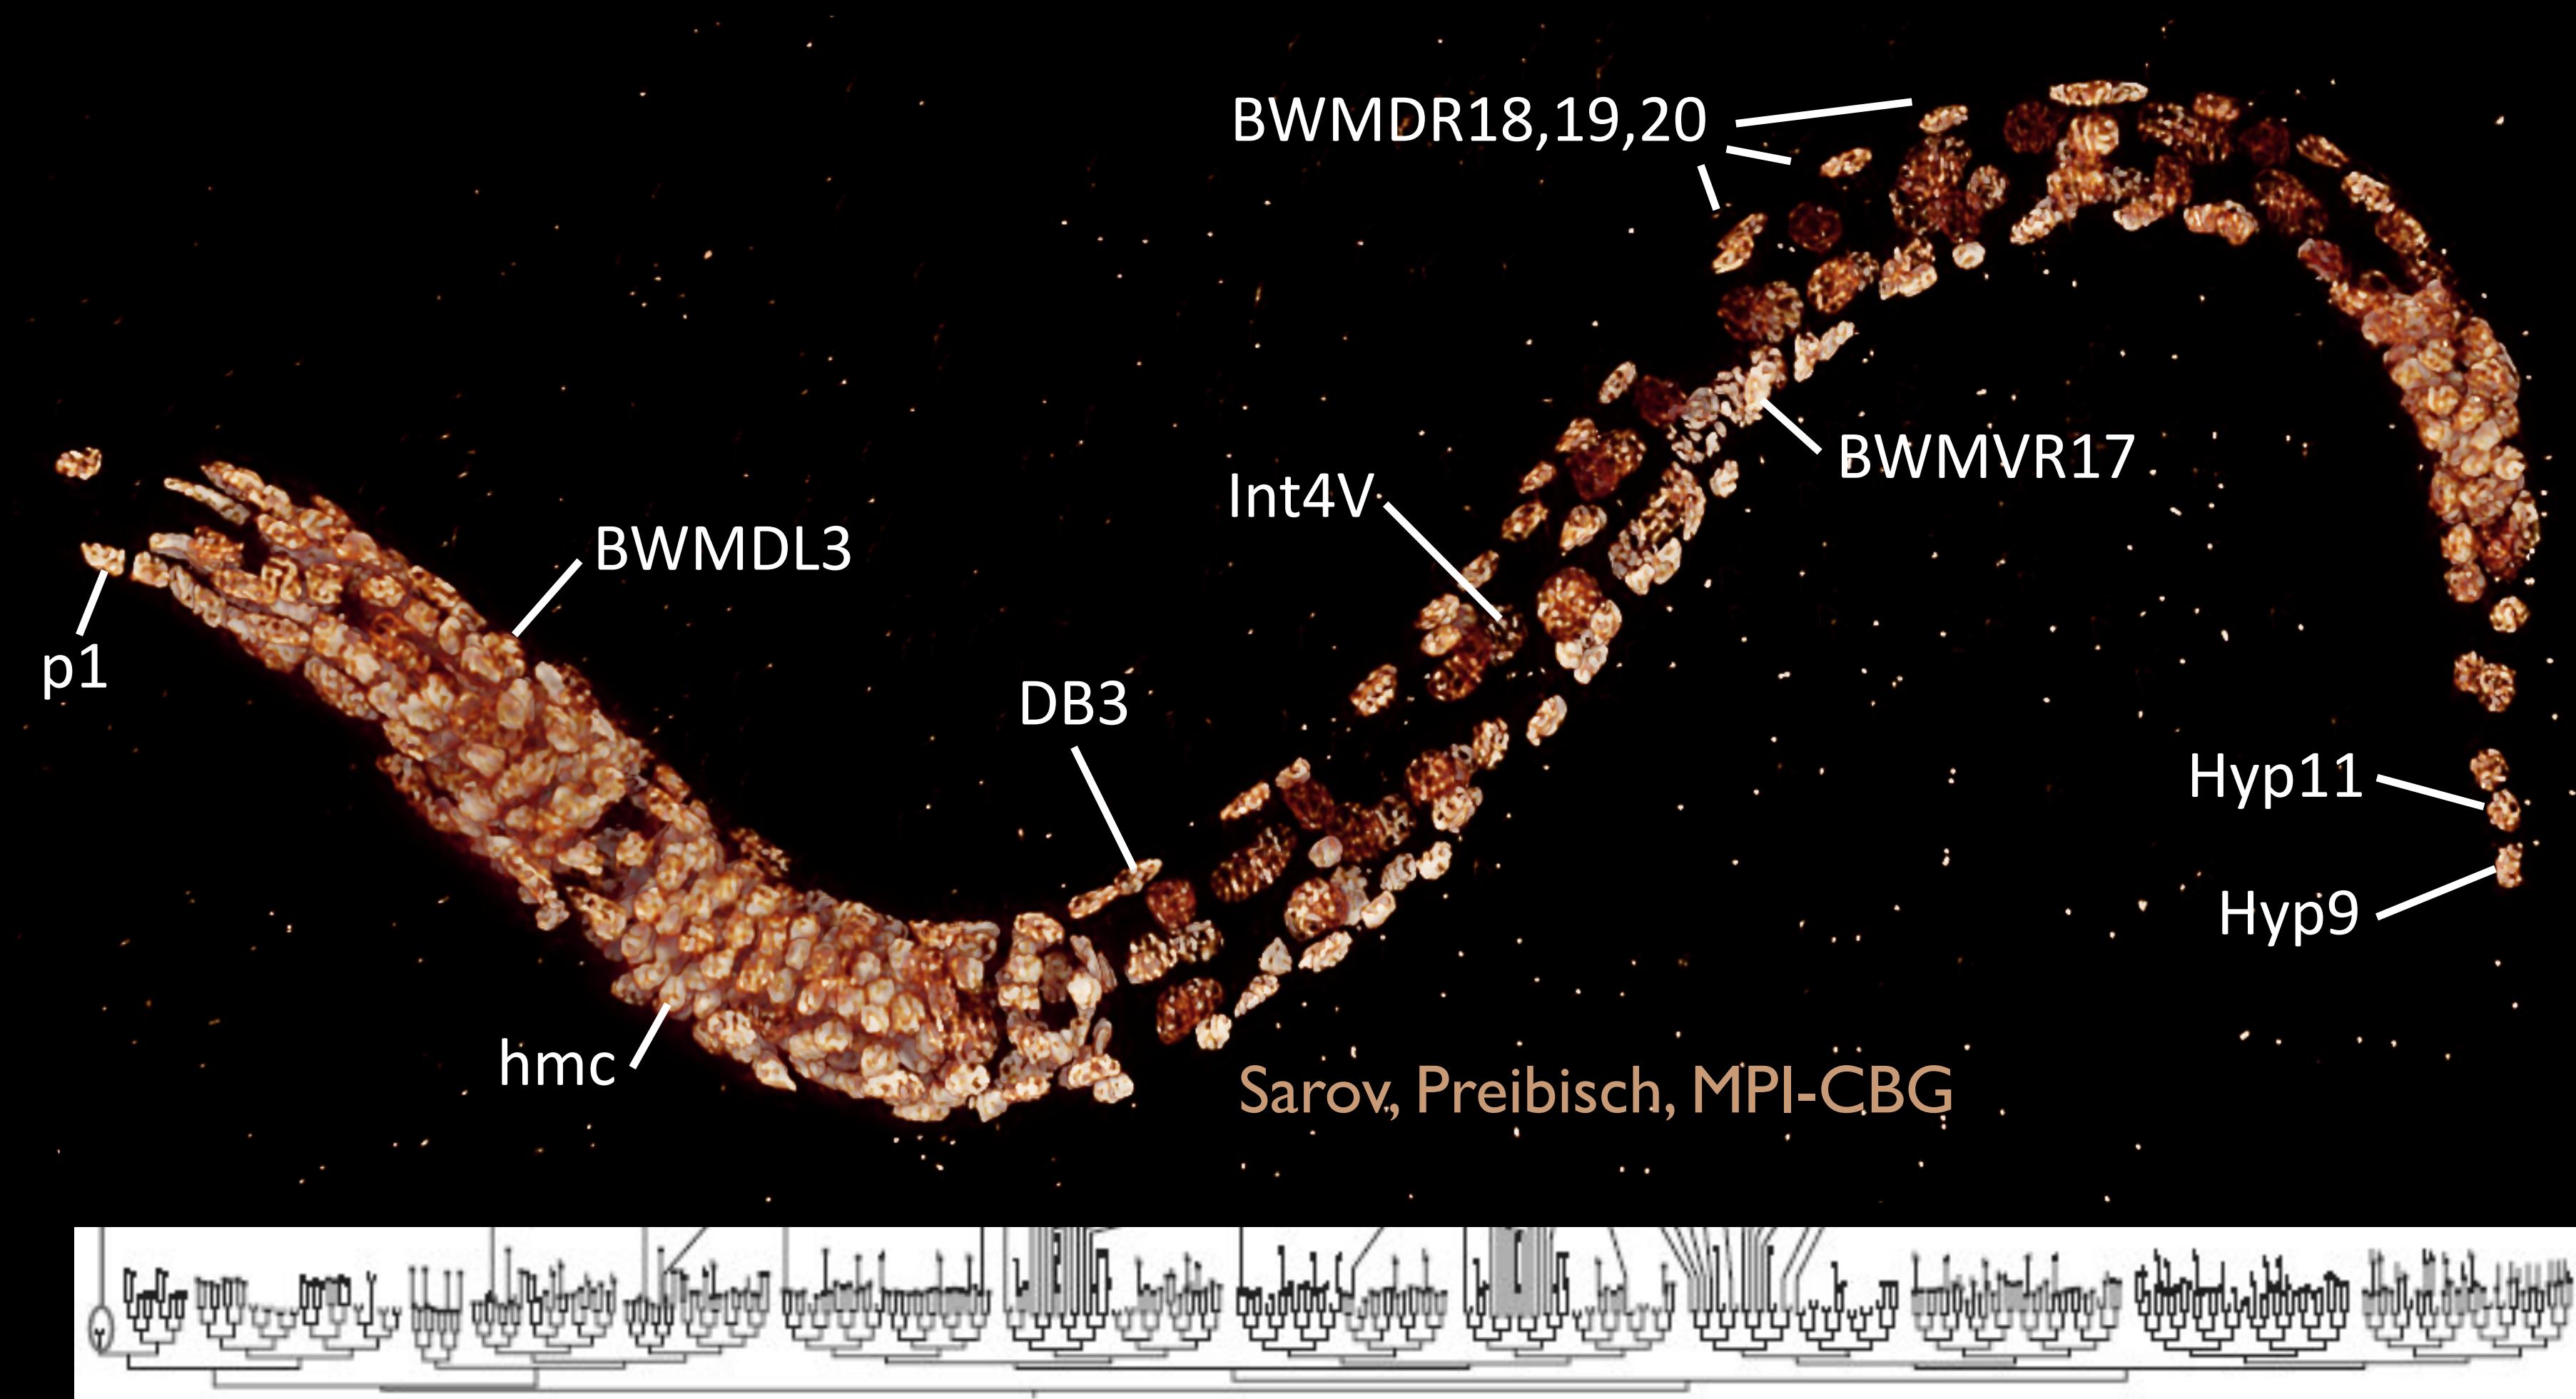
https://github.com/CVxTz/medical_image_segmentation

Semantic Parts Segmentation (aka Object Parsing)



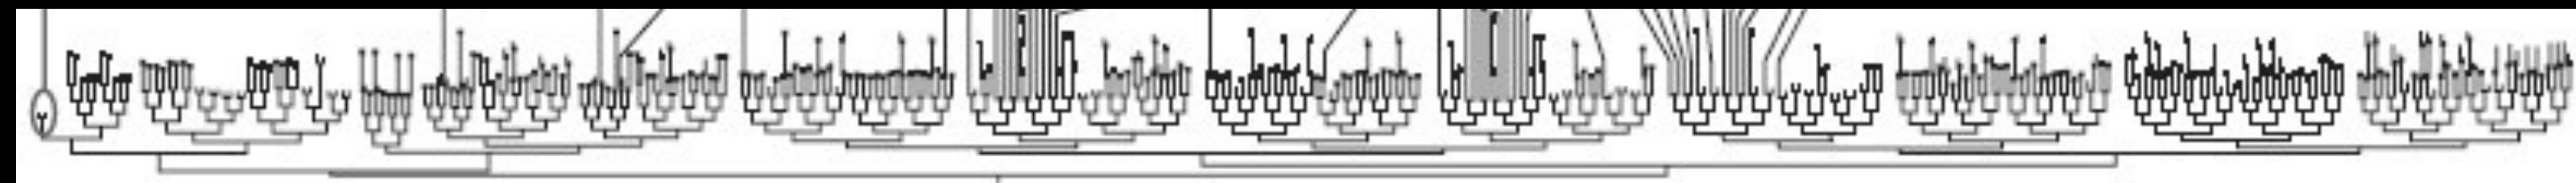
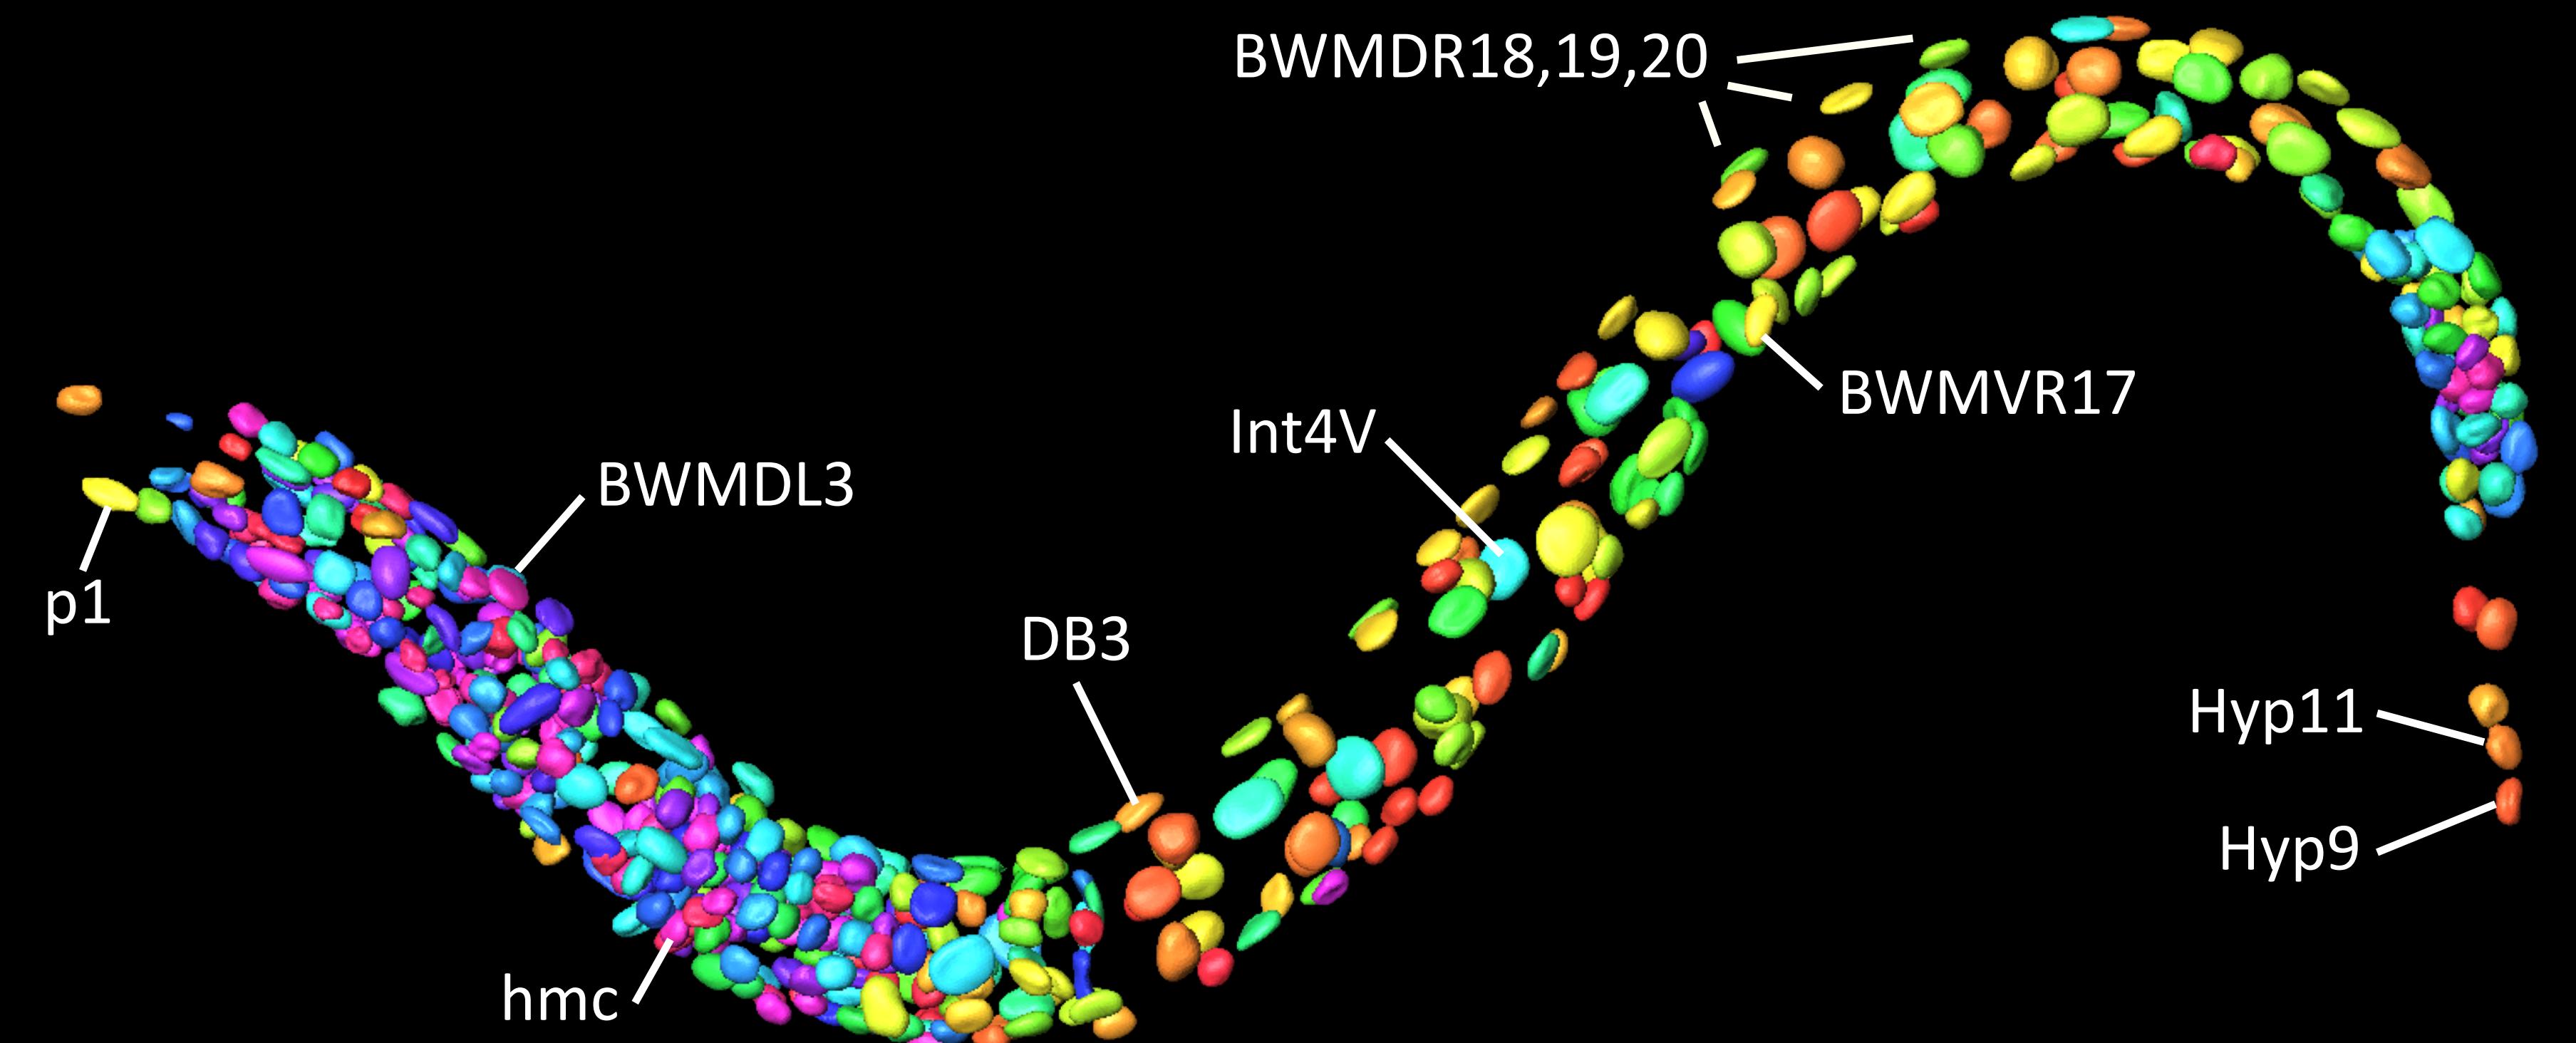
<https://www.arxiv-vanity.com/papers/1611.06612/>

Some more segmentation examples ...



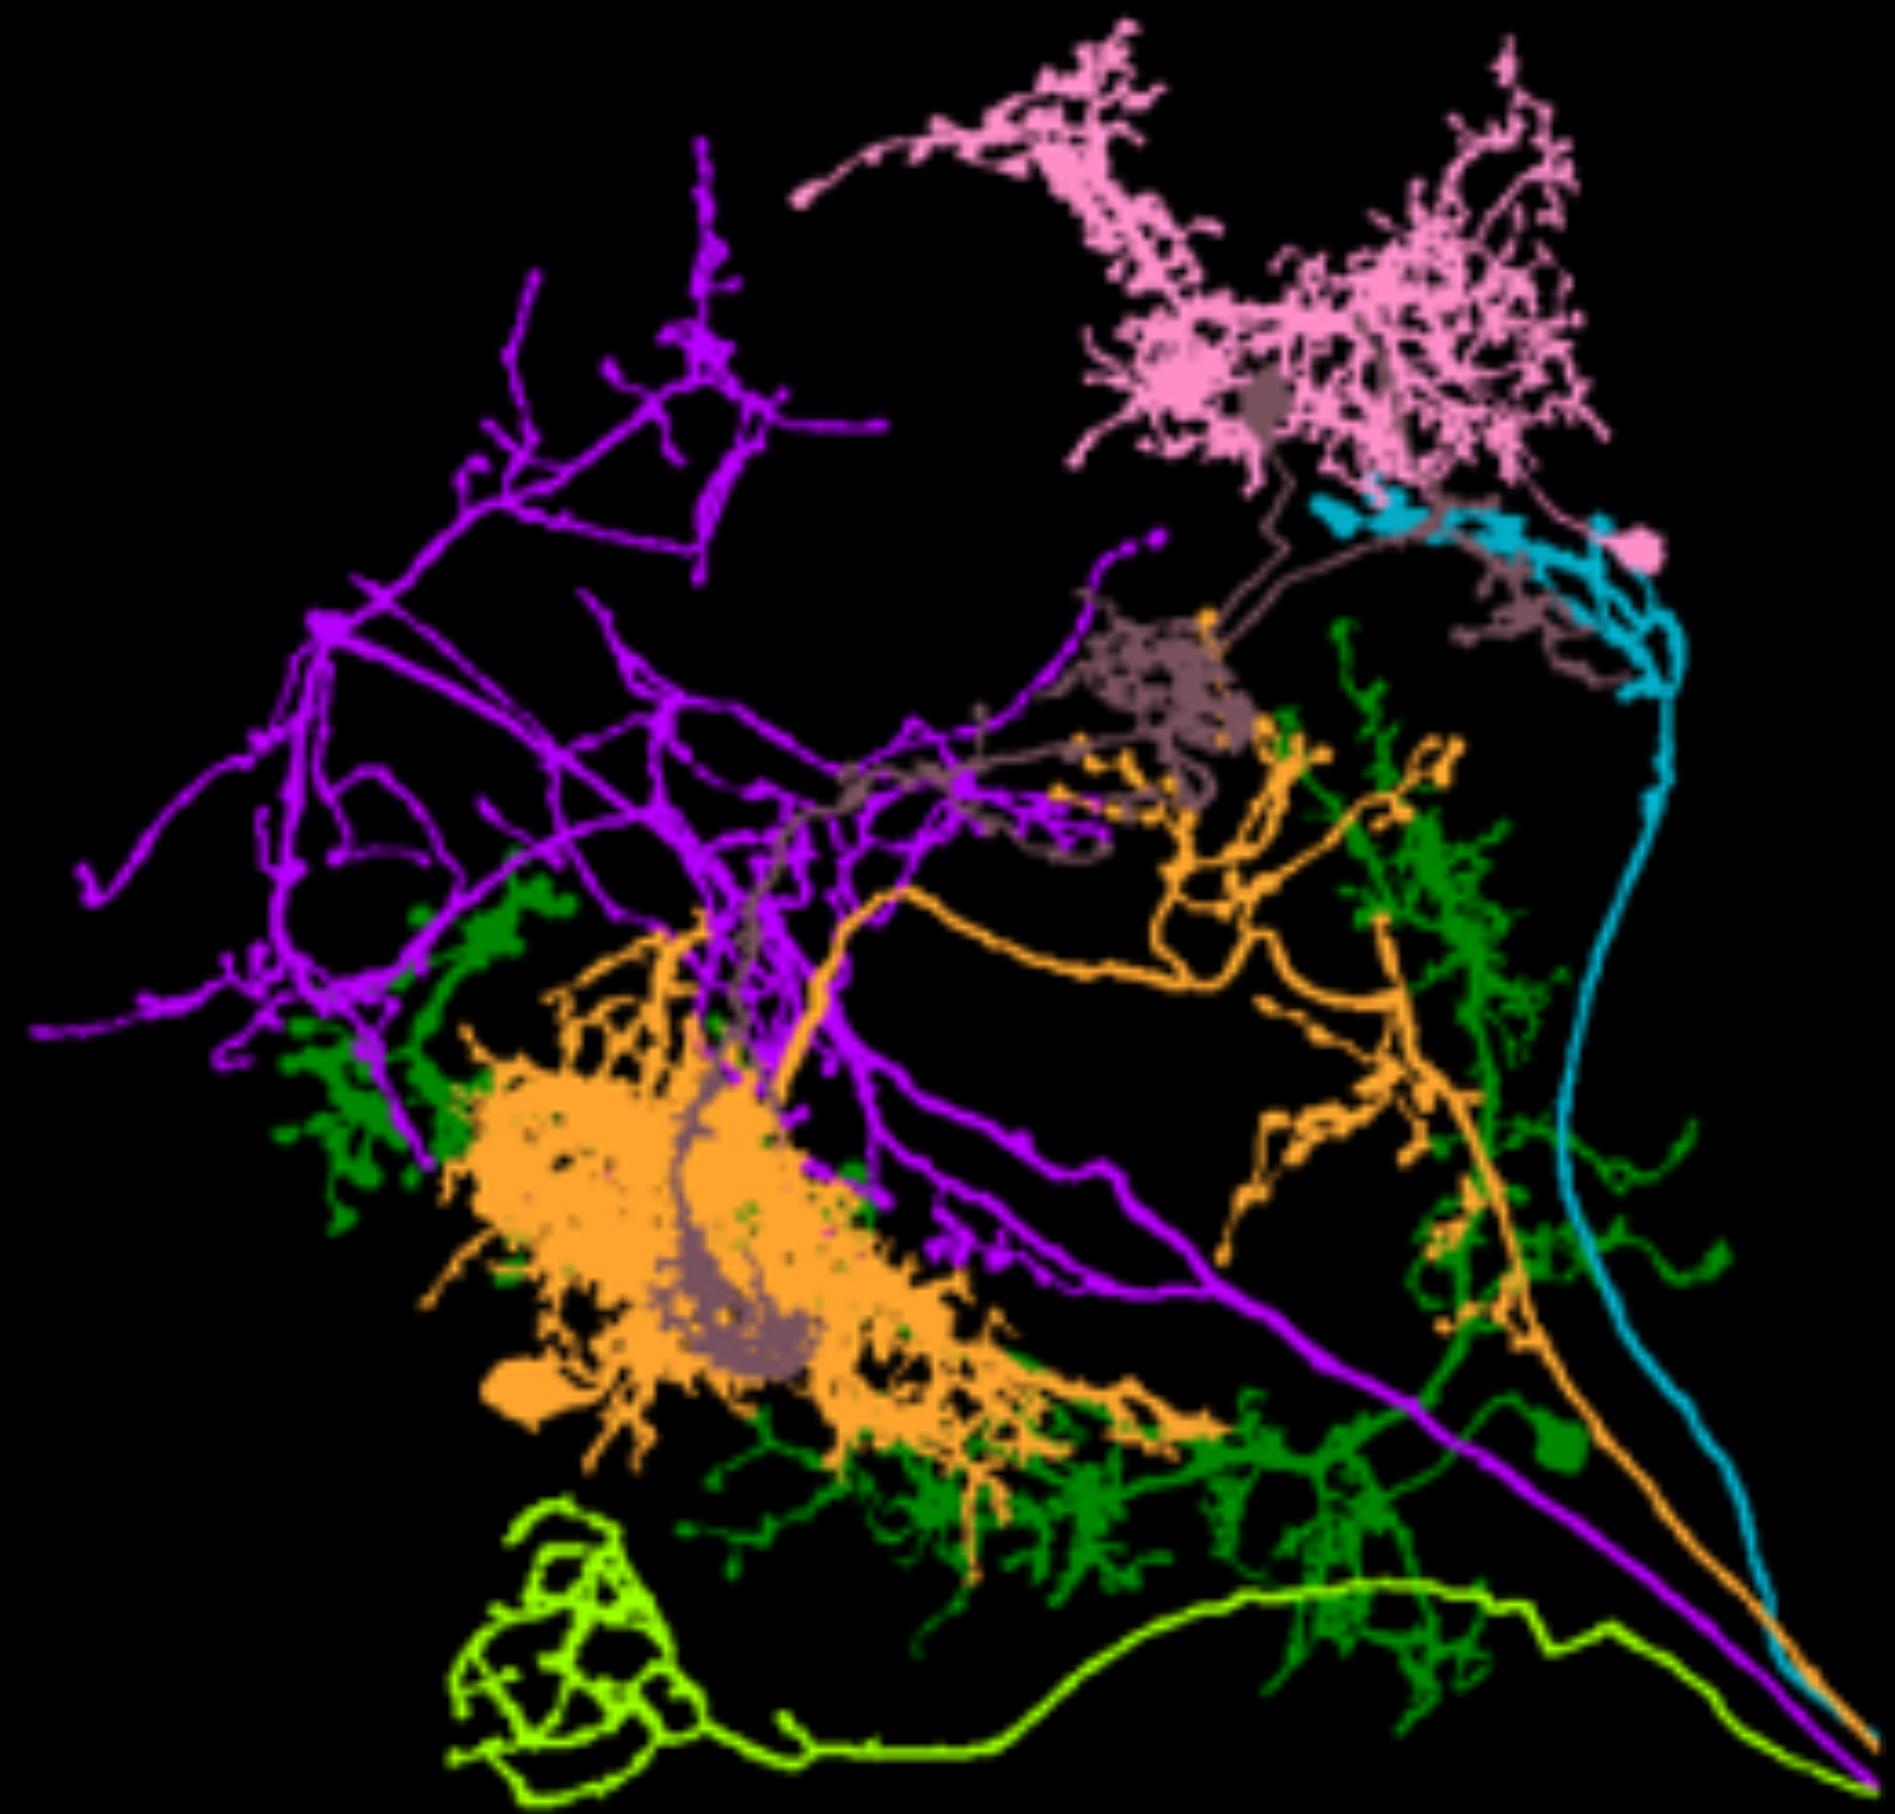
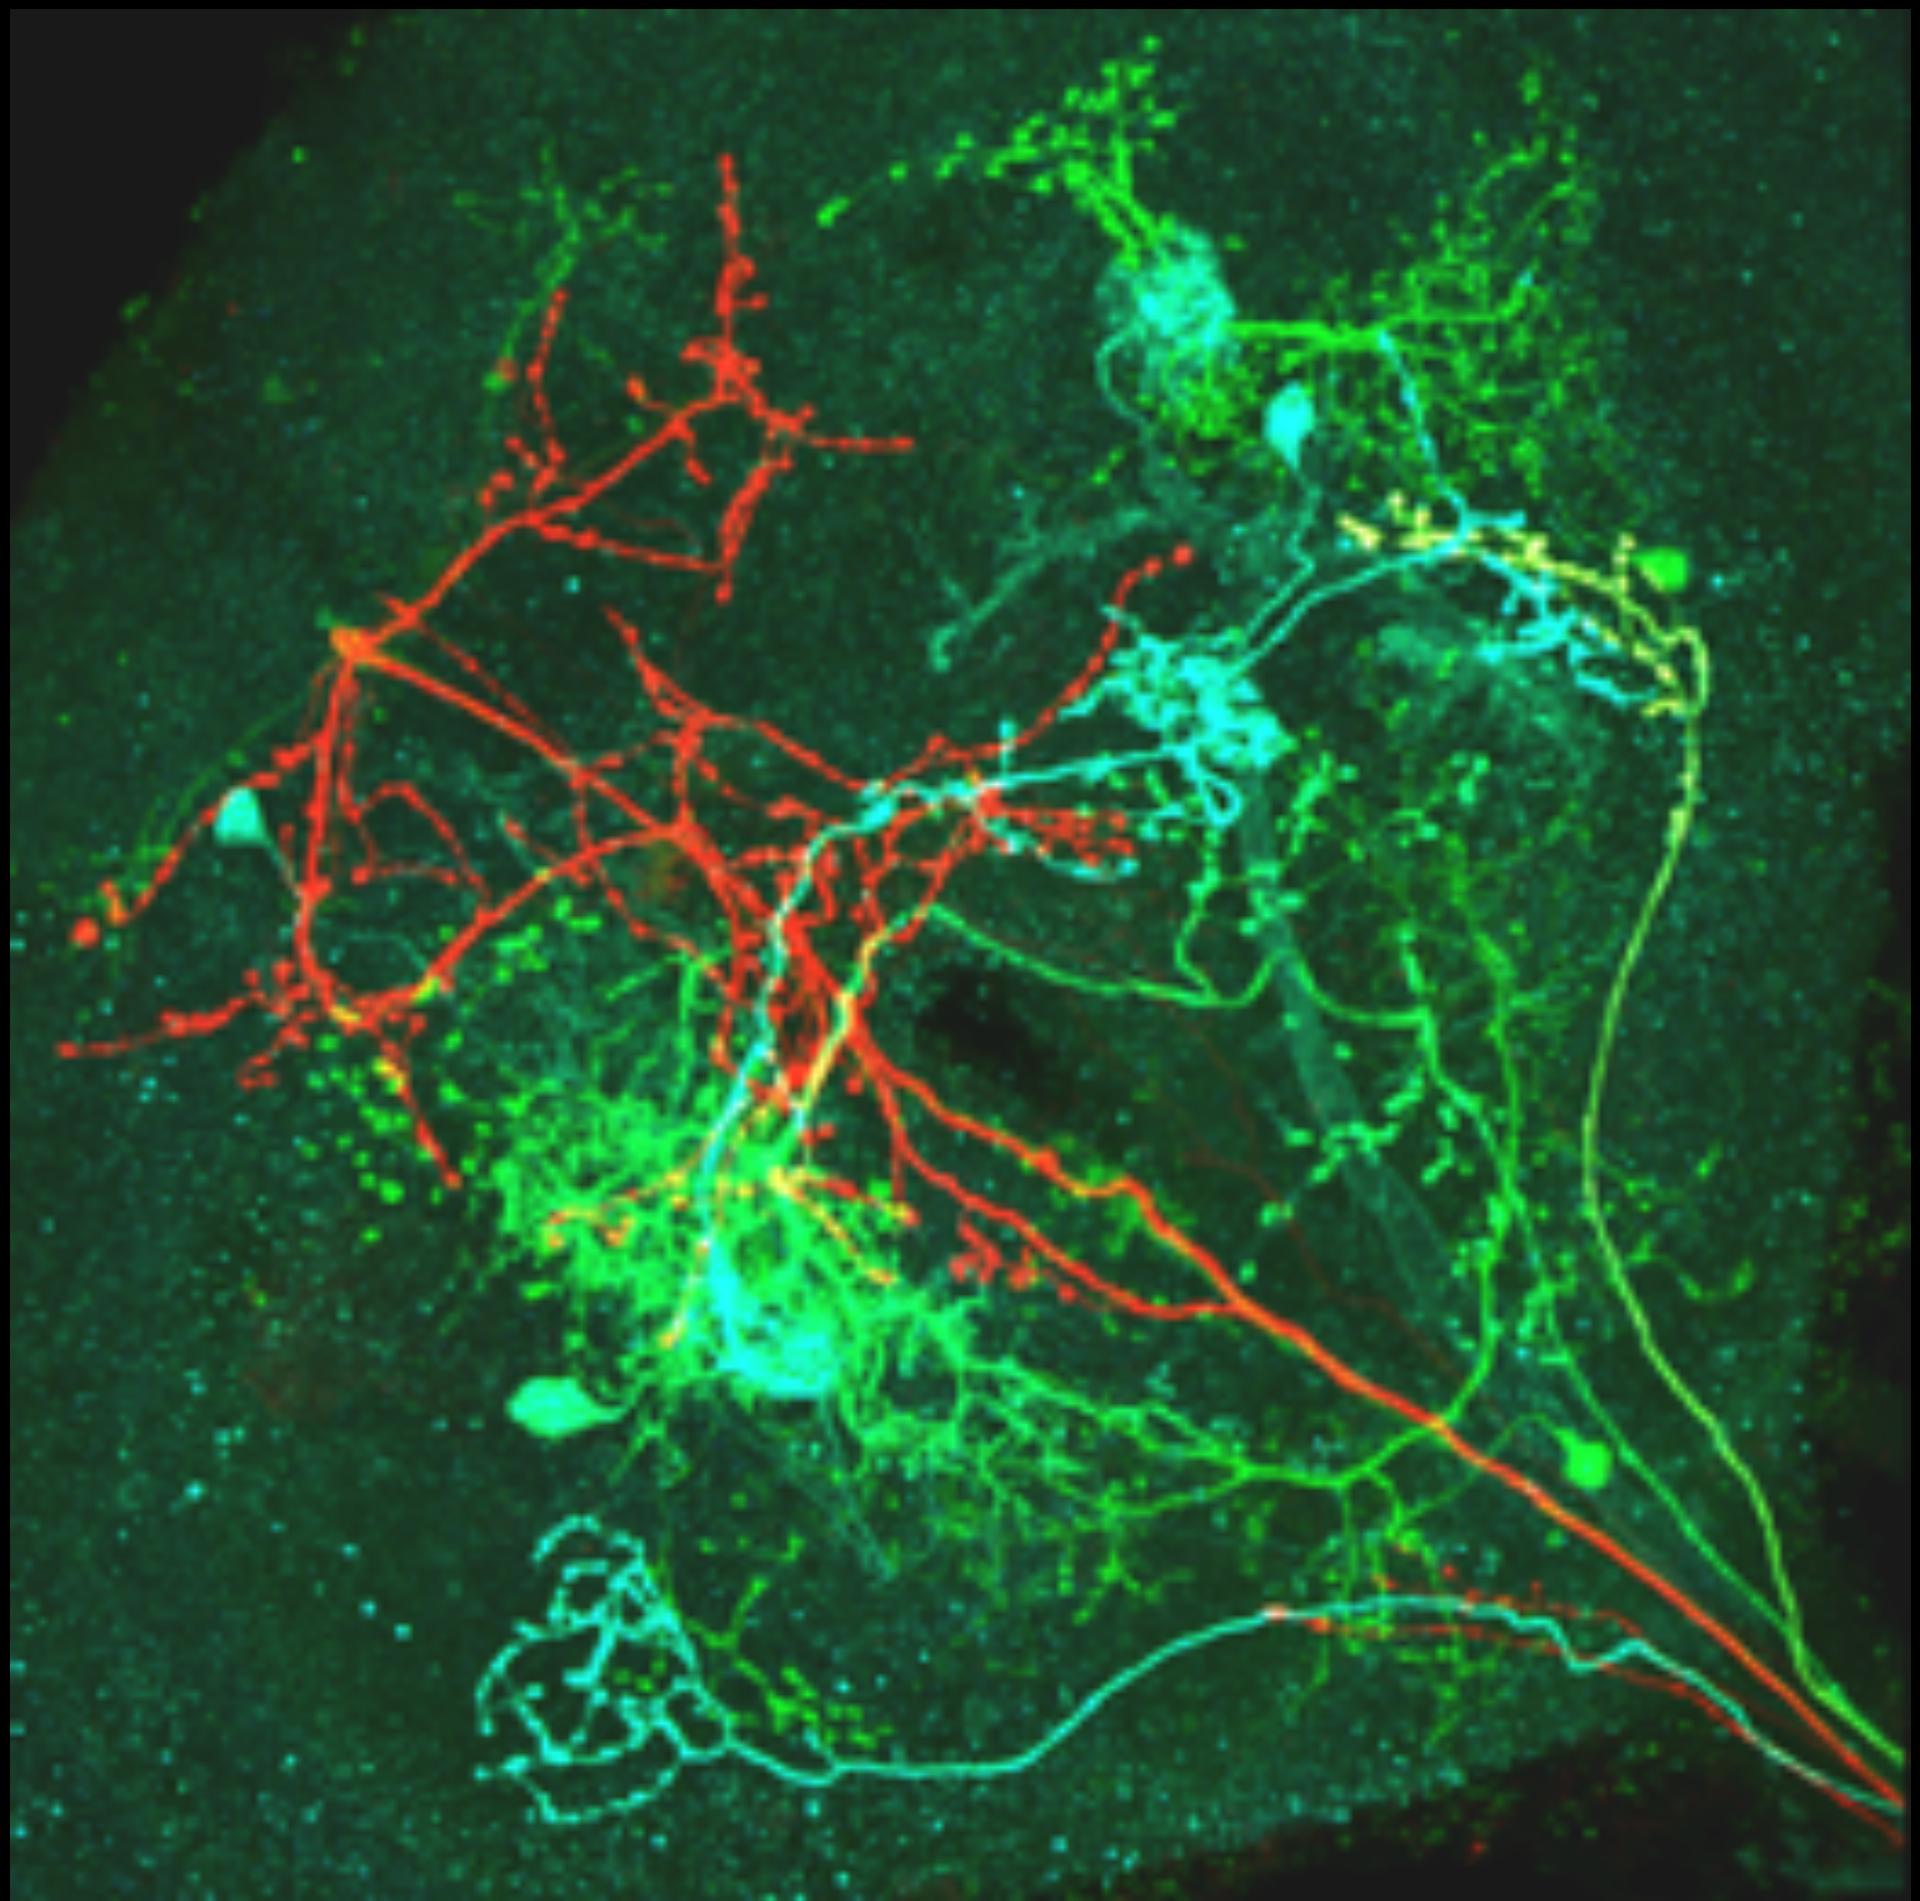
Sulston & Horvitz, Dev. Biol. 56 (1), 1977

Some more segmentation examples ...



Sulston & Horvitz, Dev. Biol. 56 (1), 1977

Some more segmentation examples ...



FlyLight Project, HHMI Janelia

Other Image-related Learning Tasks

- Object tracking in video
- Regression tasks on images:
image-wise, pixel-wise
- Image Generation
from other image, from text,
from set of conditions, ...
- Image Captioning



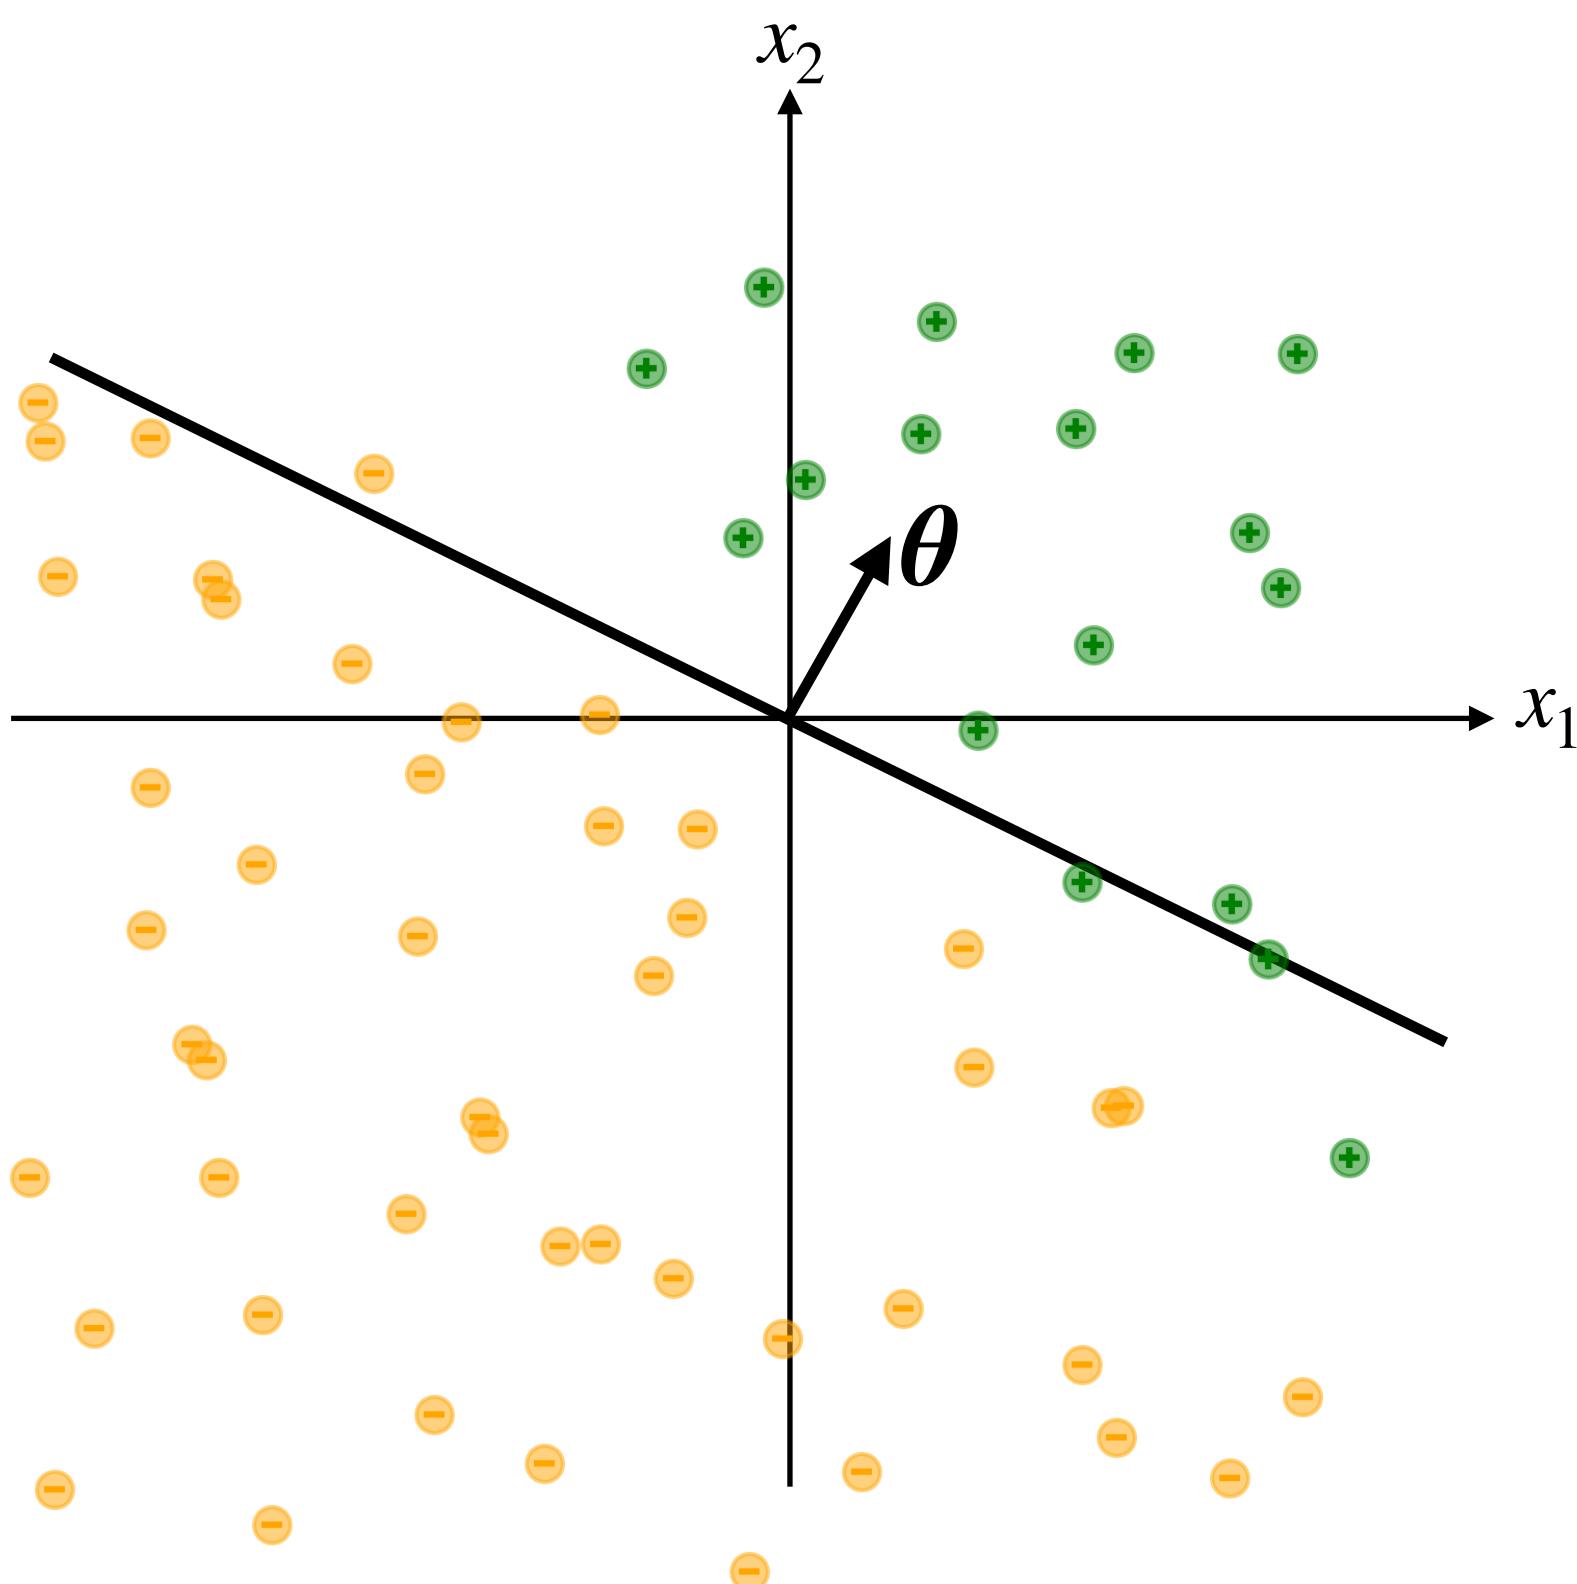
<https://stablediffusionweb.com/>

Basics: Neural Networks

- Multi-layer Perceptrons
- Backpropagation
- ... on Images

The Perceptron [McCulloch and Pitts, 1943]

The Perceptron

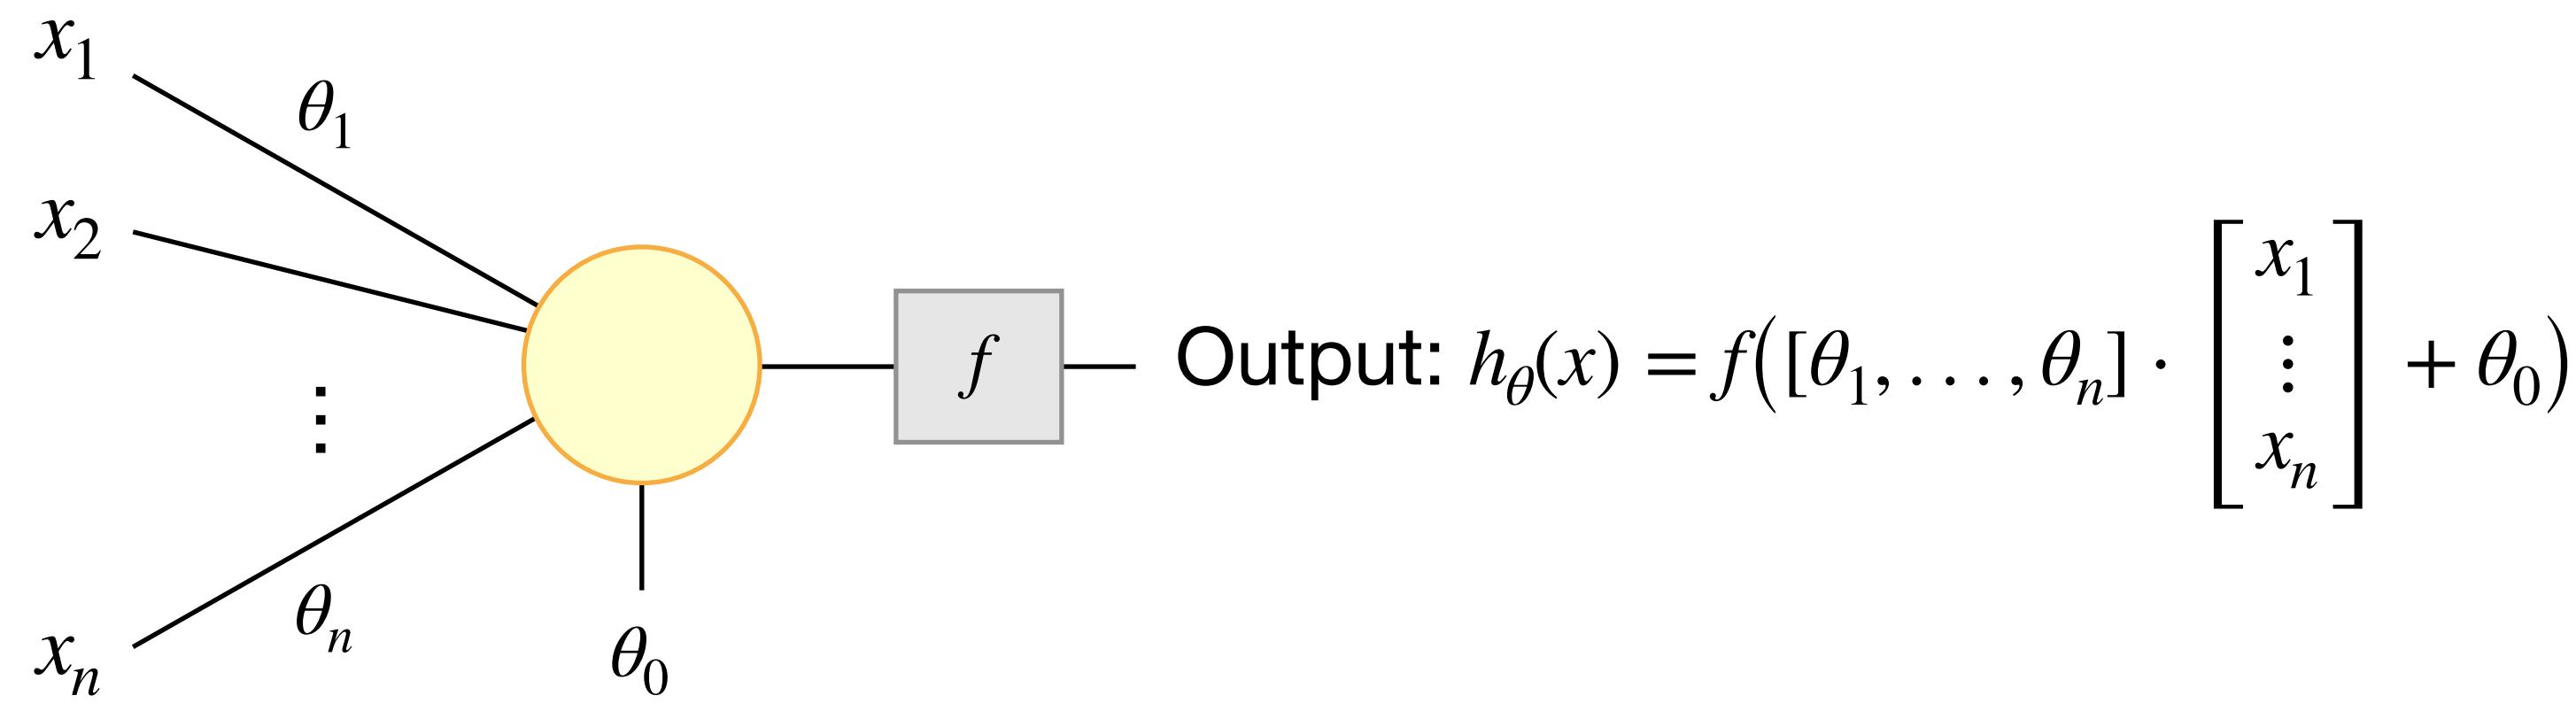


Algorithm : The Perceptron Learning Algorithm

```
 $P \leftarrow$  samples with label 1  
 $N \leftarrow$  samples with label  $-1$   
Initialize  $\theta$  randomly  
while !convergence do  
    pick random sample  $x$   
    if  $x \in P$  and  $\theta^T x < 0$  then  
         $\theta = \theta + x$   
    end if  
    if  $x \in N$  and  $\theta^T x \geq 0$  then  
         $\theta = \theta - x$   
    end if  
end while  
// the algorithm converges when all samples are classified correctly
```

The Perceptron

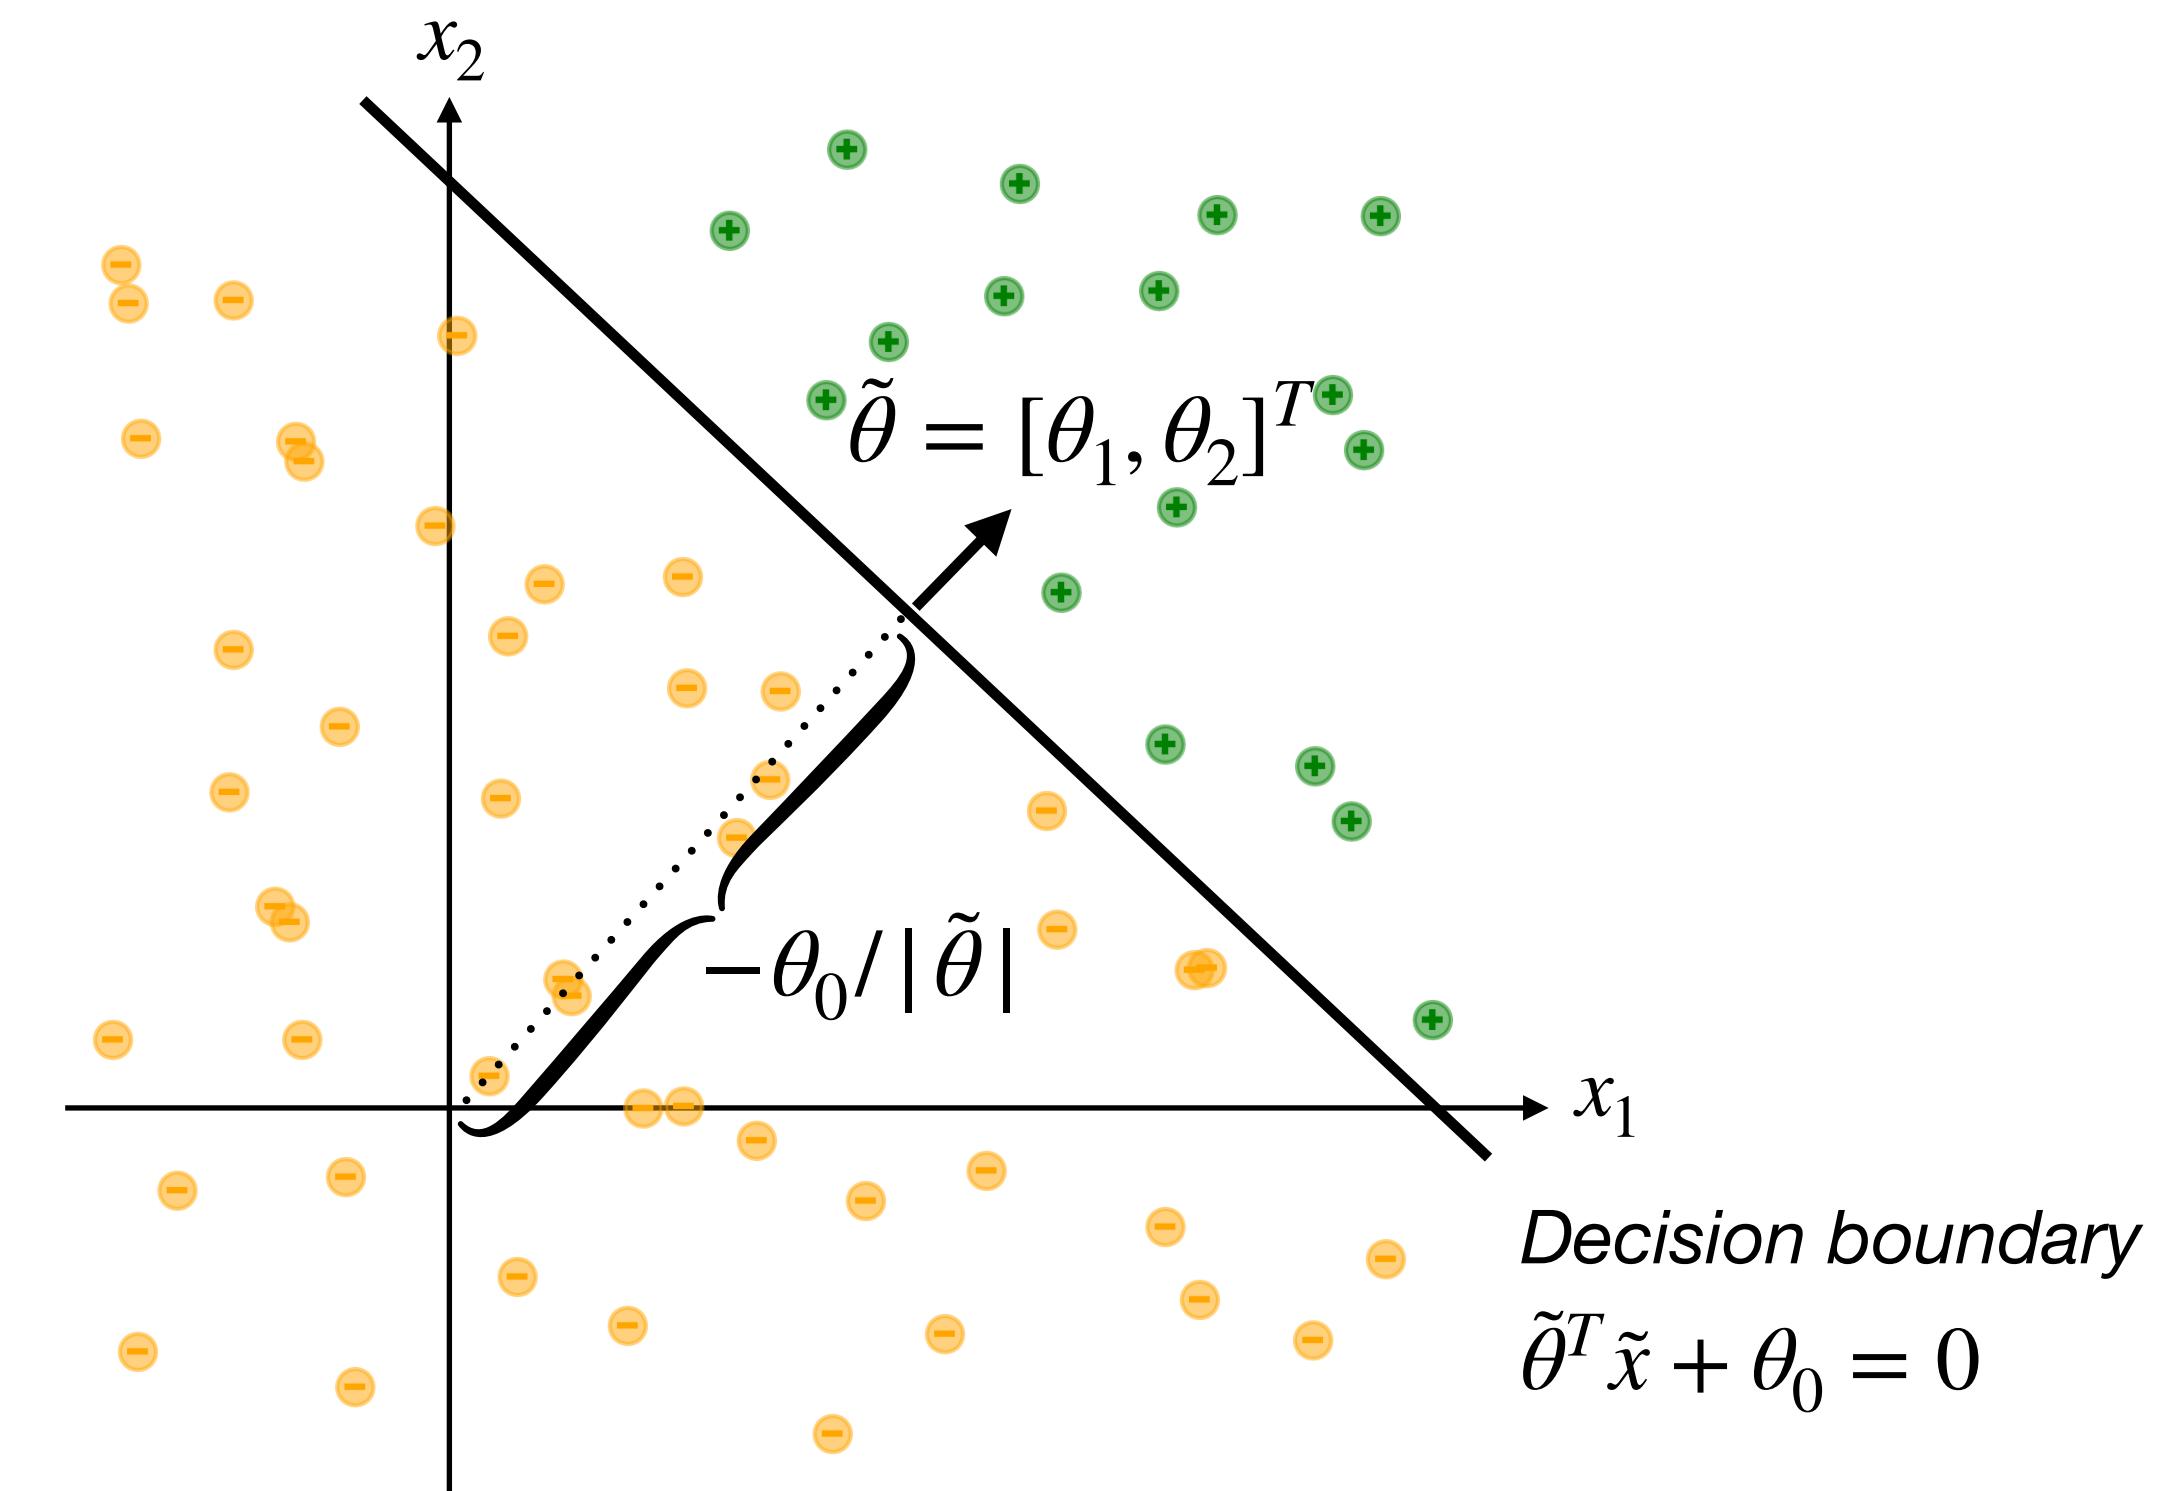
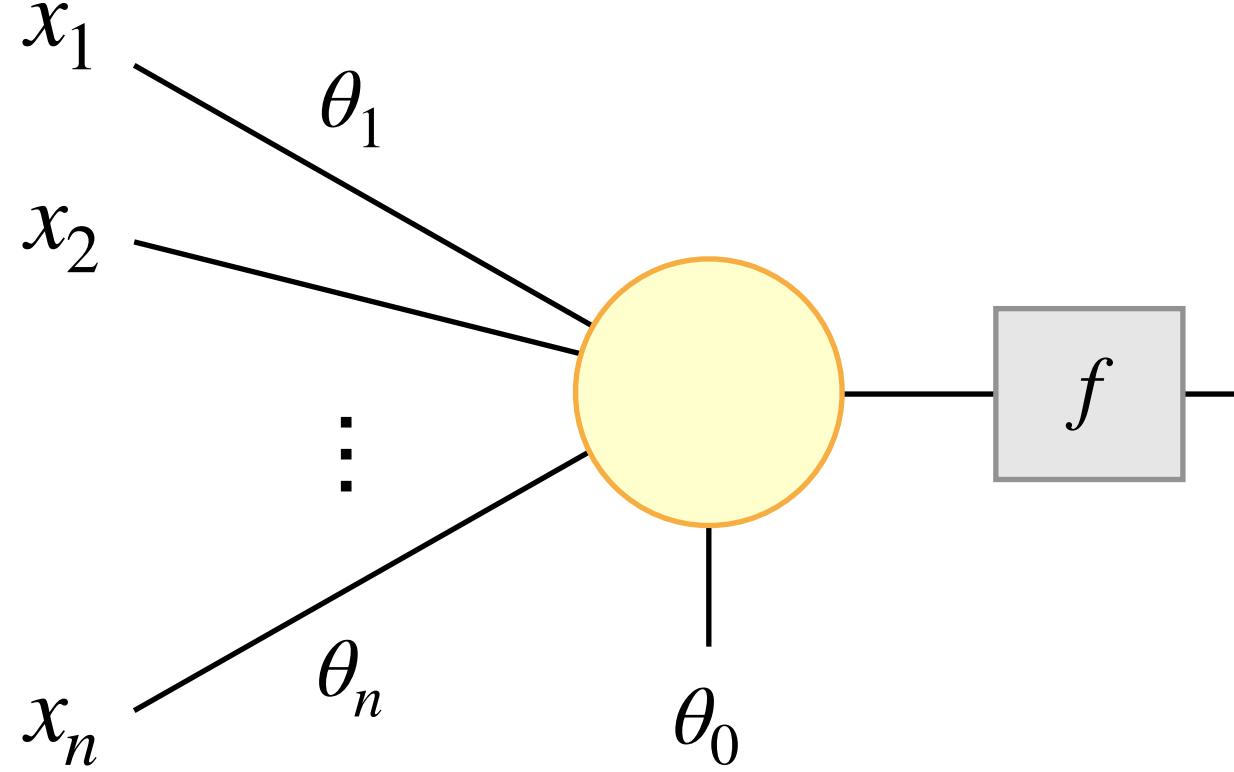
Linear classifier, building block of neural networks



Input: $\tilde{x} = [x_1, \dots, x_n]^T \in \mathbb{R}^n$
Label: $y = \pm 1$
Weights: $\tilde{\theta} = [\theta_1, \dots, \theta_n]^T$
Bias: θ_0
Activation func: 1) $f(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ -1 & \text{if } x < 0 \end{cases}$
equivalent 2) $f(x) = \tanh(x)$
3) $f(x) = \frac{1}{1 + e^{-x}}$

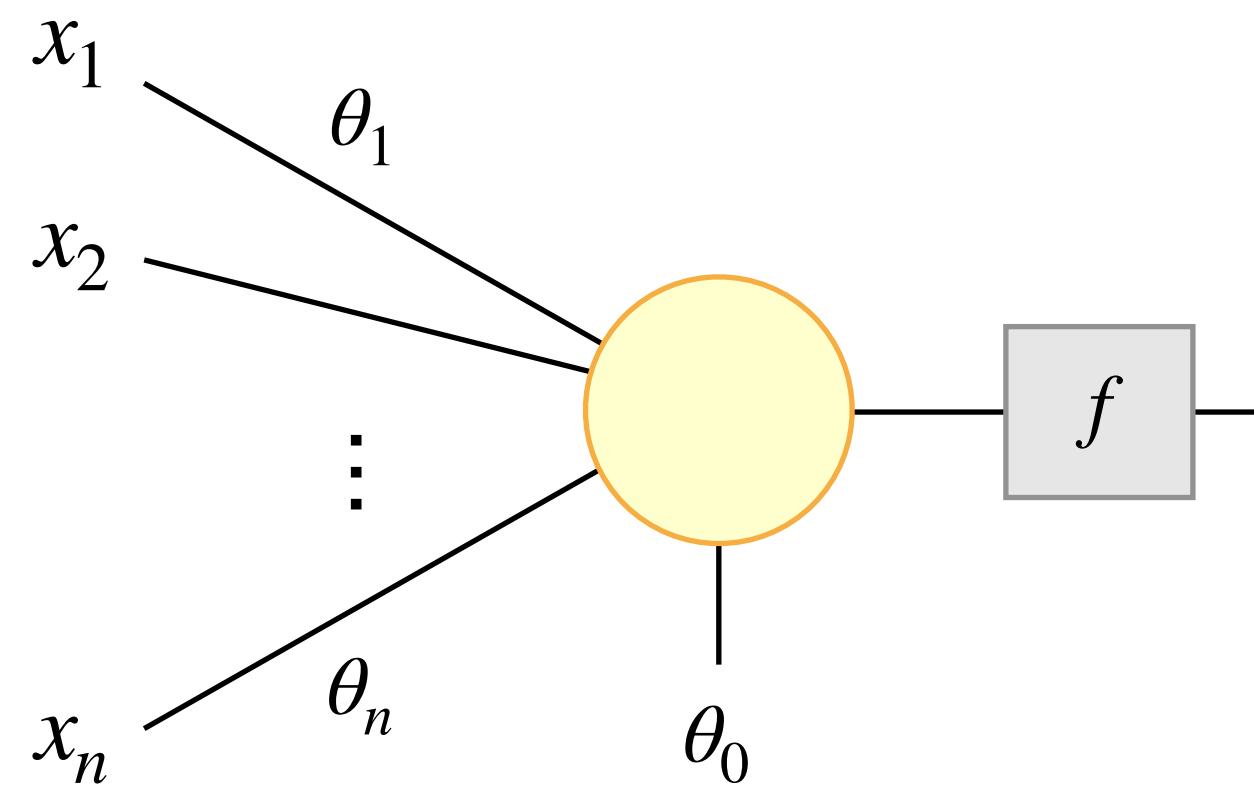
The Perceptron

Linear classifier, building block of neural networks



The Perceptron

Linear classifier, building block of neural networks

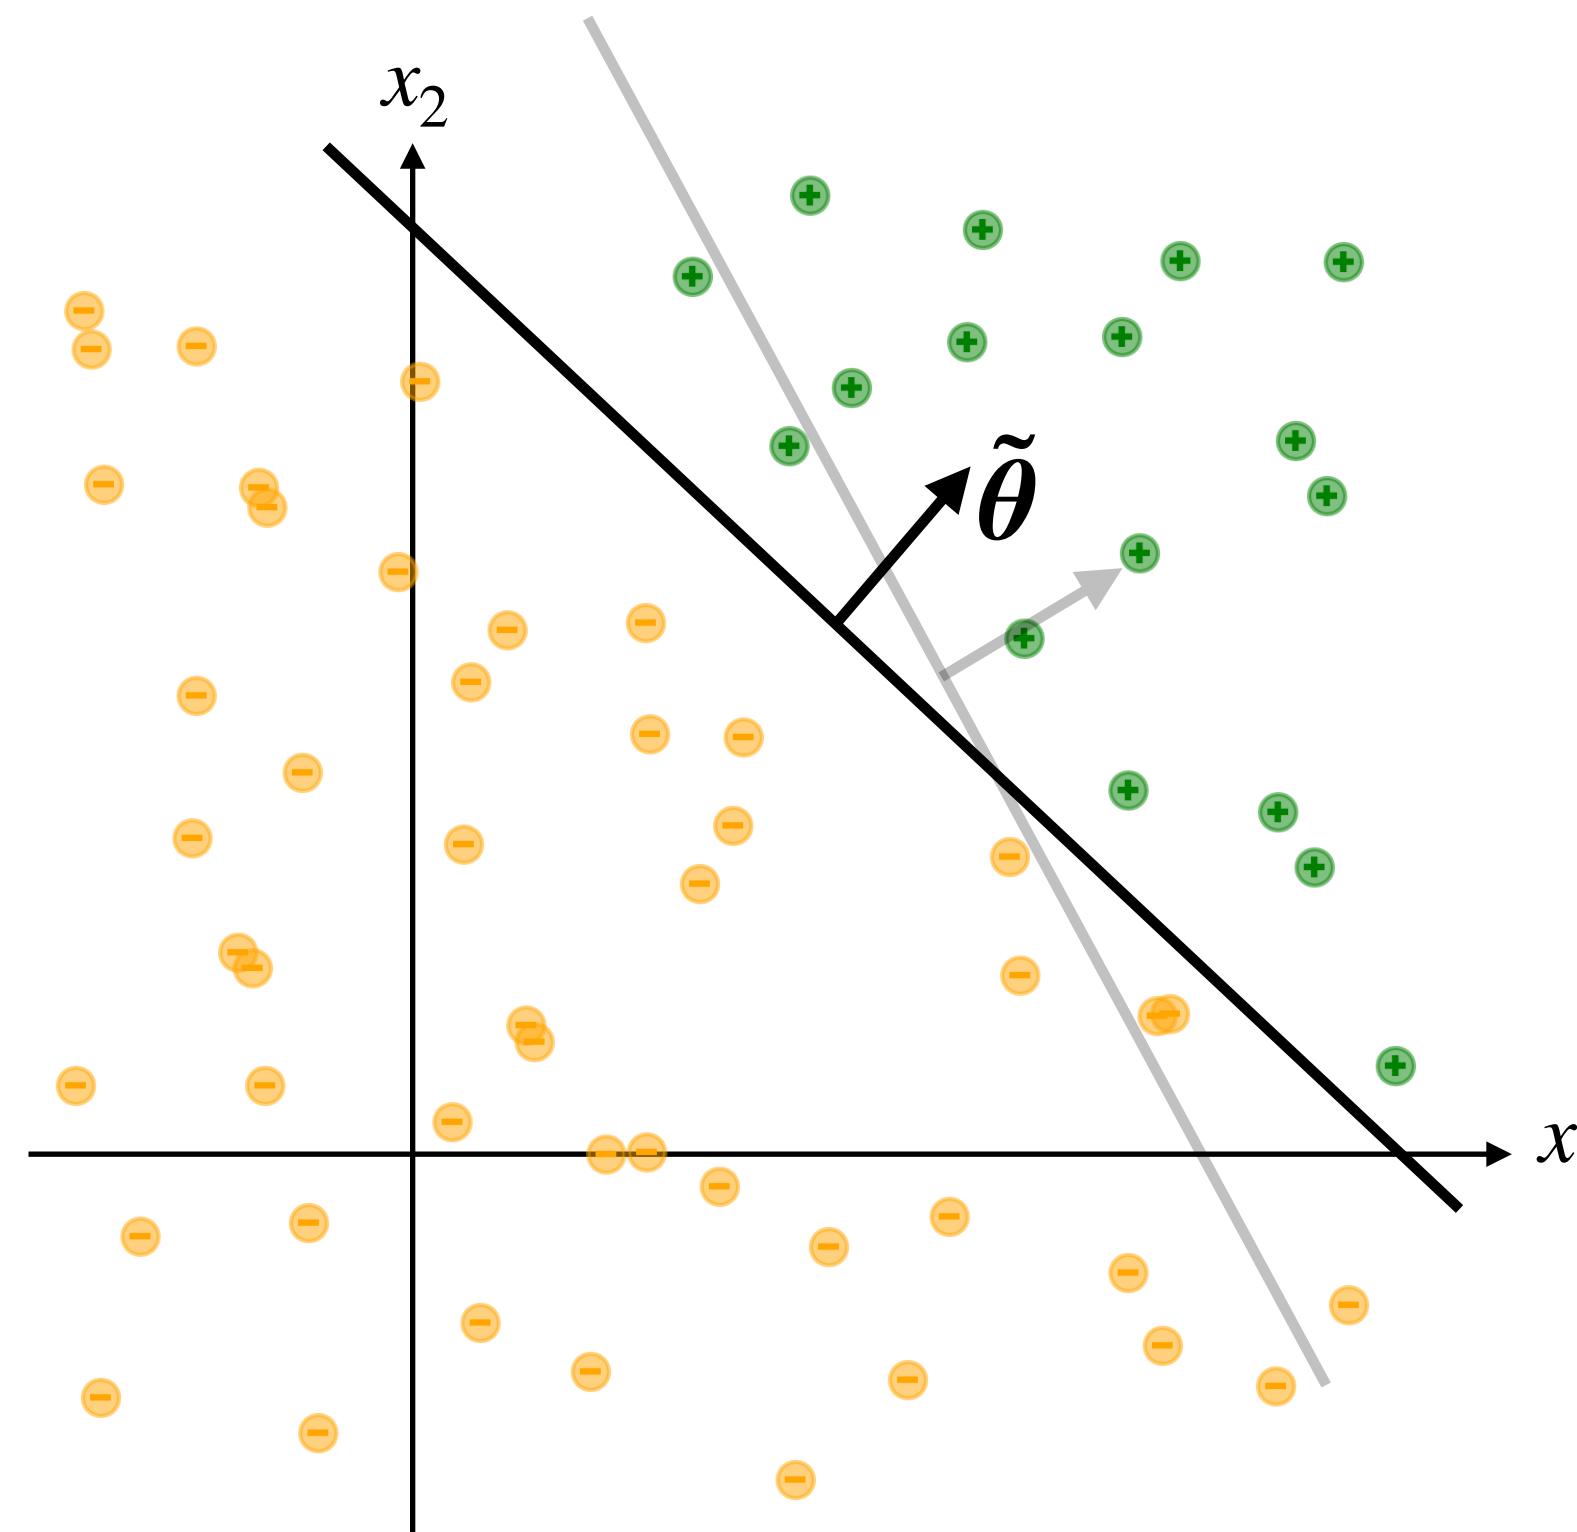


Recap: Concise notation in **homogeneous coordinates**

$$\begin{aligned} \text{Output: } h_\theta(x) &= f([\theta_1, \dots, \theta_n] \cdot [x_1, \dots, x_n]^T + \theta_0) =: f(\tilde{\theta}^T \tilde{x} + \theta_0) \\ &= f([\theta_0, \theta_1, \dots, \theta_n] \cdot [1, x_1, \dots, x_n]^T) \\ &=: f(\boldsymbol{\theta}^T \mathbf{x}) \end{aligned}$$

Loss function (aka cost function)

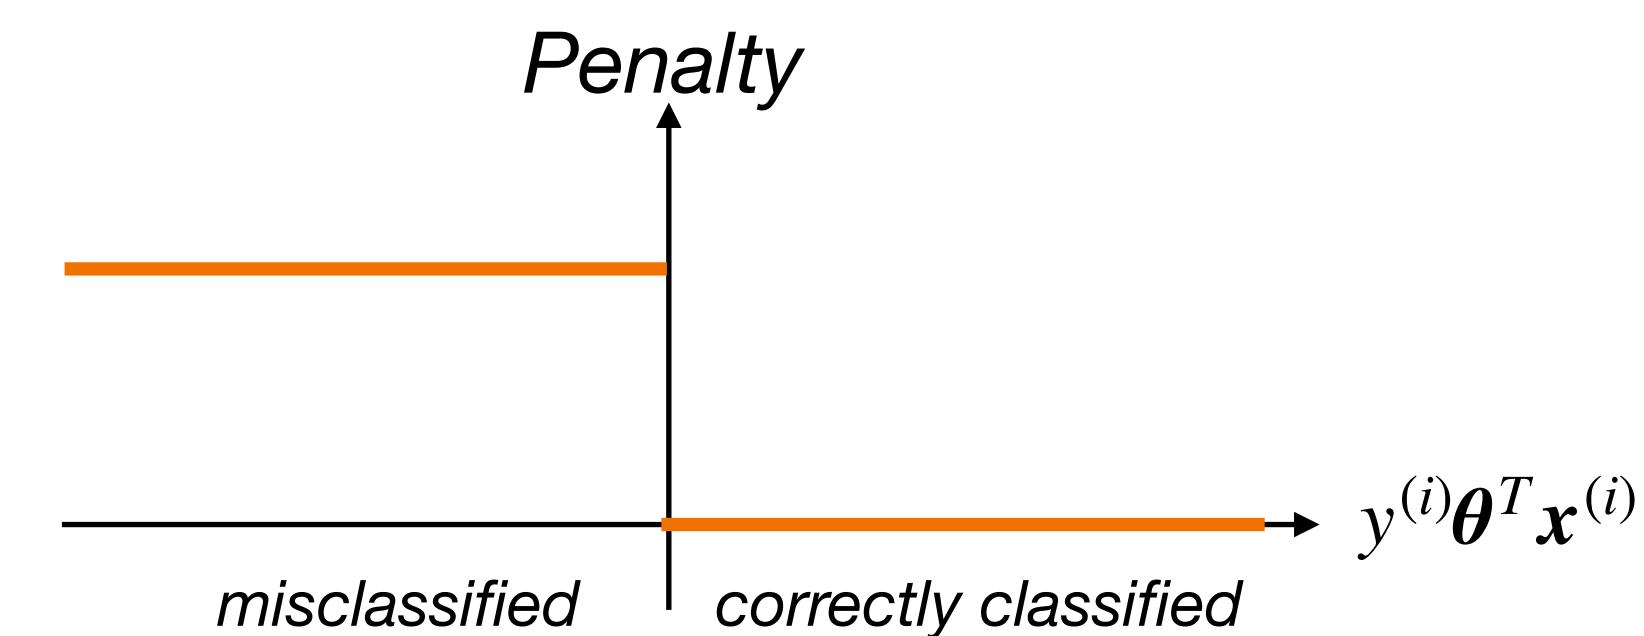
Goal: Find an optimal solution



A) The 0-1 loss:

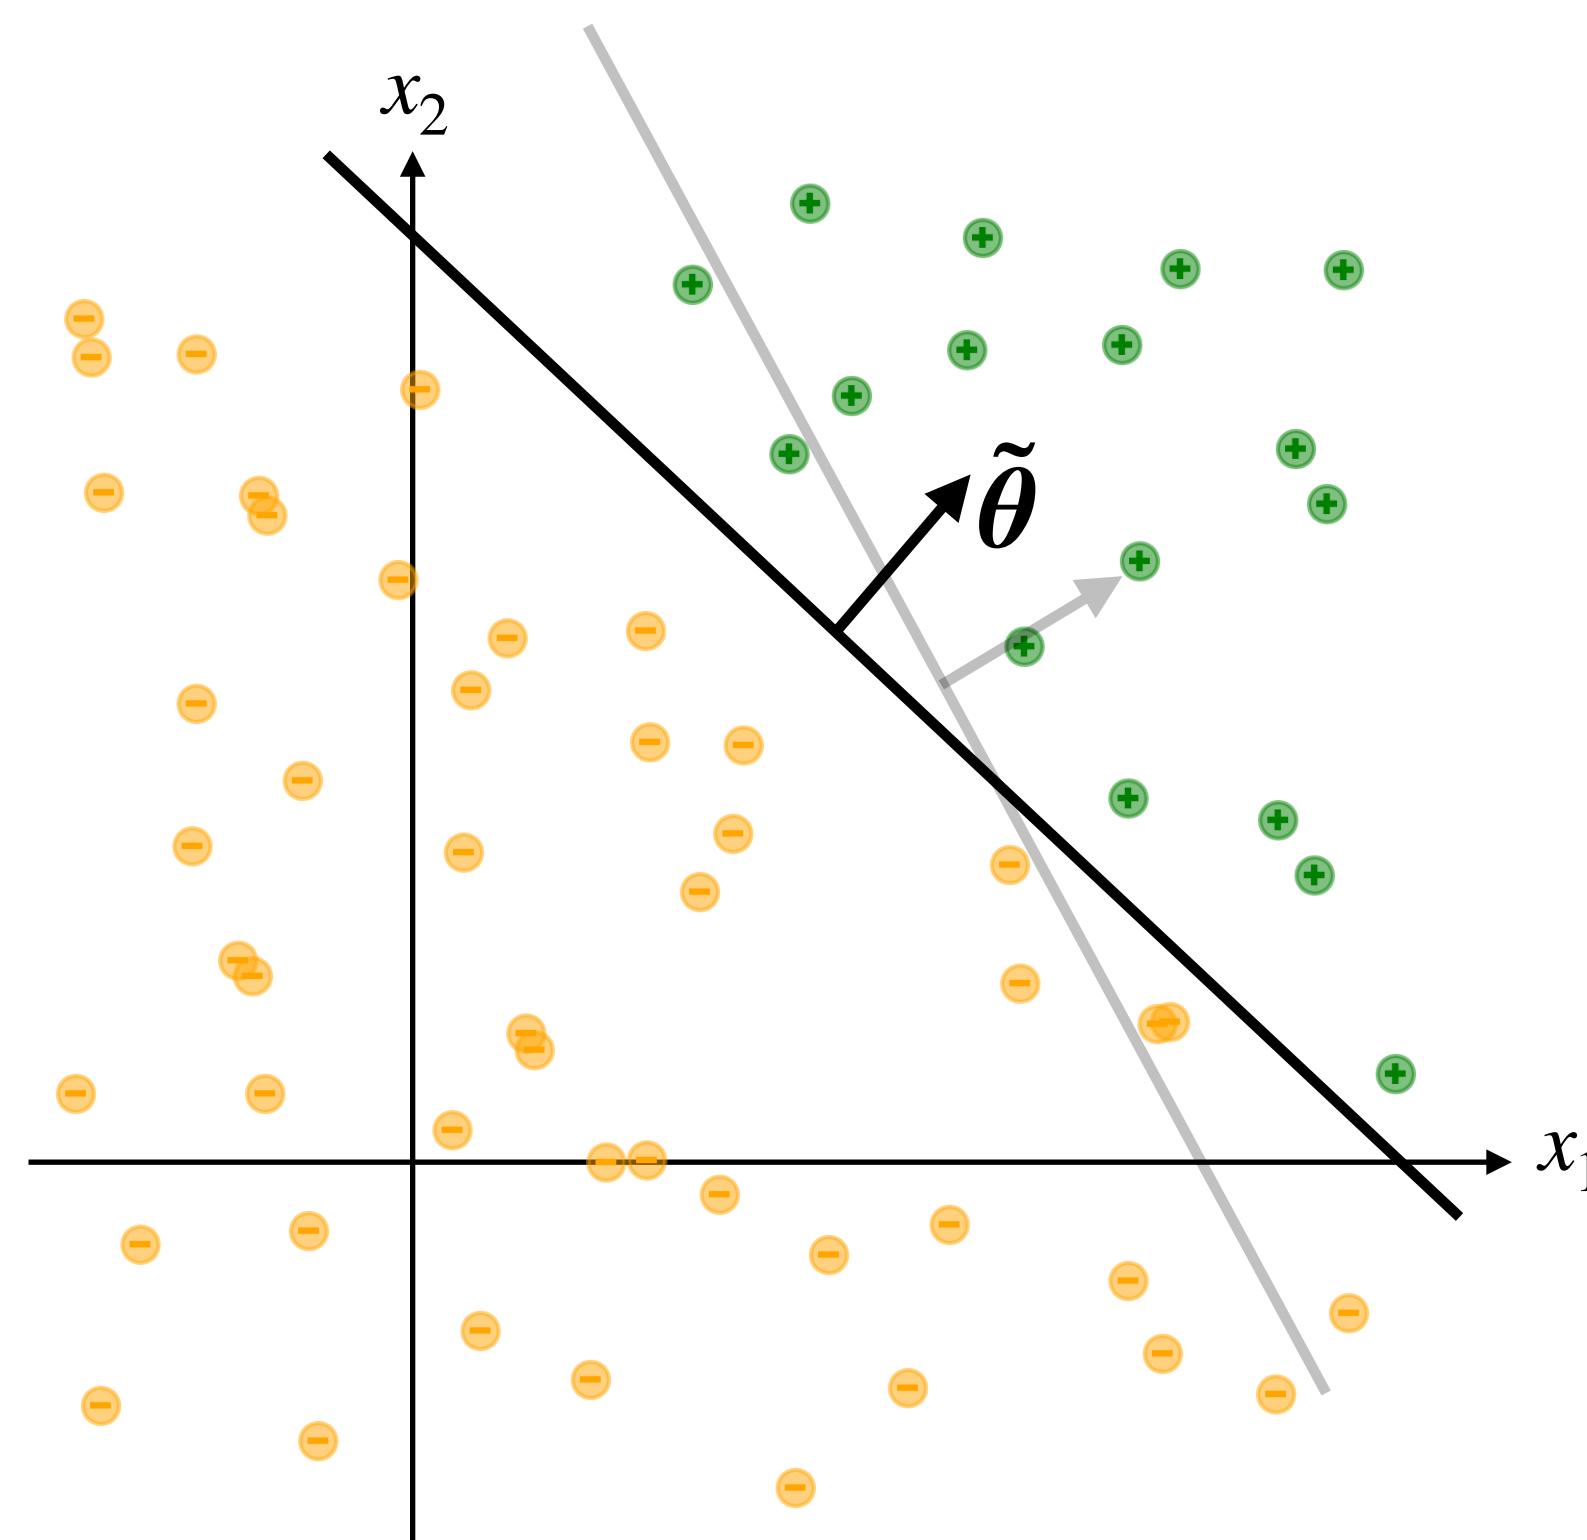
Heaviside function

$$\varepsilon(\theta) = \sum_{i=1}^m H(-y^{(i)}\theta^T x^{(i)})$$

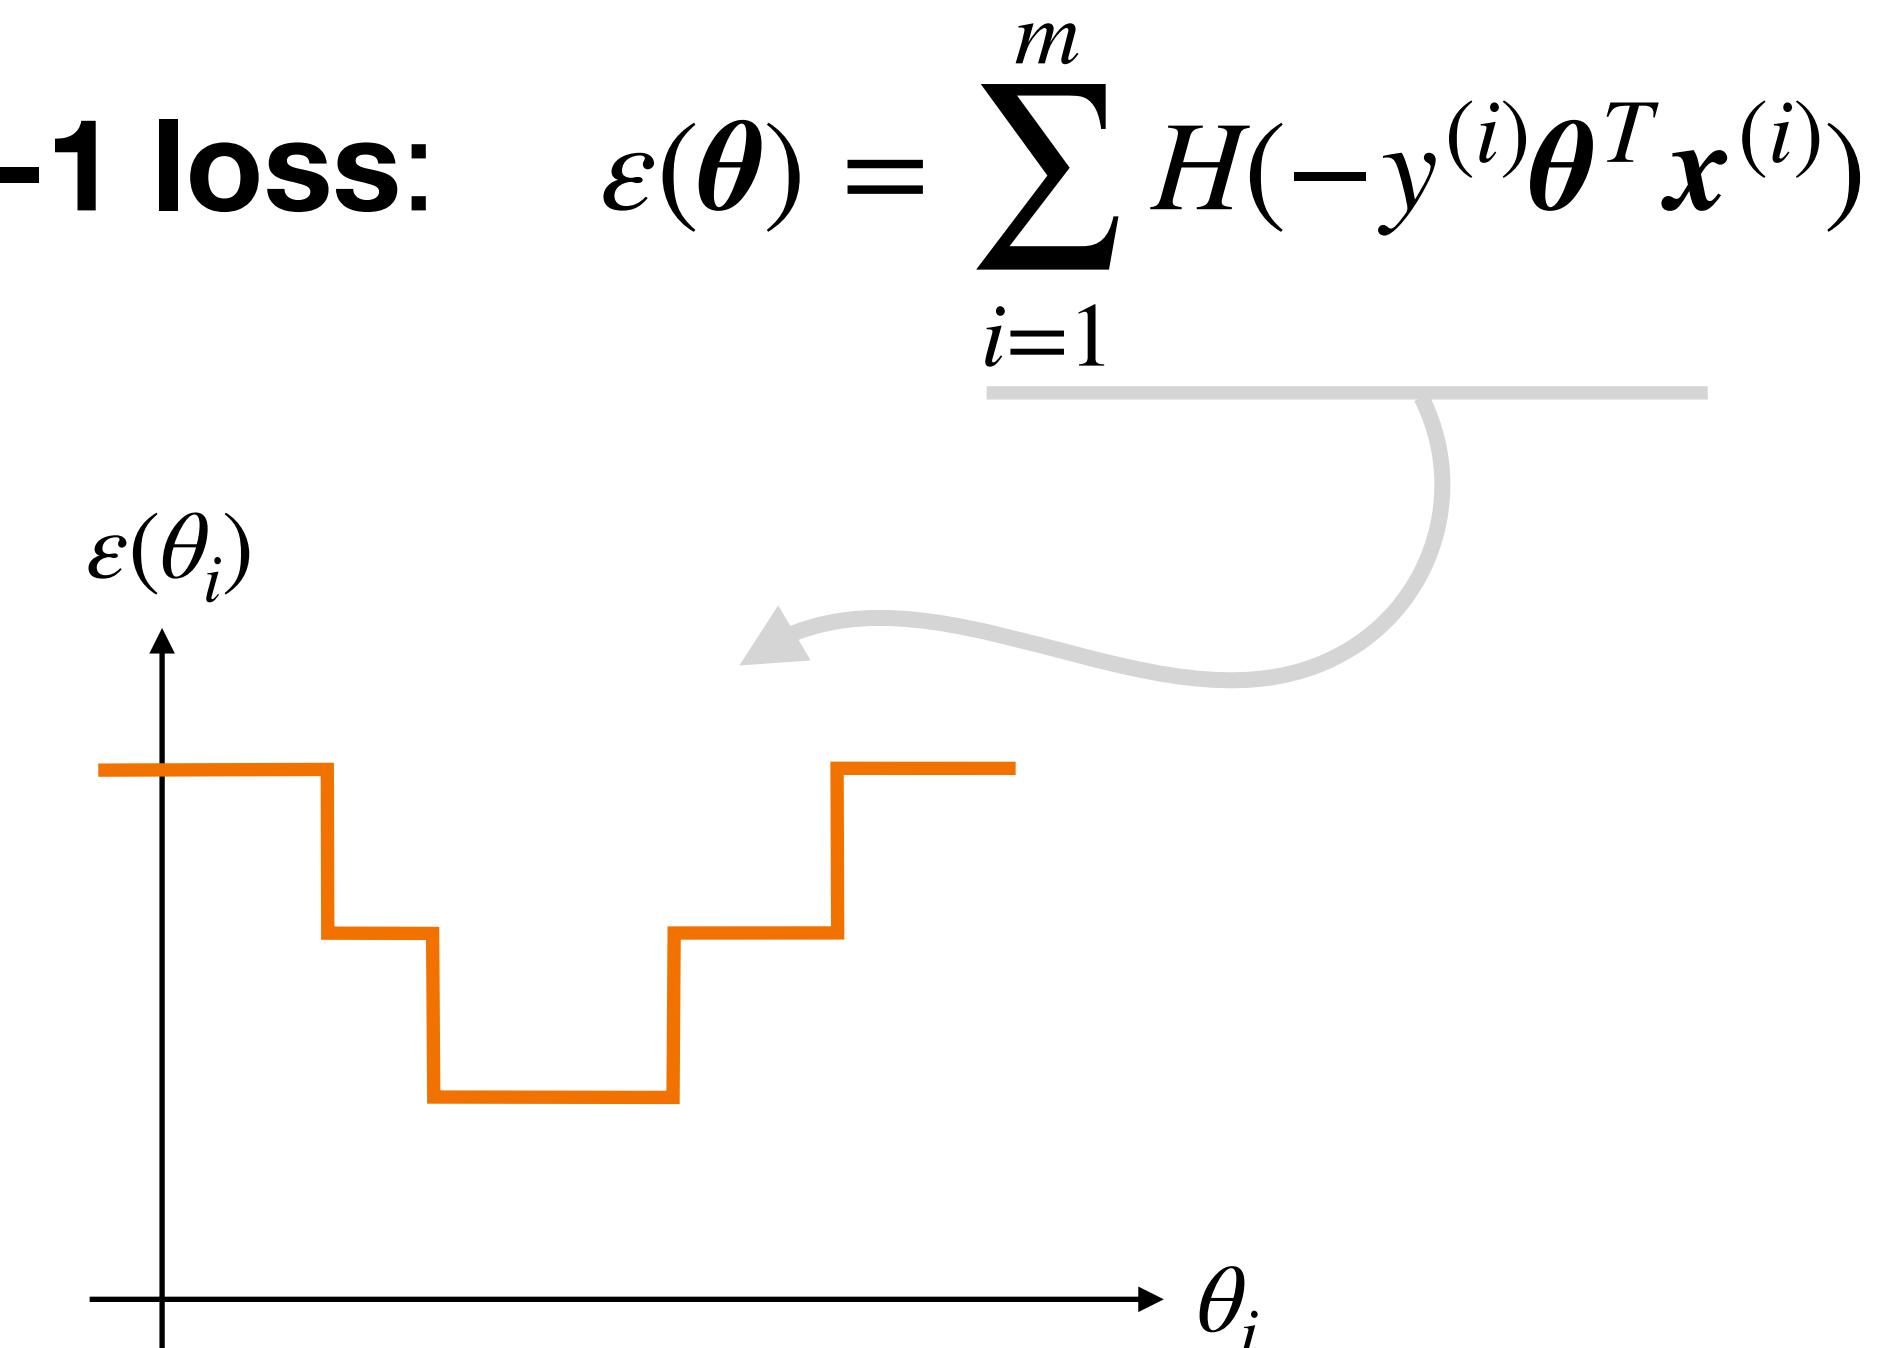


Loss function

Goal: Find an optimal solution

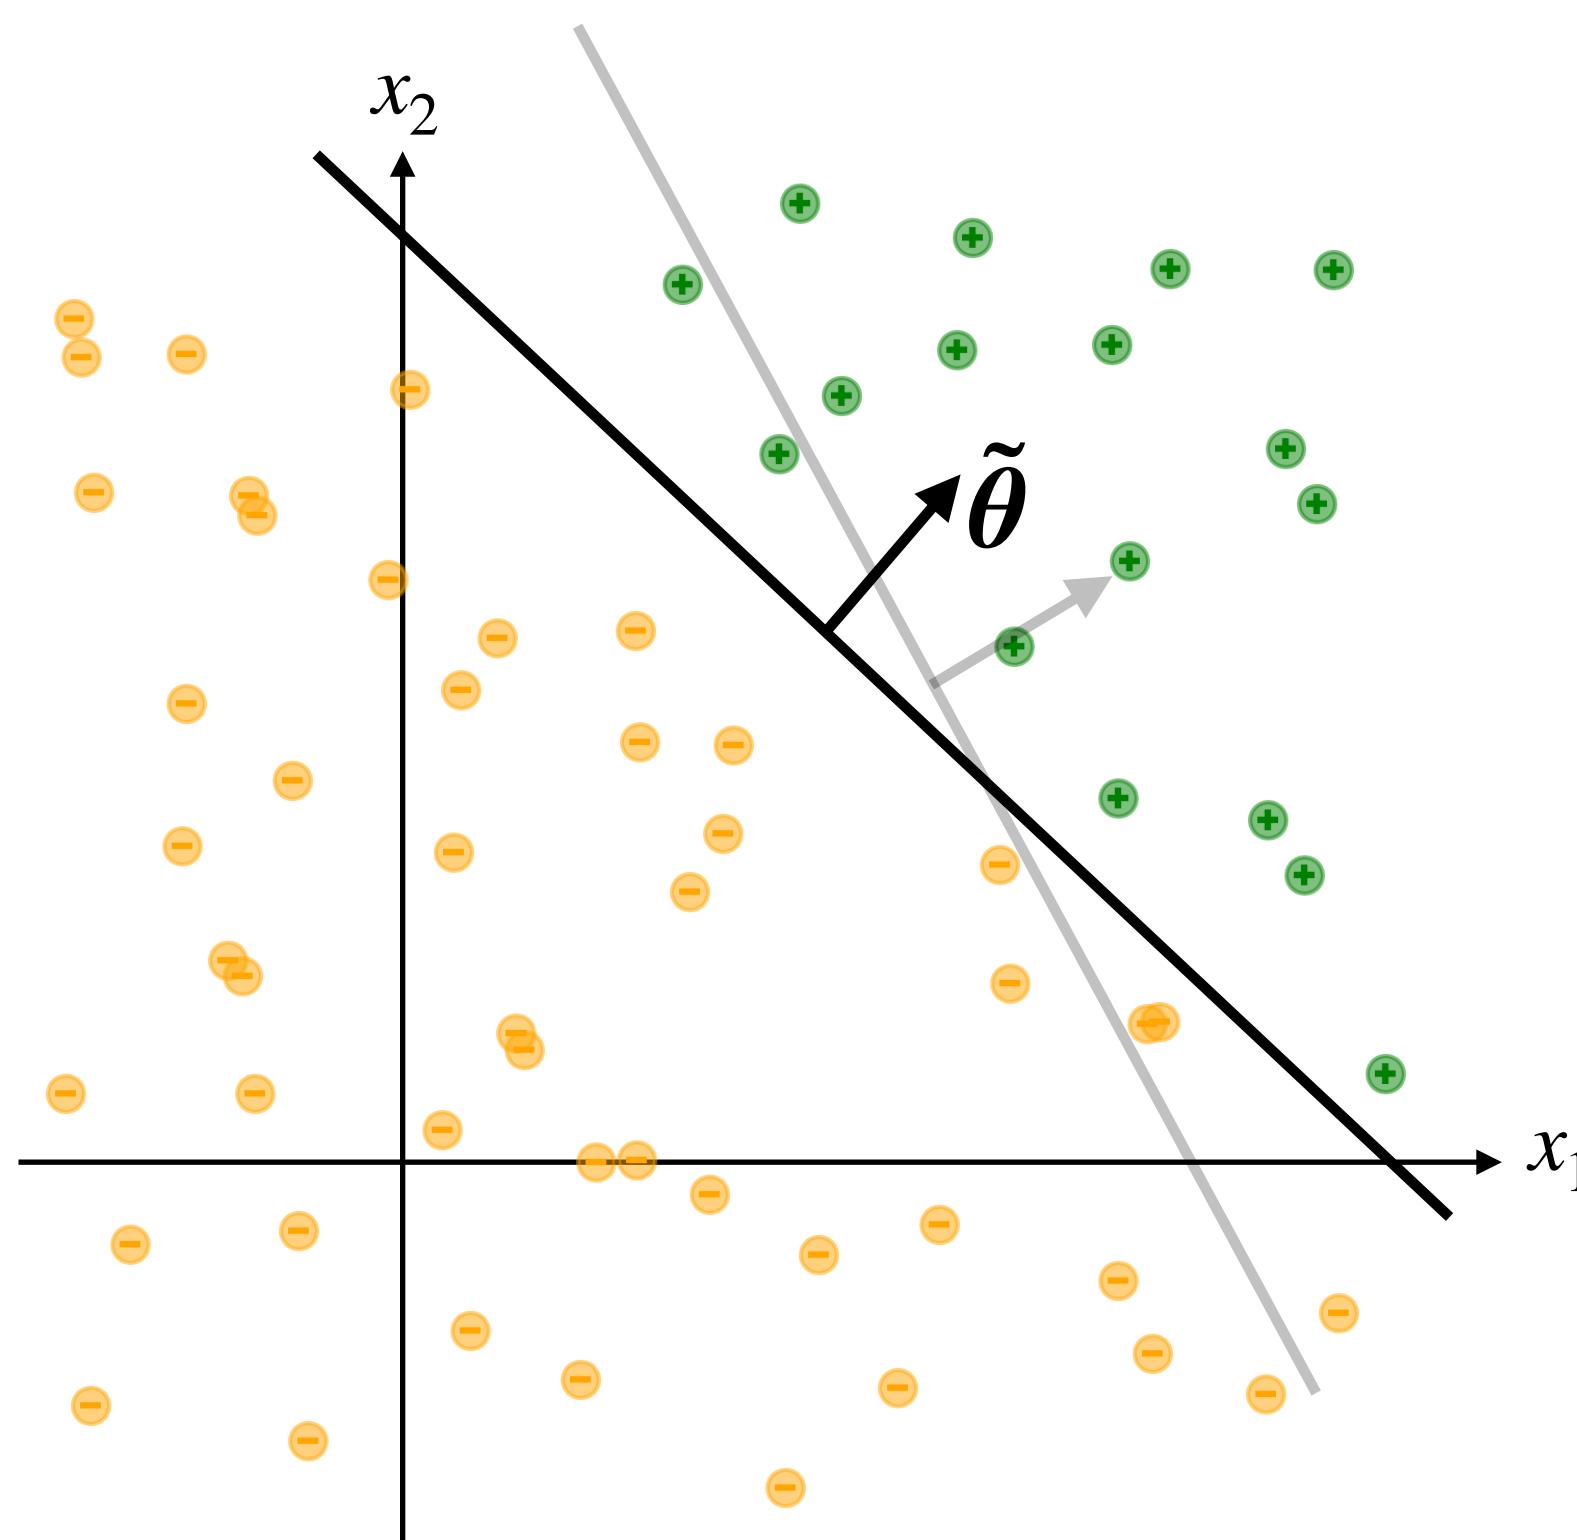


A) The 0-1 loss: $\varepsilon(\theta) = \sum_{i=1}^m H(-y^{(i)}\theta^T x^{(i)})$



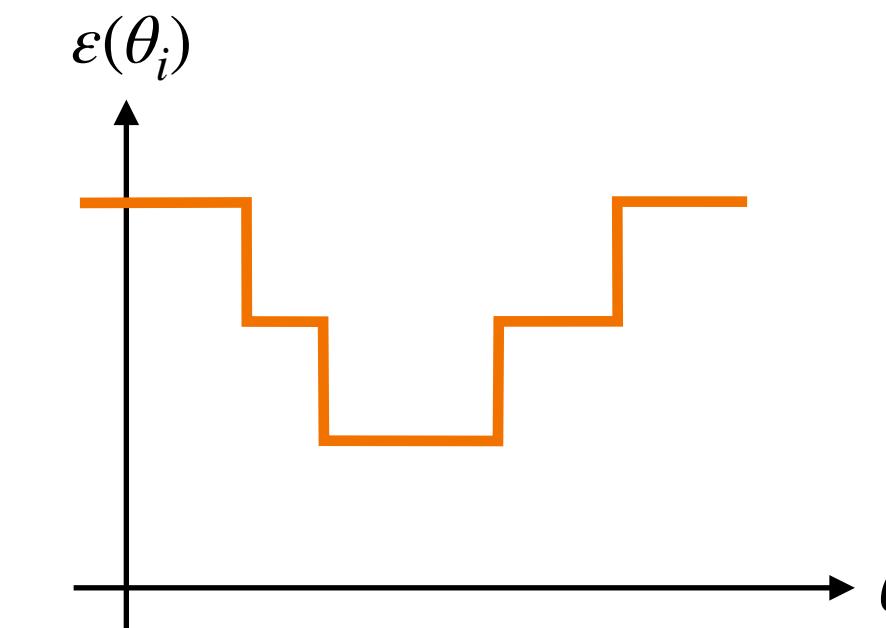
Loss function

Goal: Find an optimal solution



A) The 0-1 loss: $\varepsilon(\theta) = \sum_{i=1}^m H(-y^{(i)}\theta^T x^{(i)})$

- We can't perform gradient descent
- Not convex
- Ambiguity of "optimal" solution



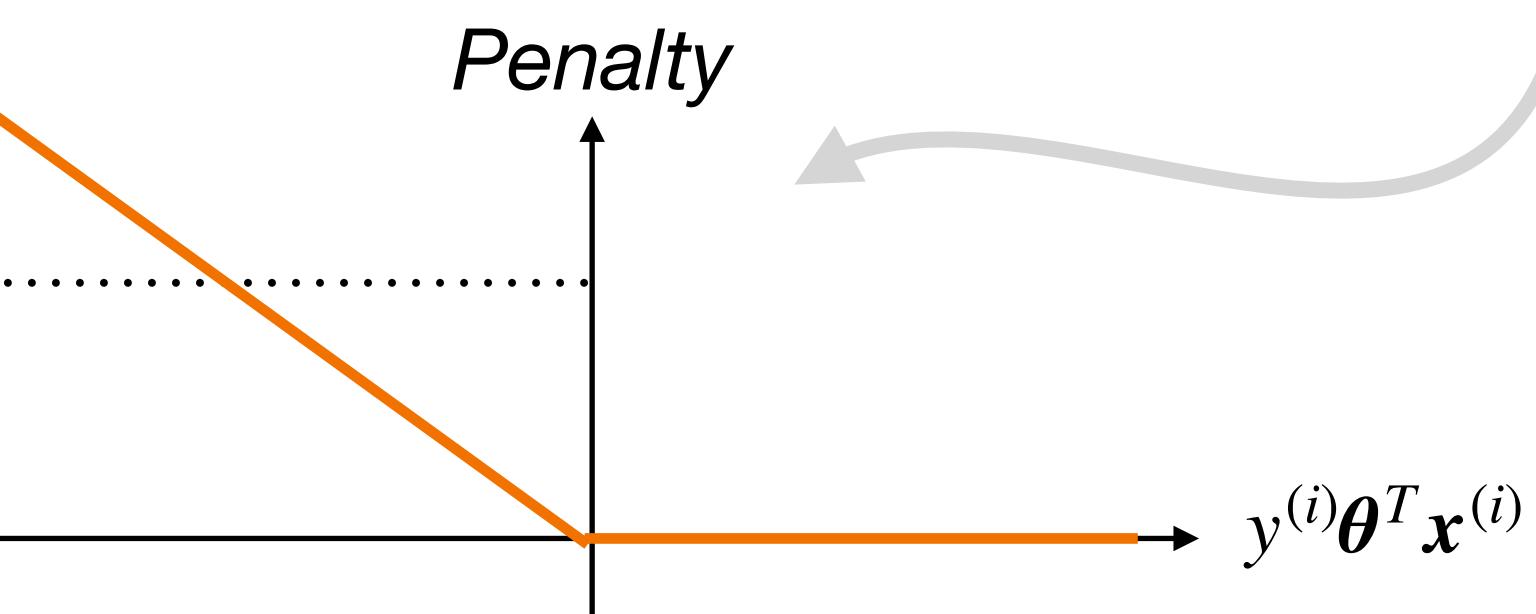
Loss function

Goal: Find an optimal solution

Algorithm : The Perceptron Learning Algorithm

```
P ← samples with label 1  
N ← samples with label 0  
Initialize  $\theta$  randomly  
while !convergence do  
    pick random sample  $x$   
    if  $x \in P$  and  $\theta^T x < 0$  then  
         $\theta = \theta + x$   
    end if  
    if  $x \in N$  and  $\theta^T x \geq 0$  then  
         $\theta = \theta - x$   
    end if  
end while  
// the algorithm converges when all samples are classified correctly
```

B) The hinge loss: $\varepsilon(\theta) = \sum_{i=1}^m \max(-y^{(i)}\theta^T x^{(i)}, 0)$

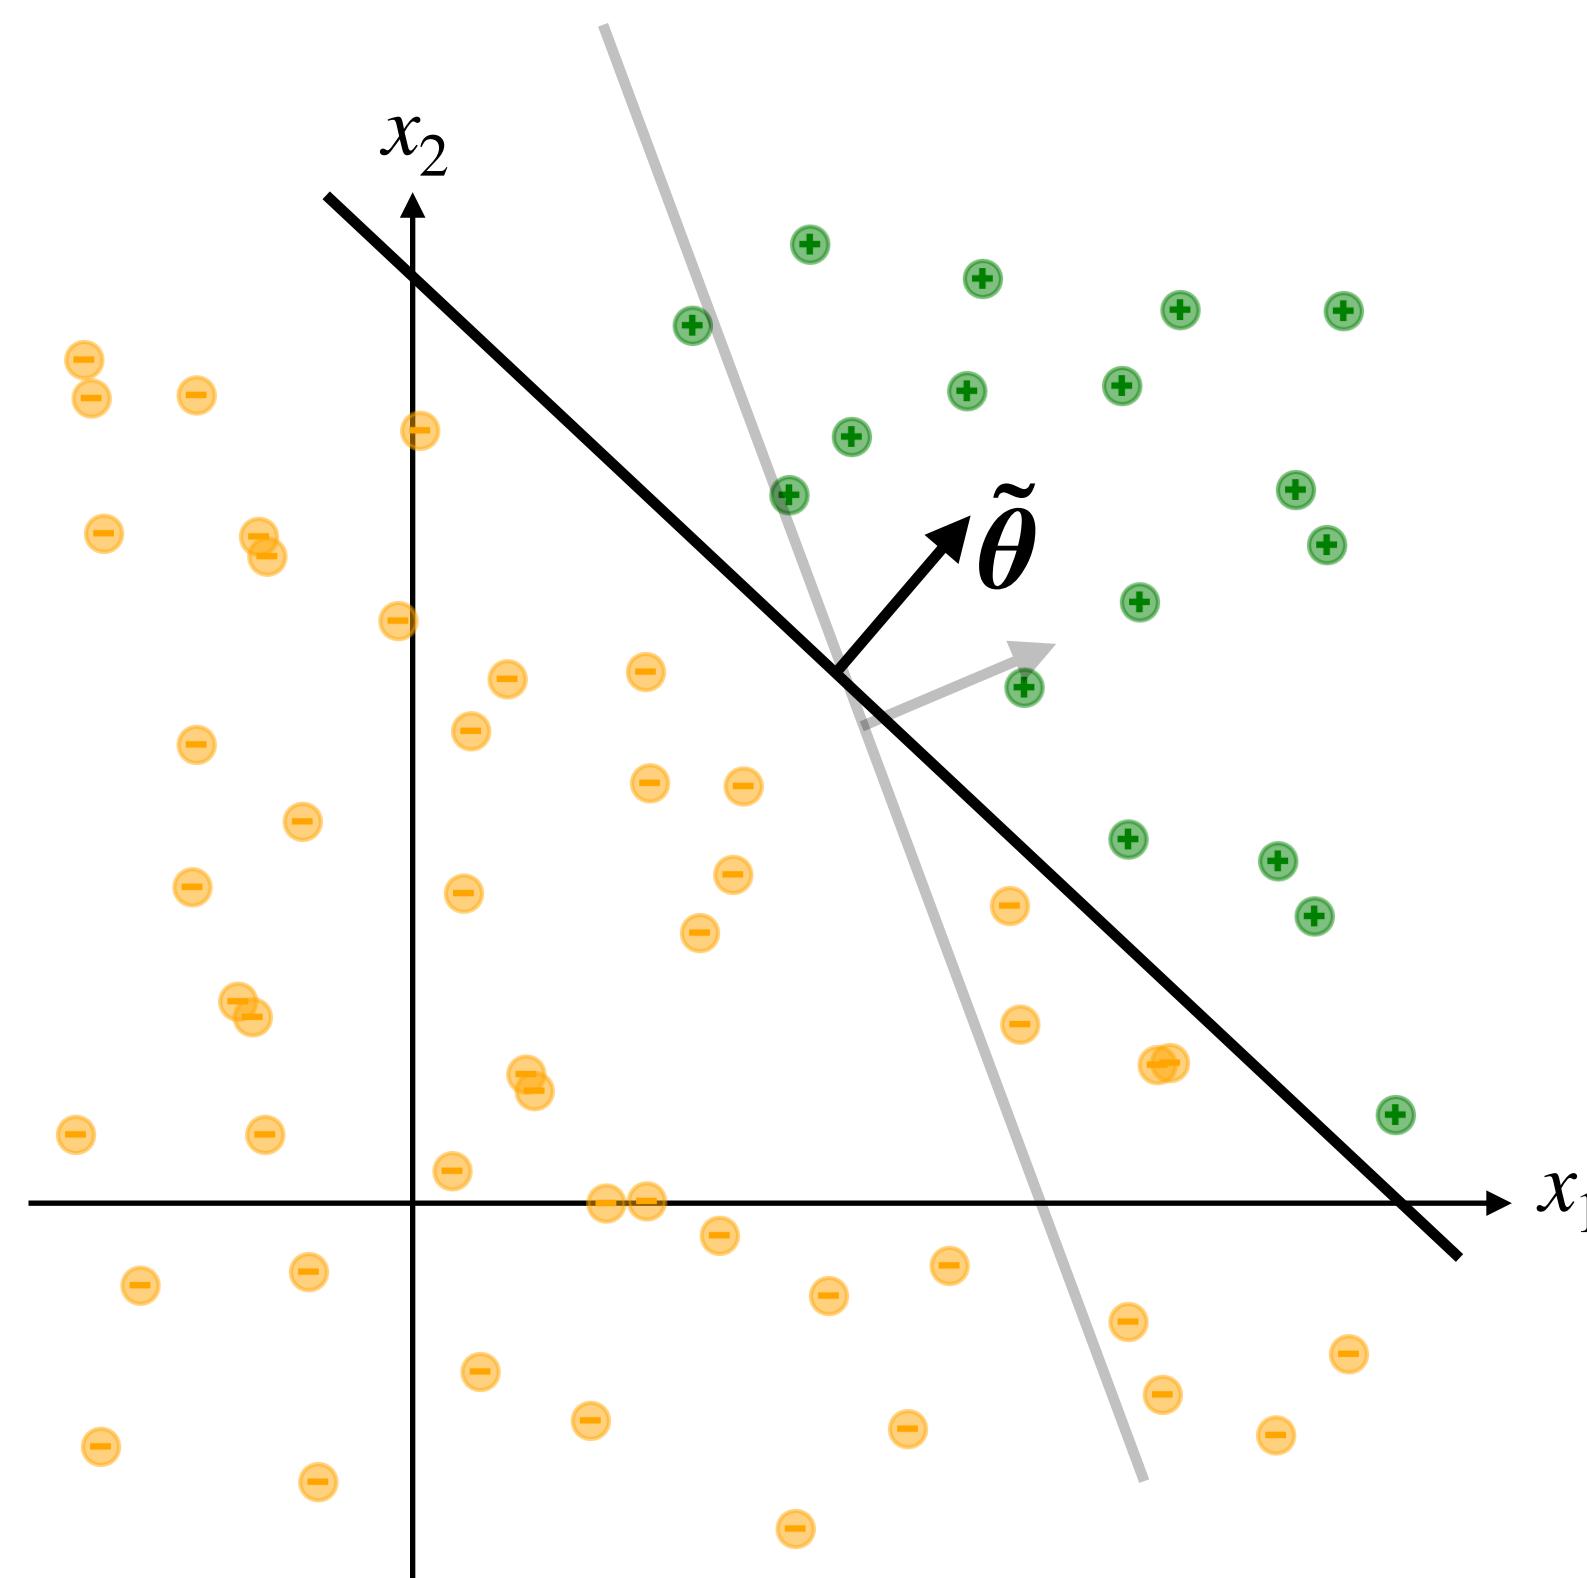


- We **can** perform gradient descent

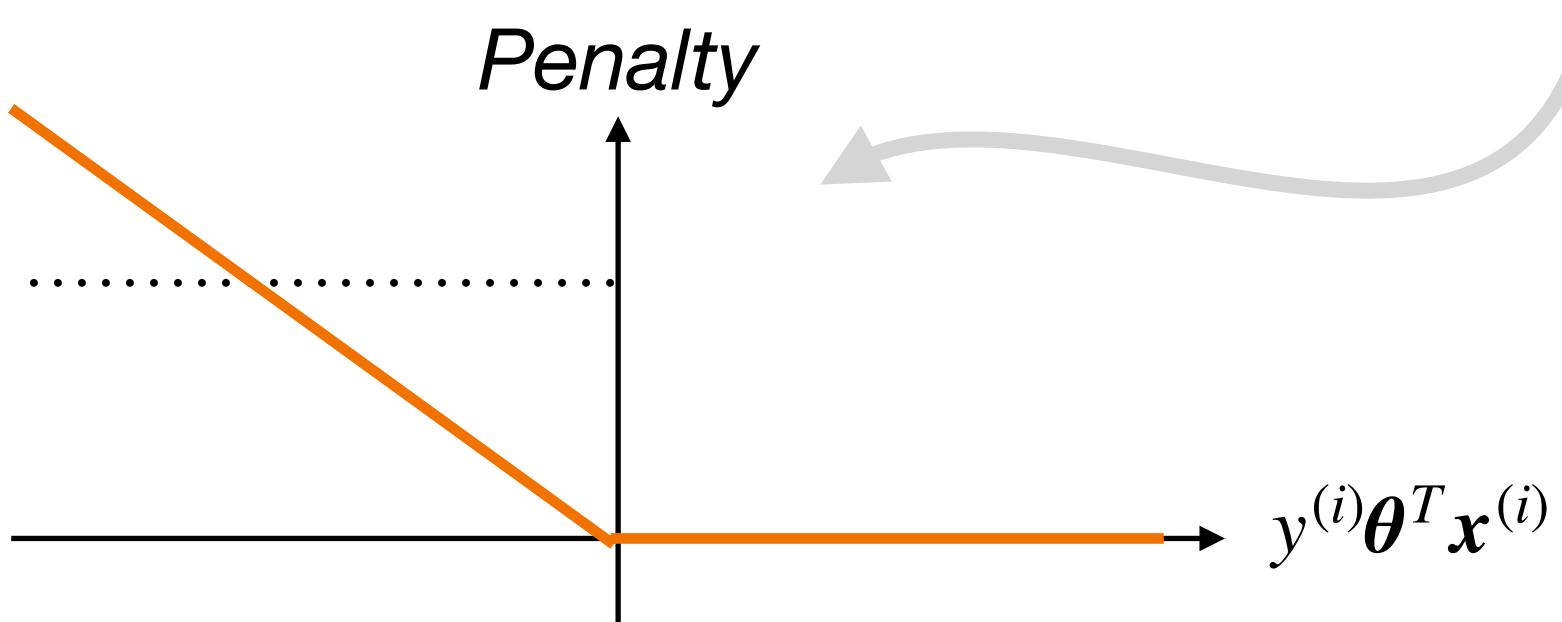
$$\theta^{t+1} = \theta^t - \lambda \frac{\partial \varepsilon(\theta)}{\partial \theta} = \theta^t + \lambda \sum_{i: \text{misclassified}} y^{(i)} x^{(i)}$$

Loss function

Goal: Find an optimal solution

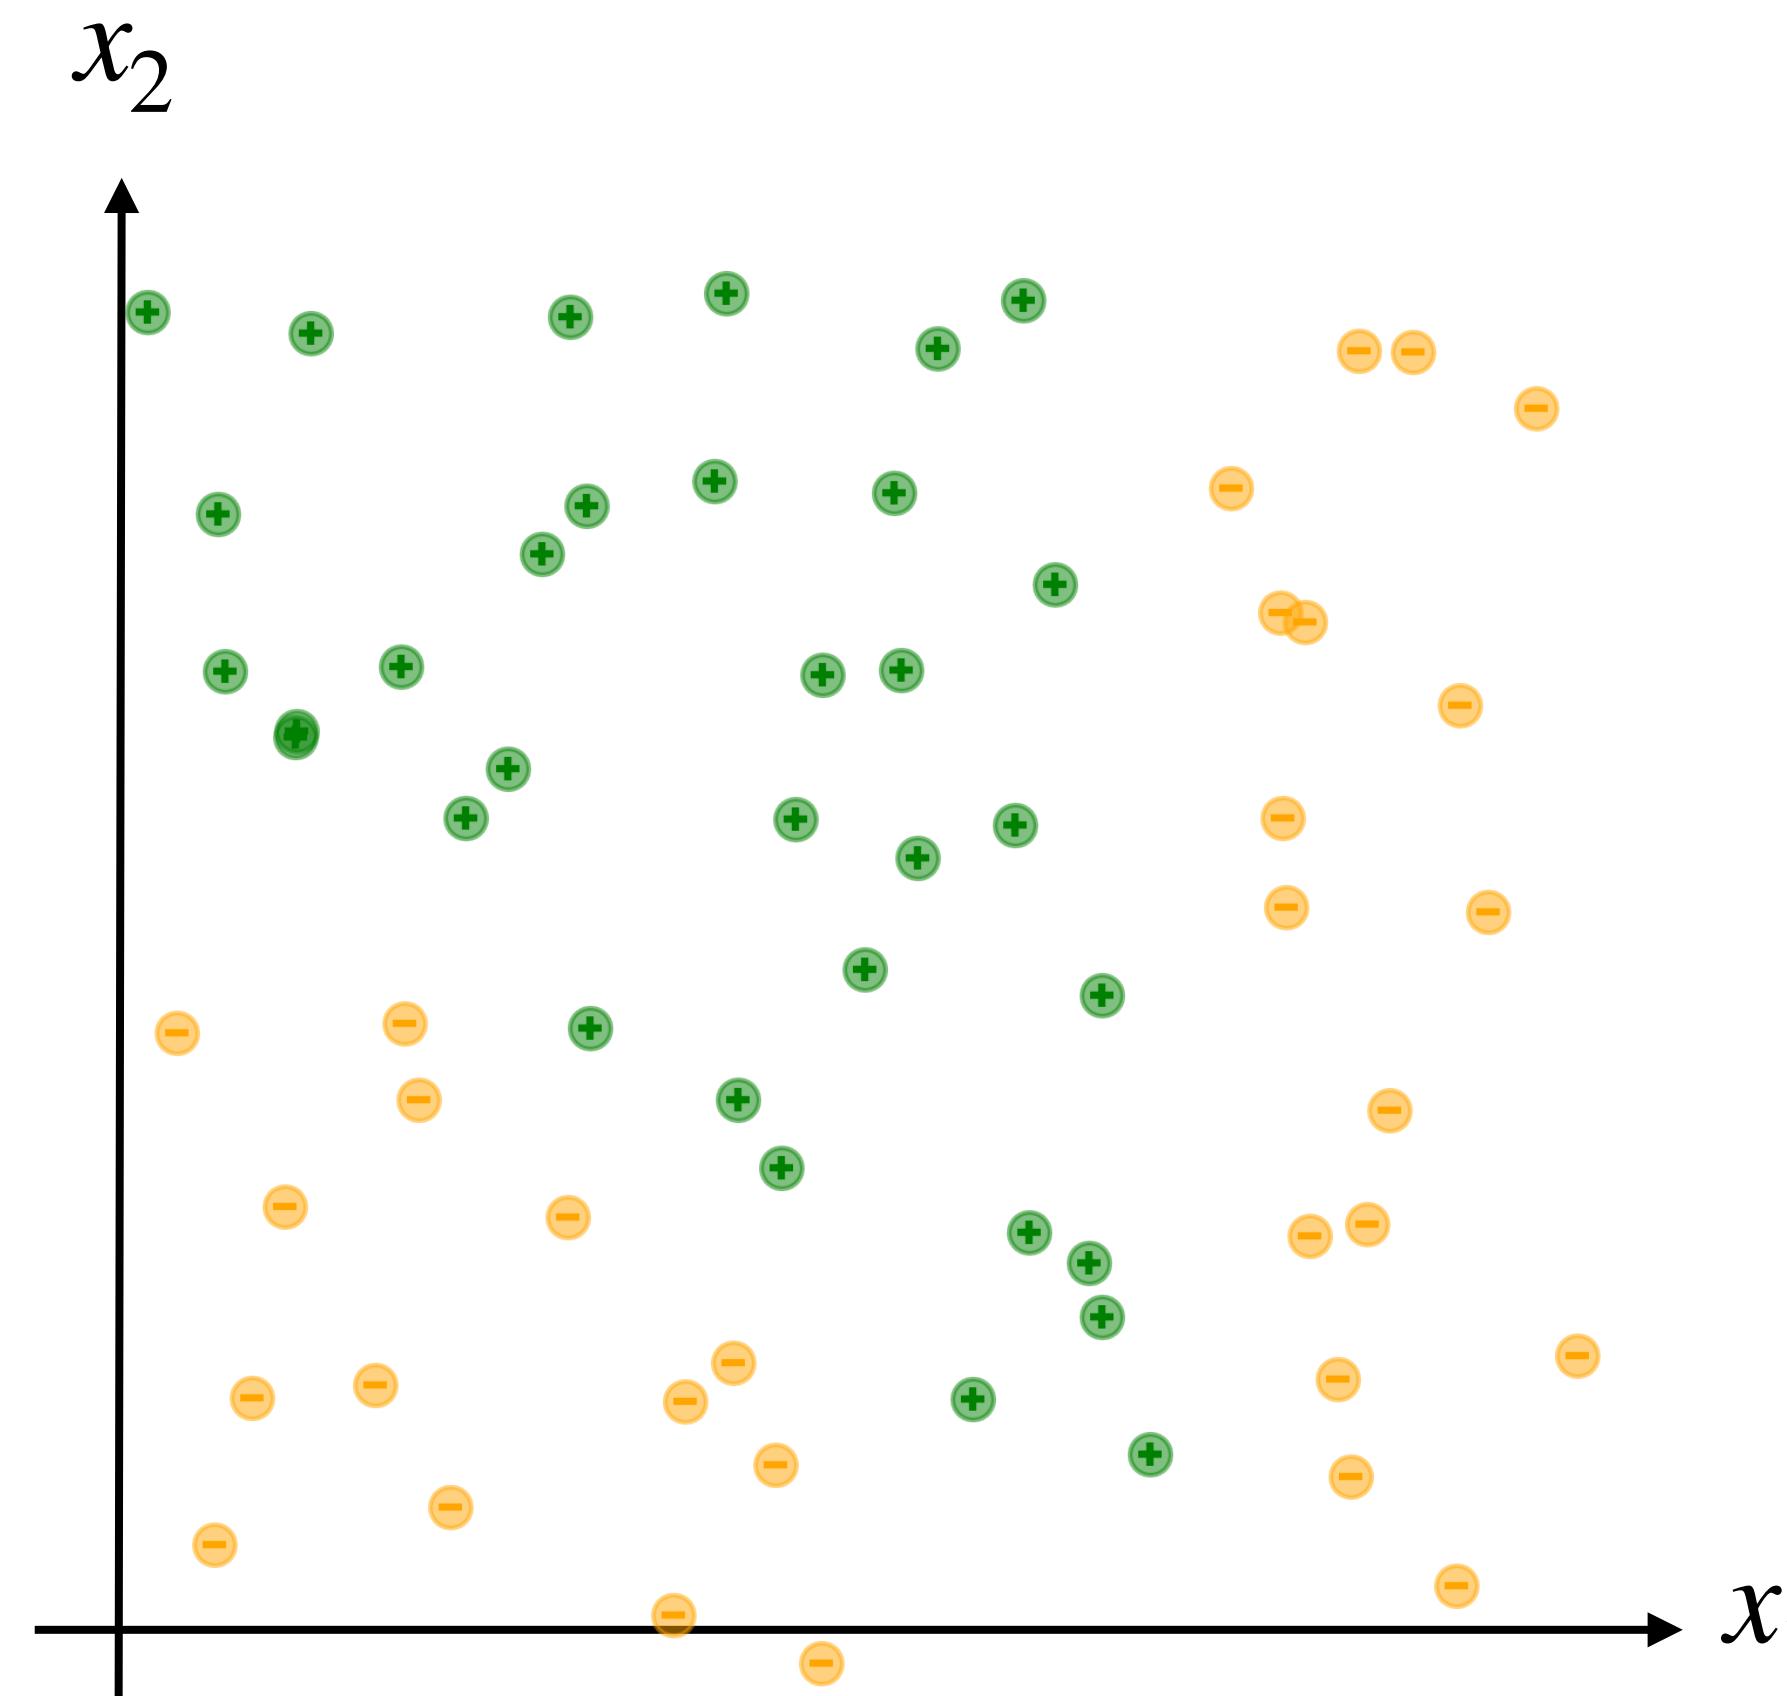


B) The hinge loss: $\varepsilon(\theta) = \sum_{i=1}^m \max(-y^{(i)}\theta^T x^{(i)}, 0)$



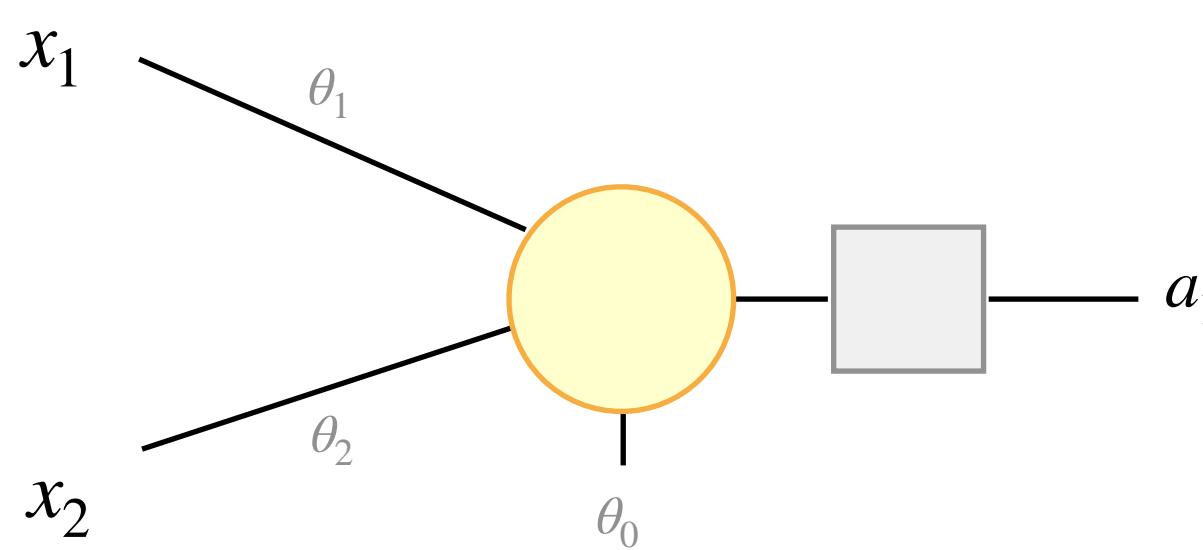
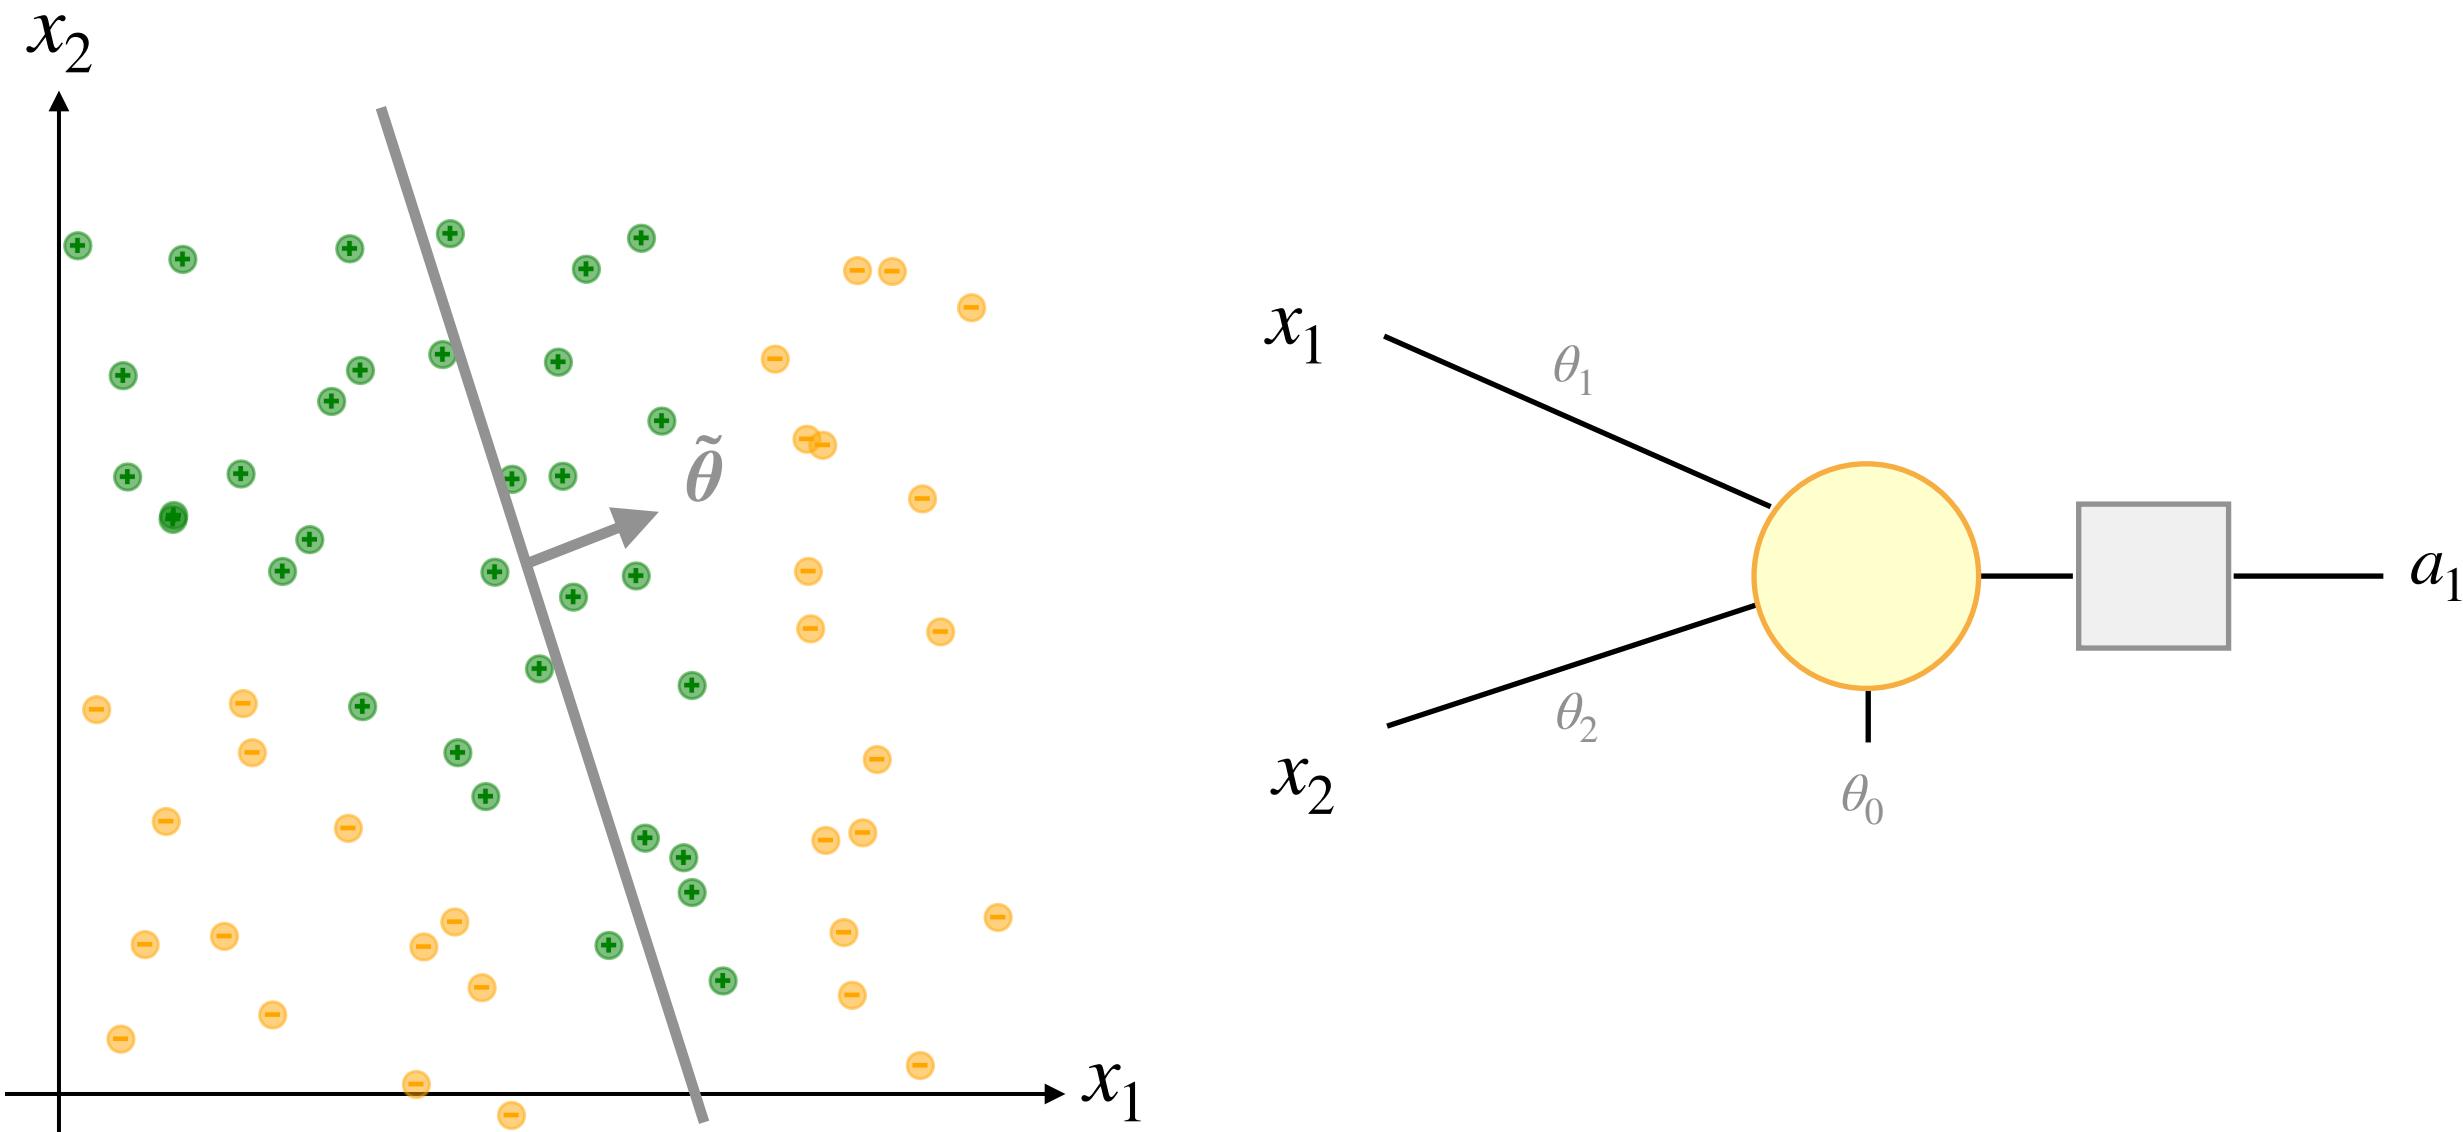
- We **can** perform gradient descent
- But: not convex, ambiguity of “optimal” solution
=> Logistic Regression :)

*What if the problem is **not** linearly separable?*



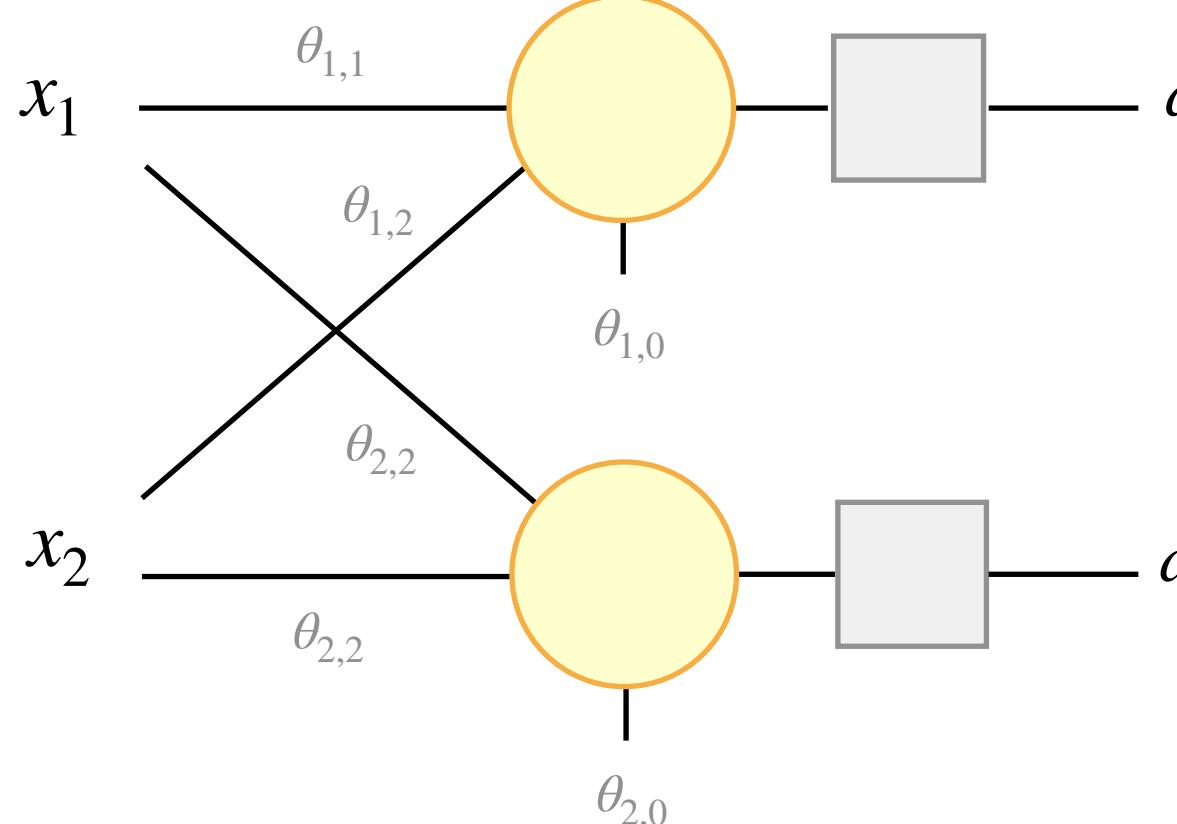
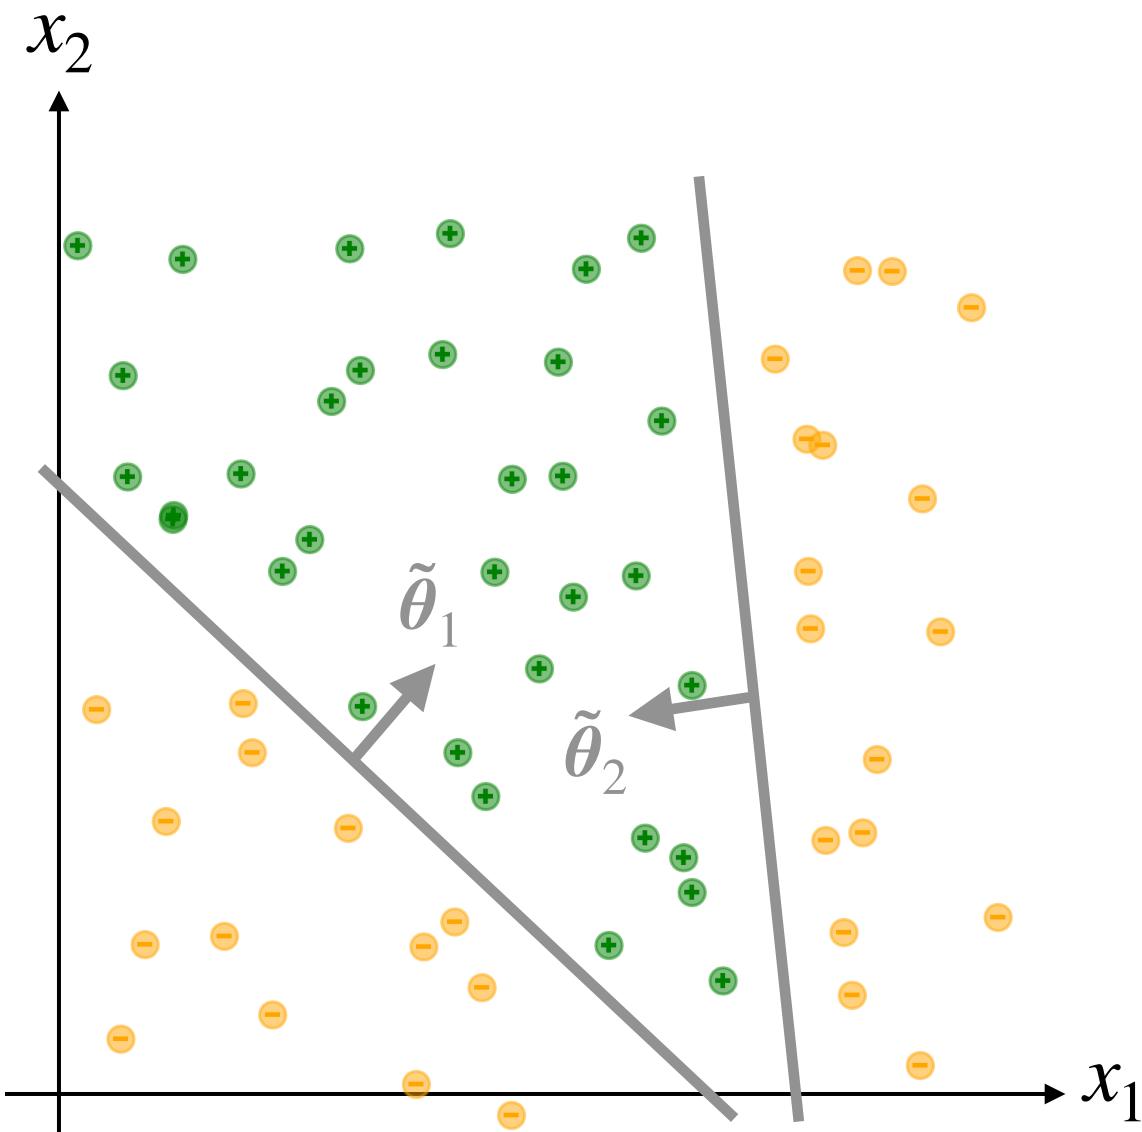
Multi-layer perceptrons (MLPs) / Feed-forward neural networks

Multi-Layer Perceptrons

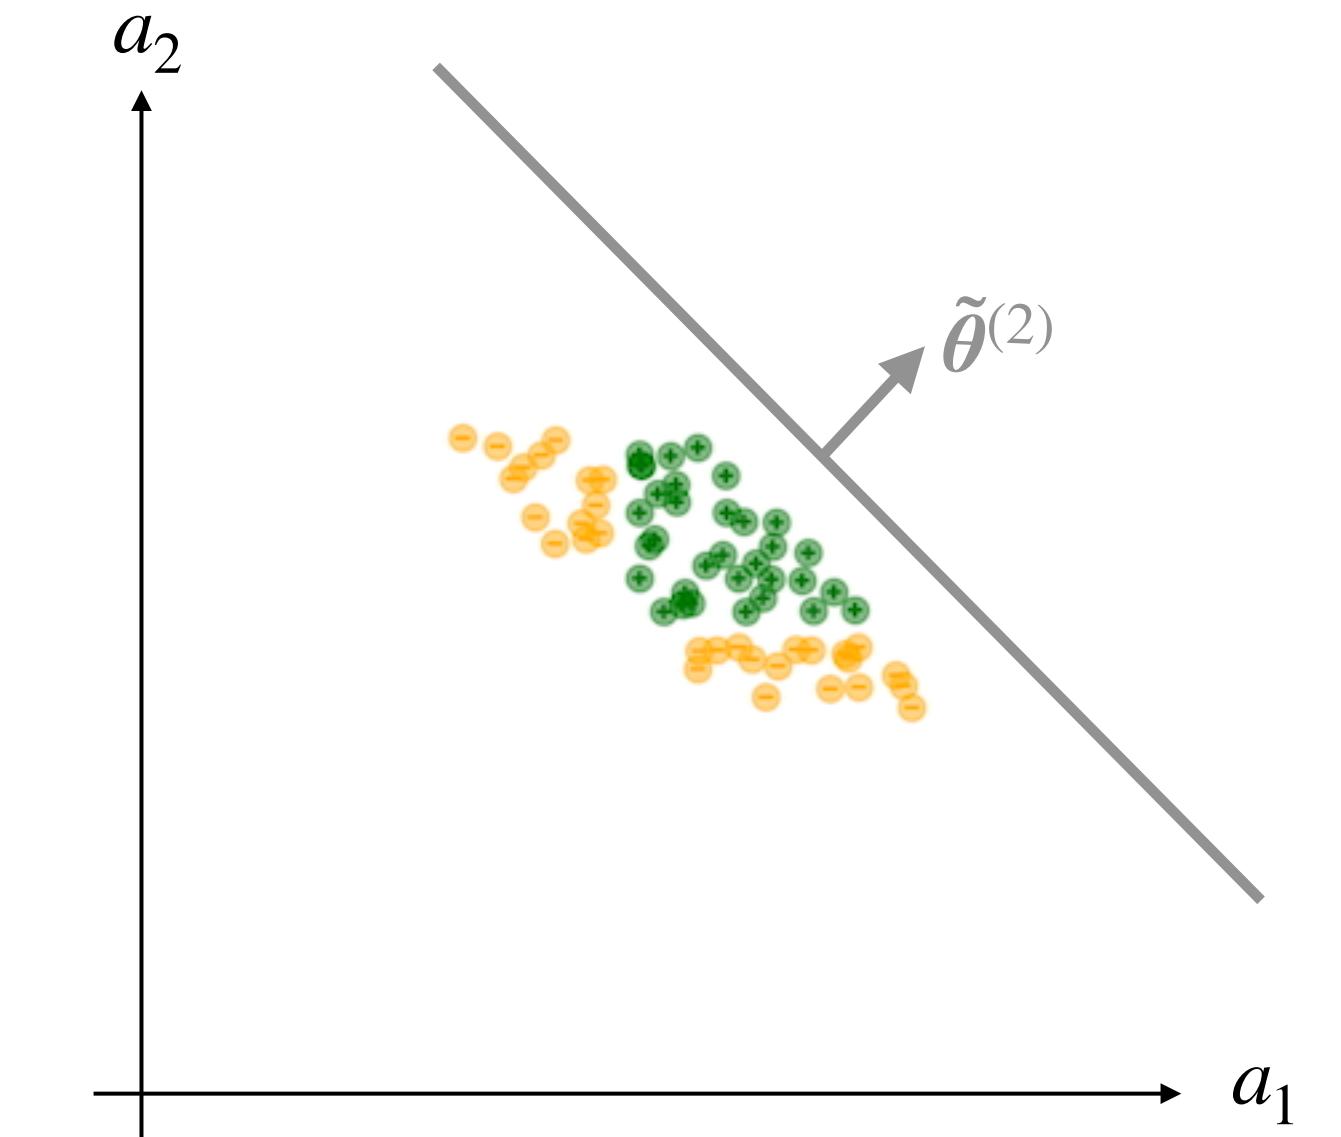
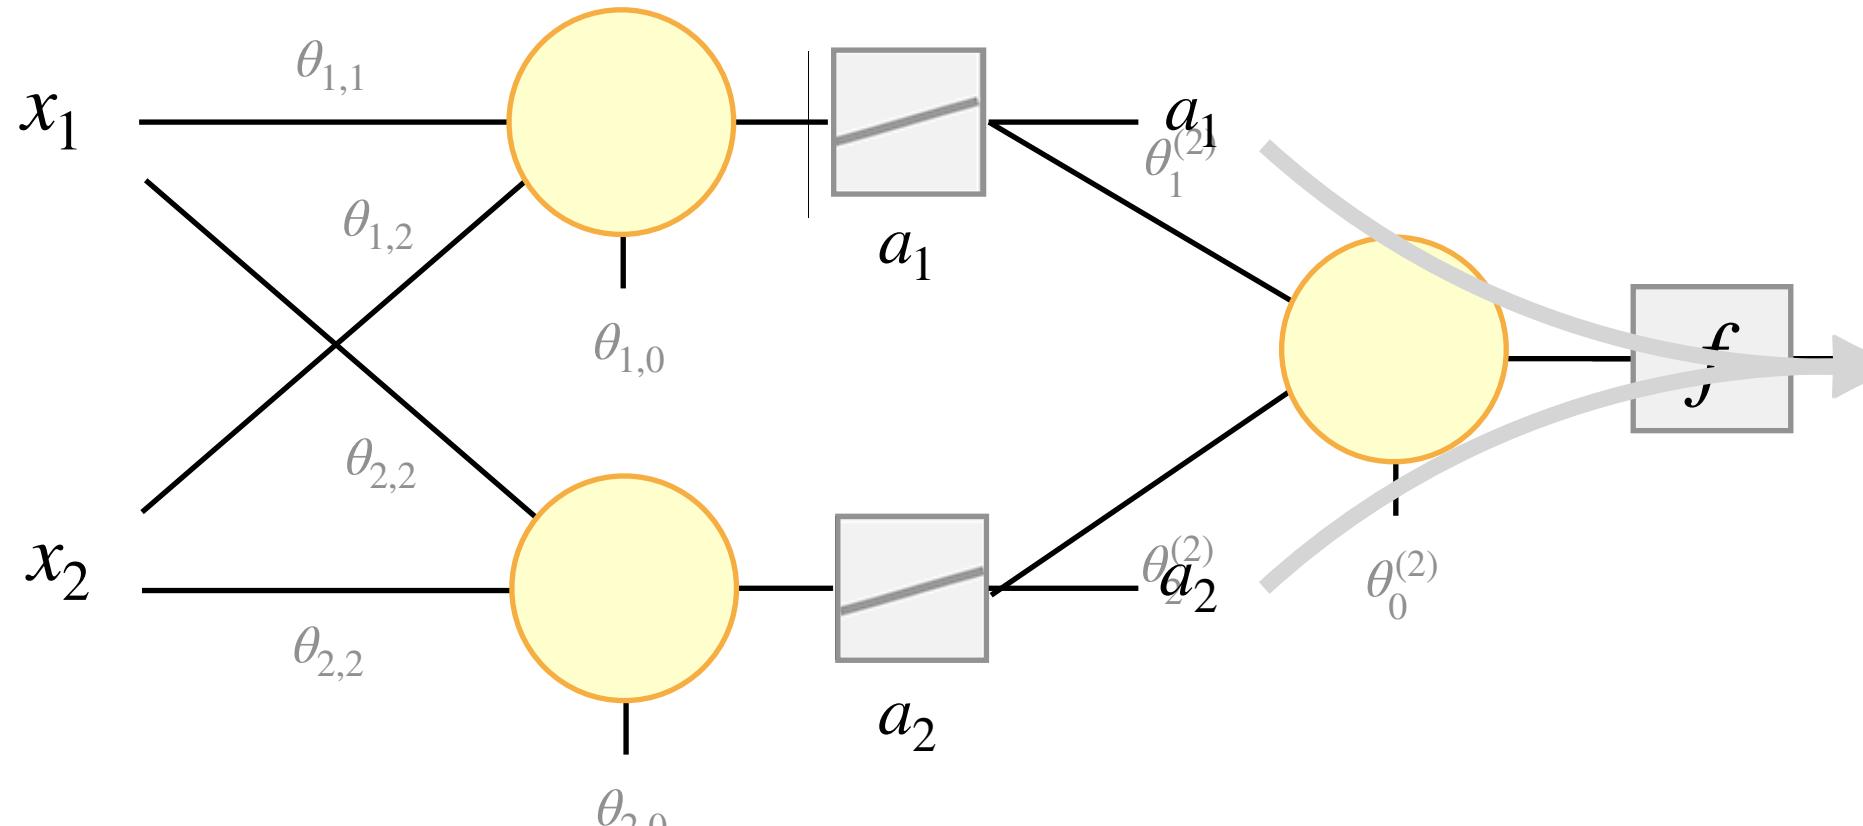
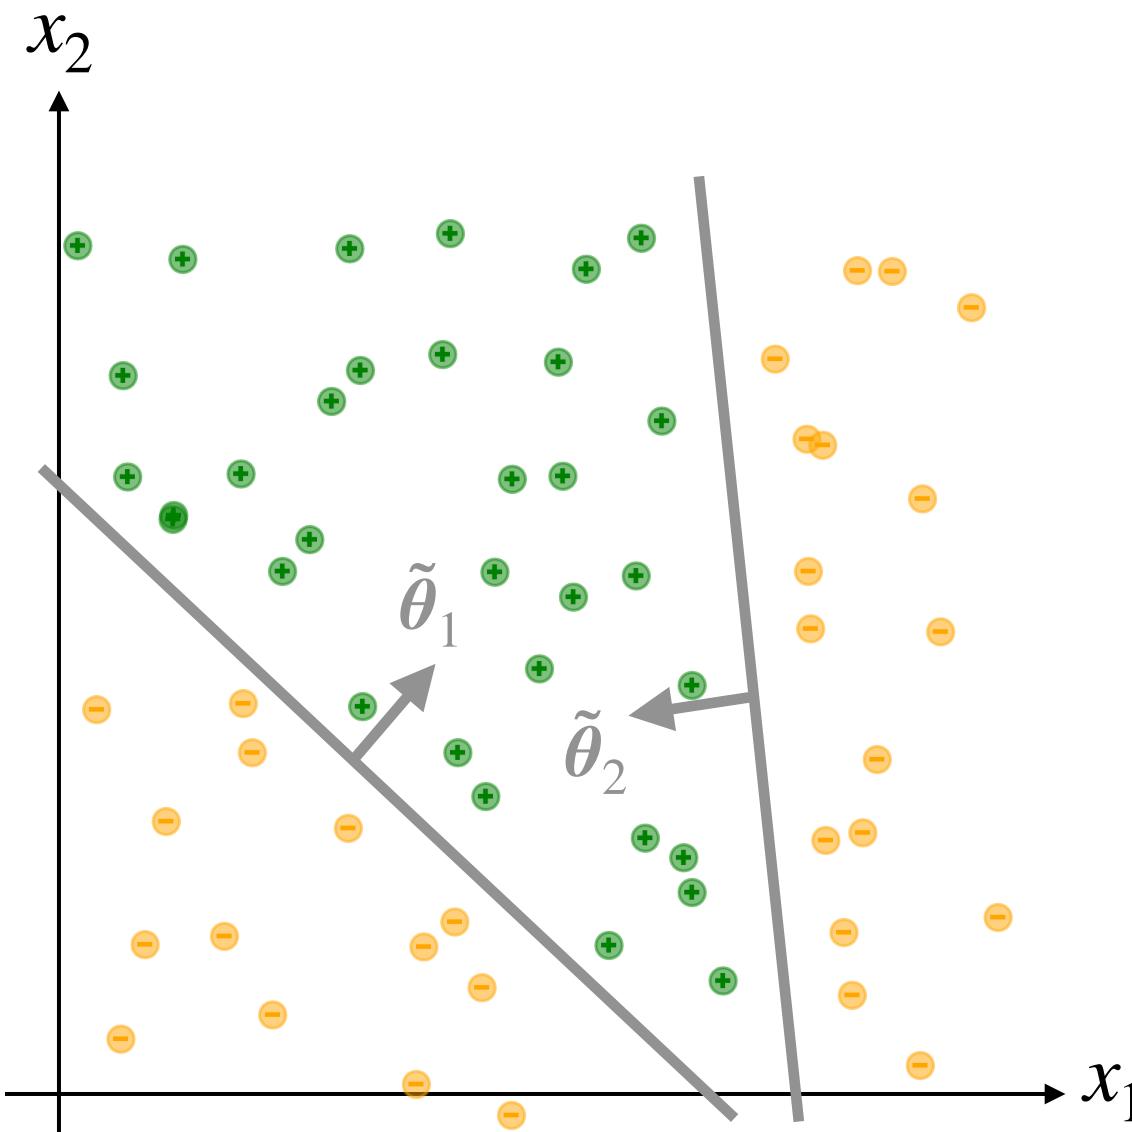


Not linearly separable

Multi-Layer Perceptrons



Multi-Layer Perceptrons

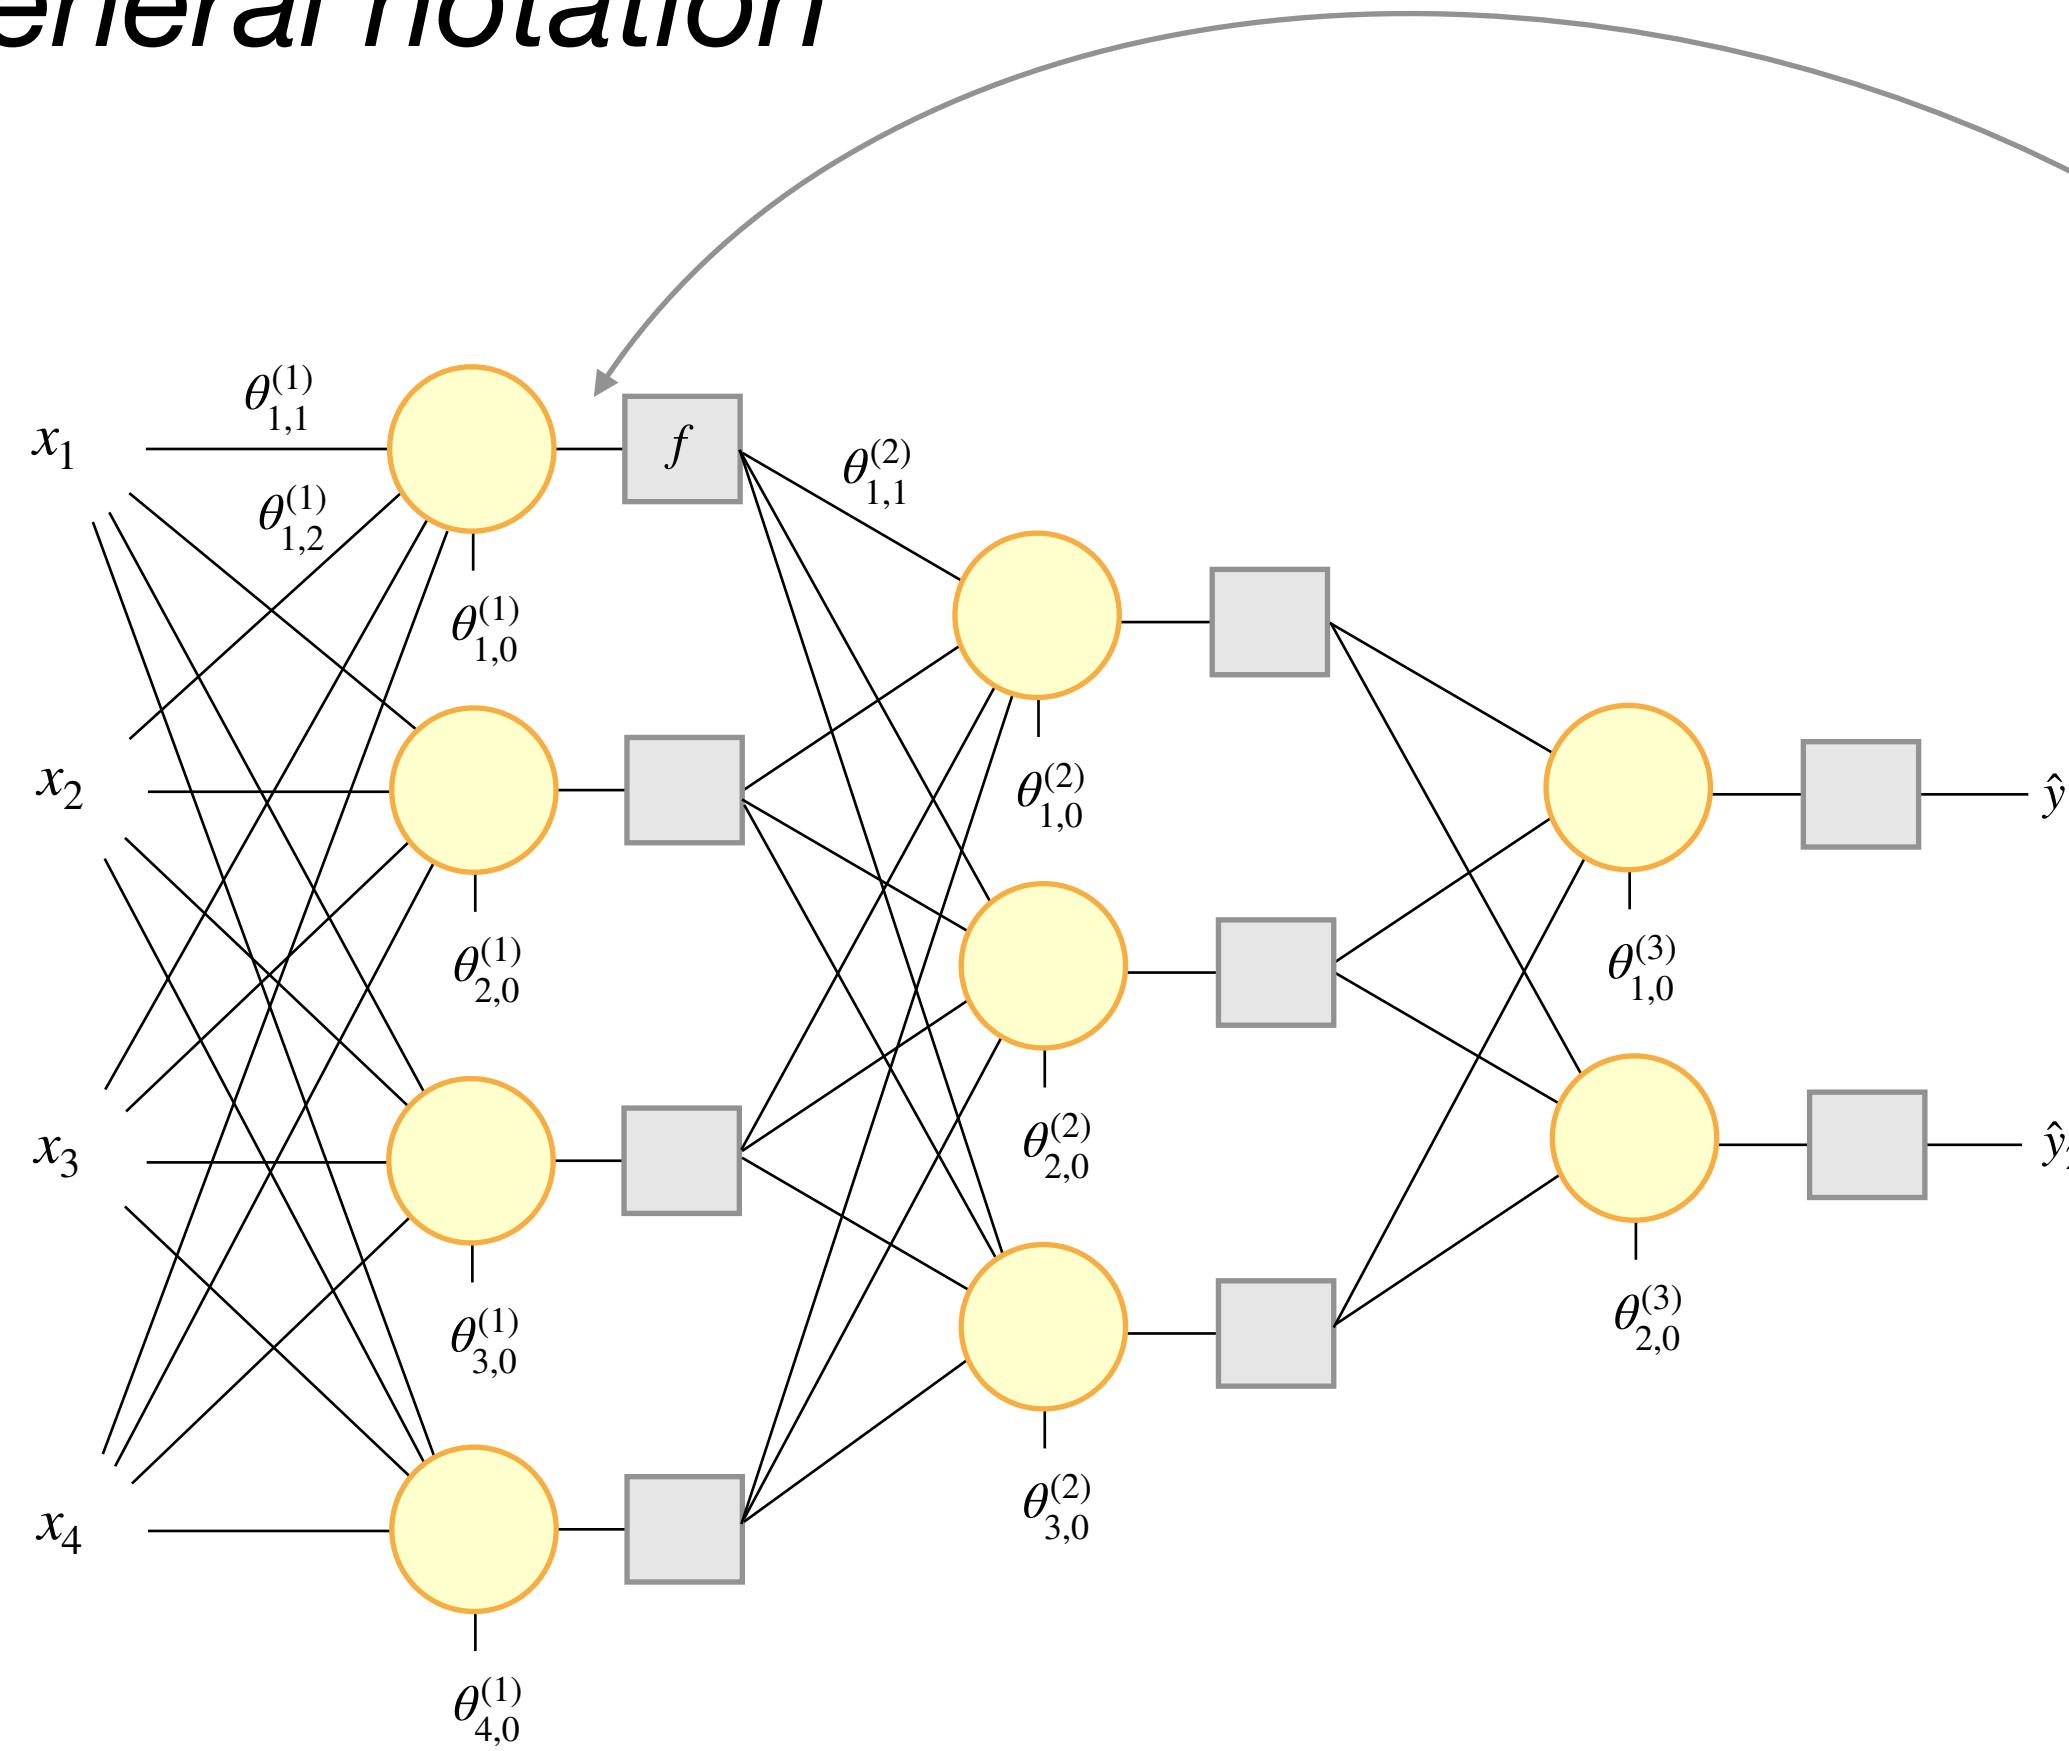


Why non-linear activation function?

NN with linear activation func $\hat{=}$ Single layer NN

Multi-Layer Perceptrons

General notation



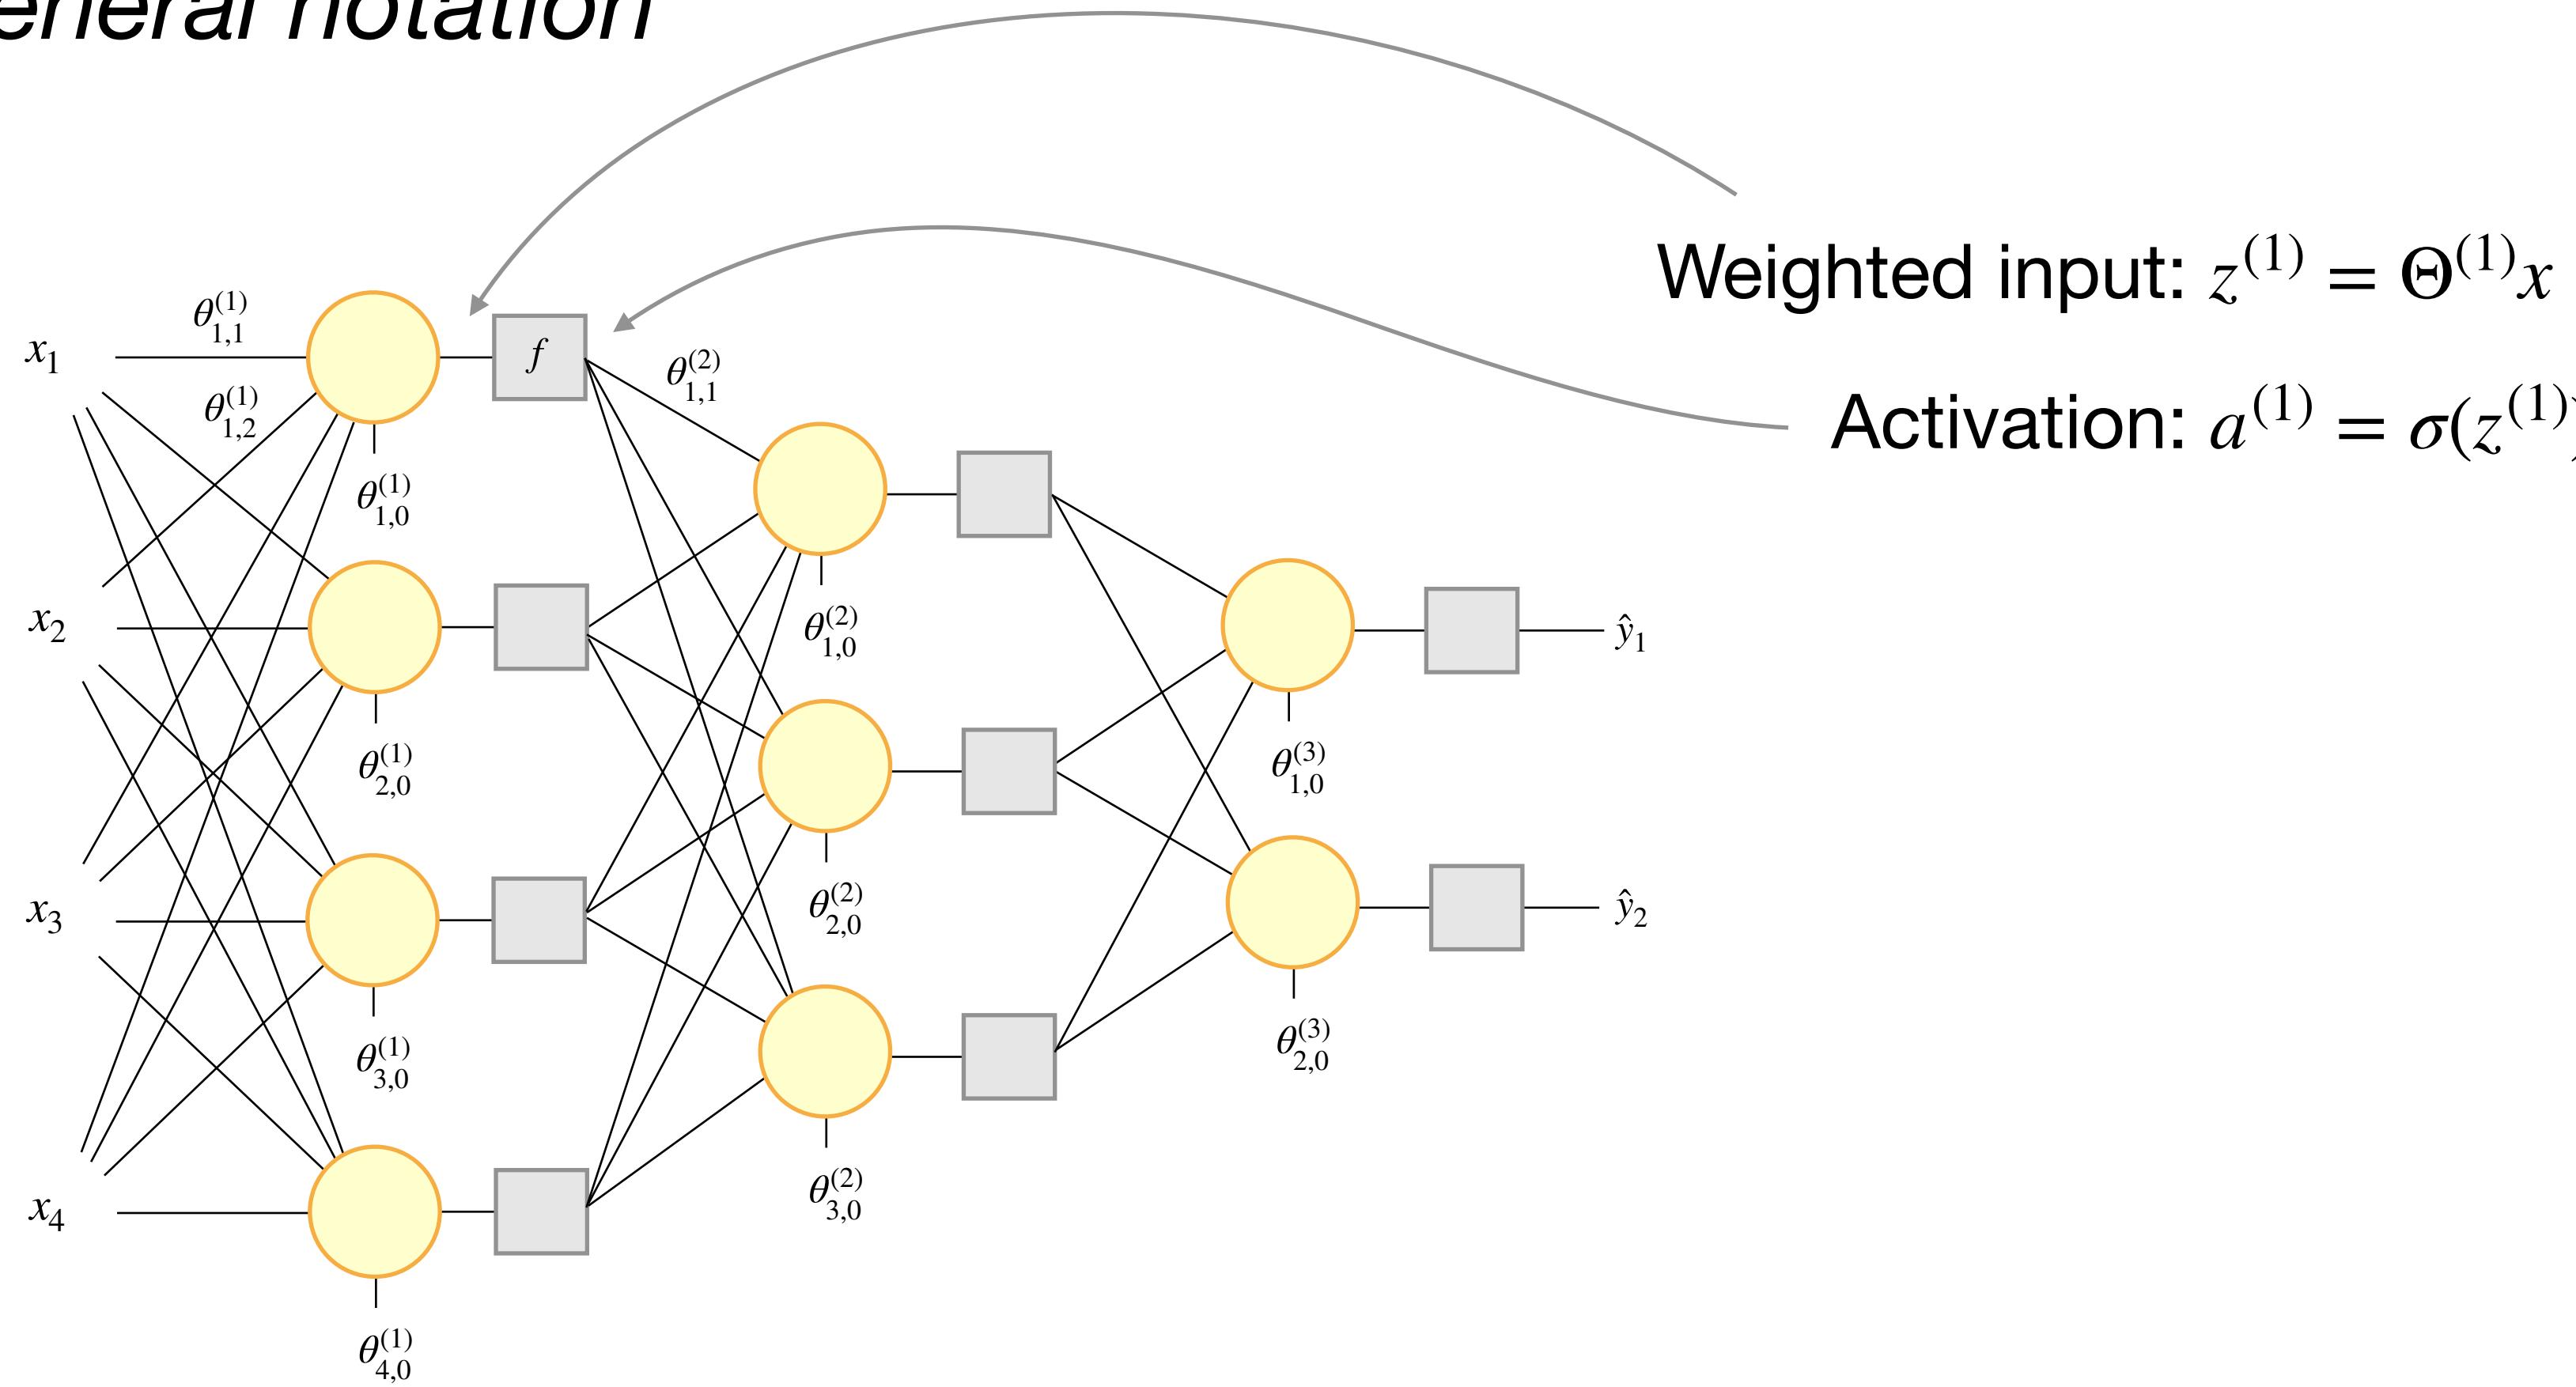
Weighted input: $z^{(1)} = \Theta^{(1)}x =$

Weights to first neuron

$$\begin{bmatrix} \theta_{1,0}^{(1)} & \theta_{1,1}^{(1)} & \dots & \theta_{1,n}^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ \theta_{k,0}^{(1)} & \theta_{k,1}^{(1)} & \dots & \theta_{k,n}^{(1)} \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$$

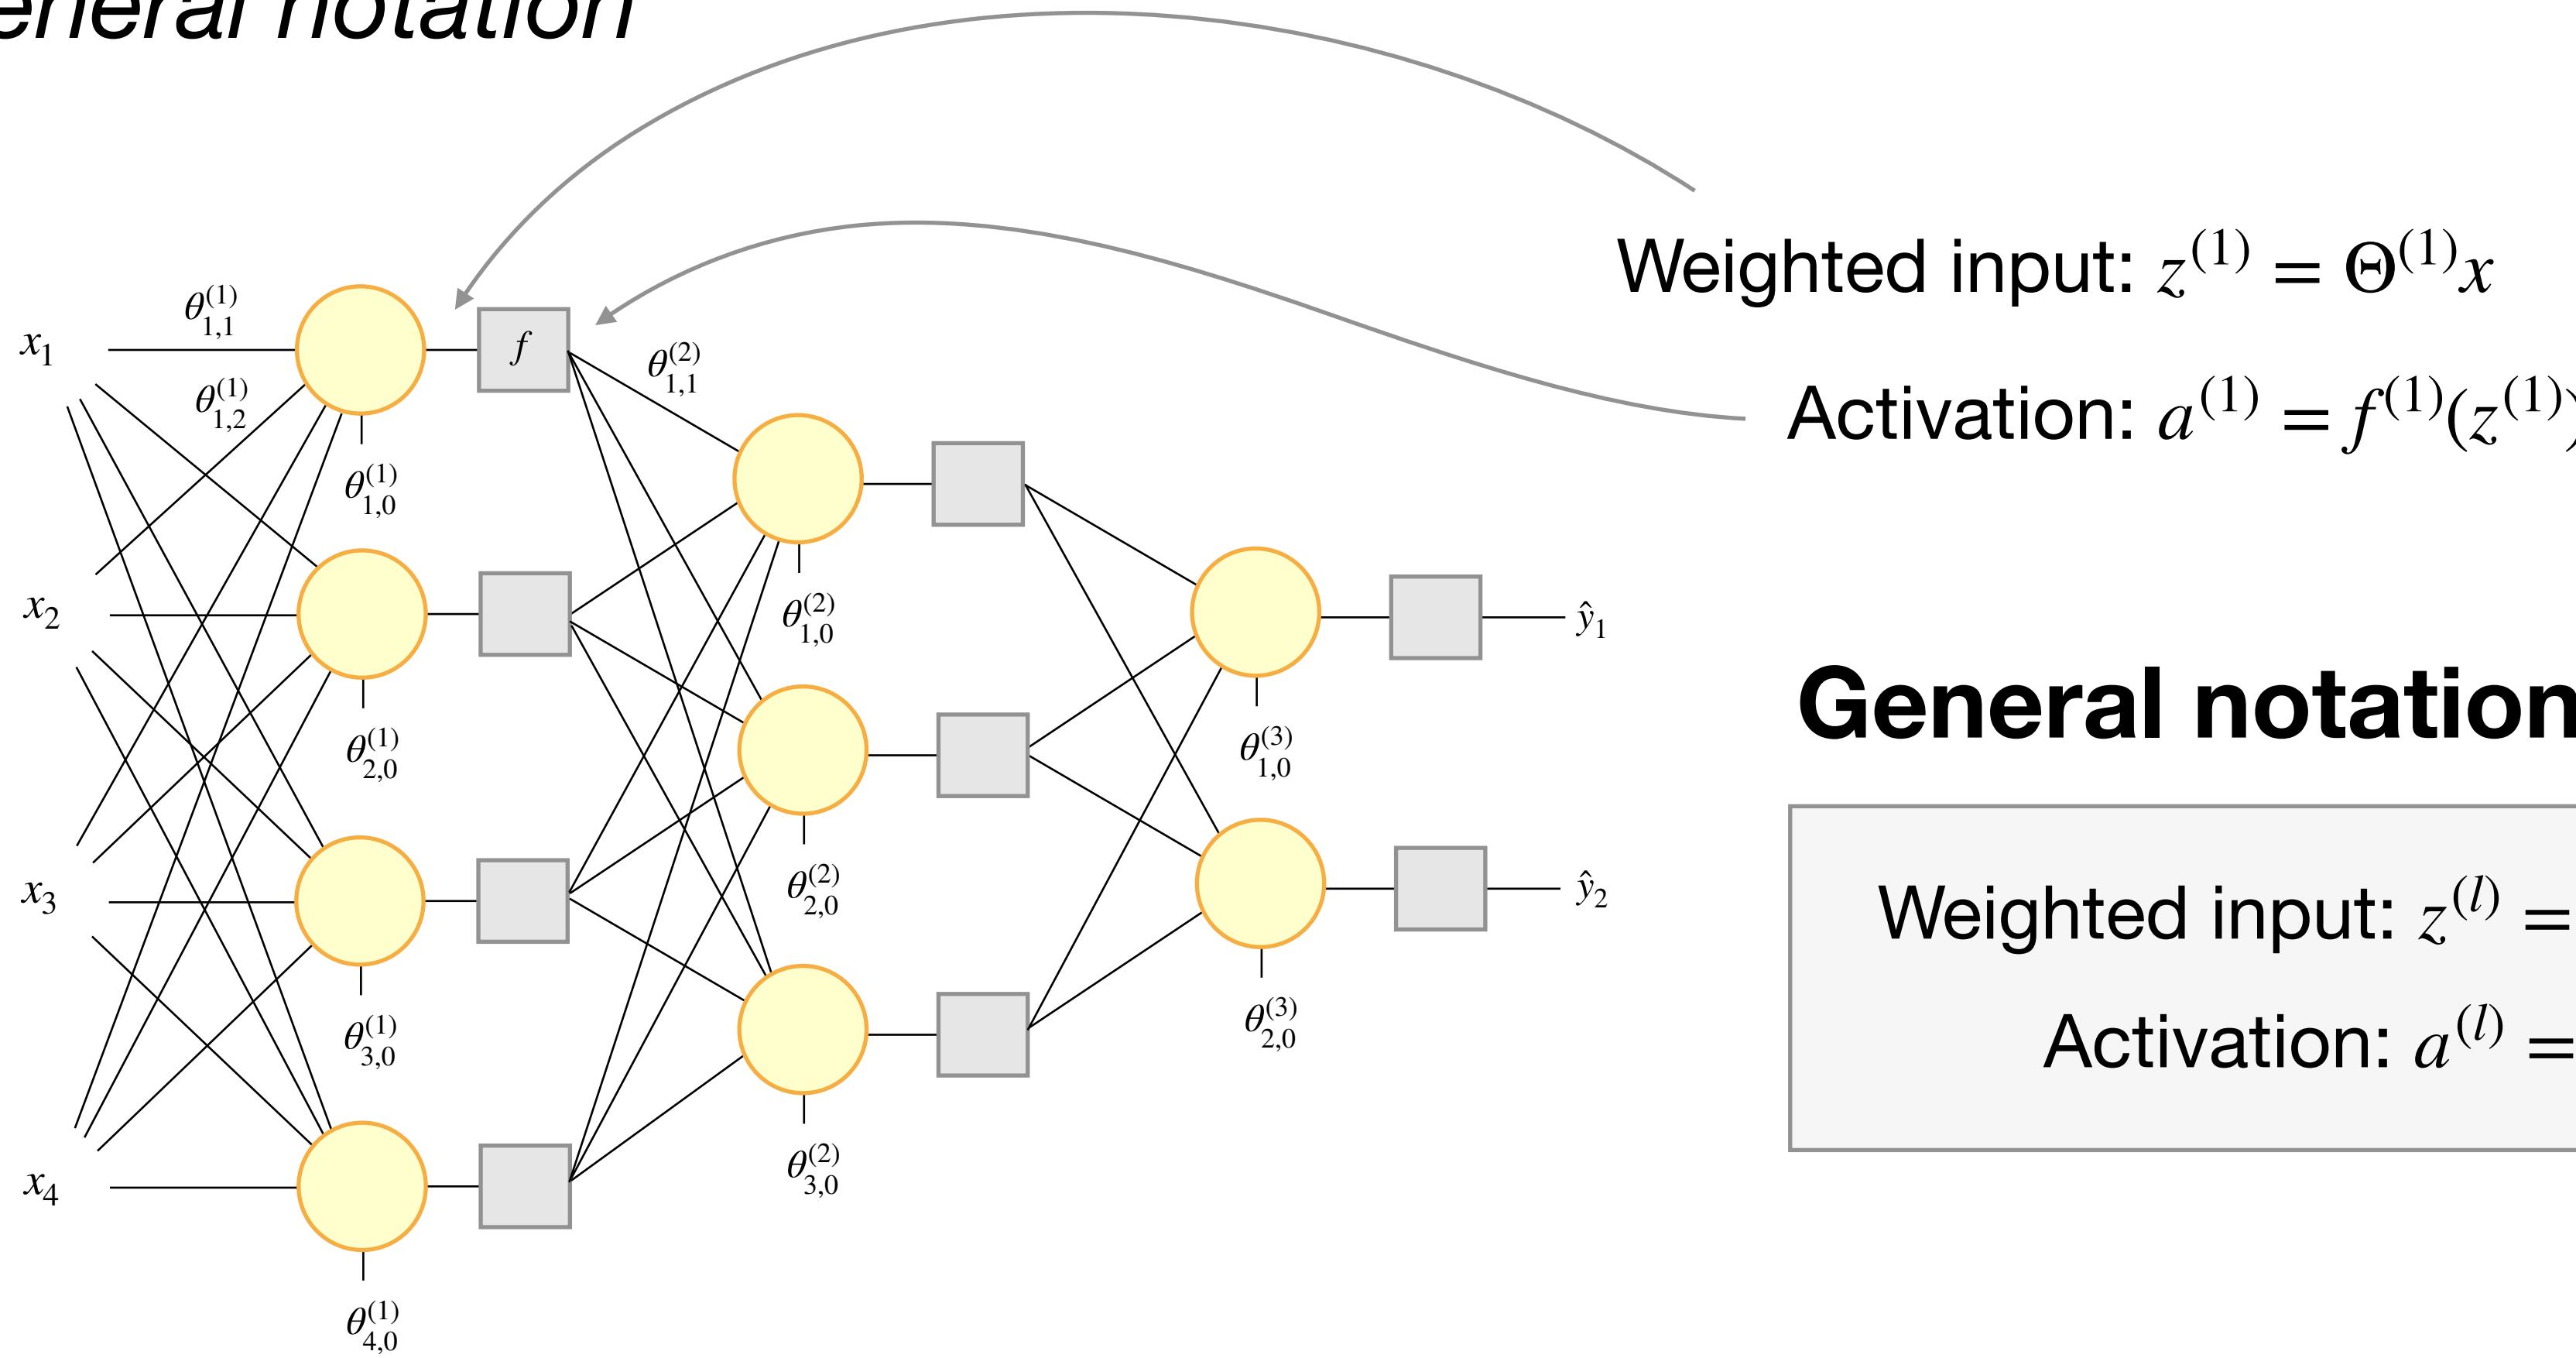
Multi-Layer Perceptrons

General notation



Multi-Layer Perceptrons

General notation



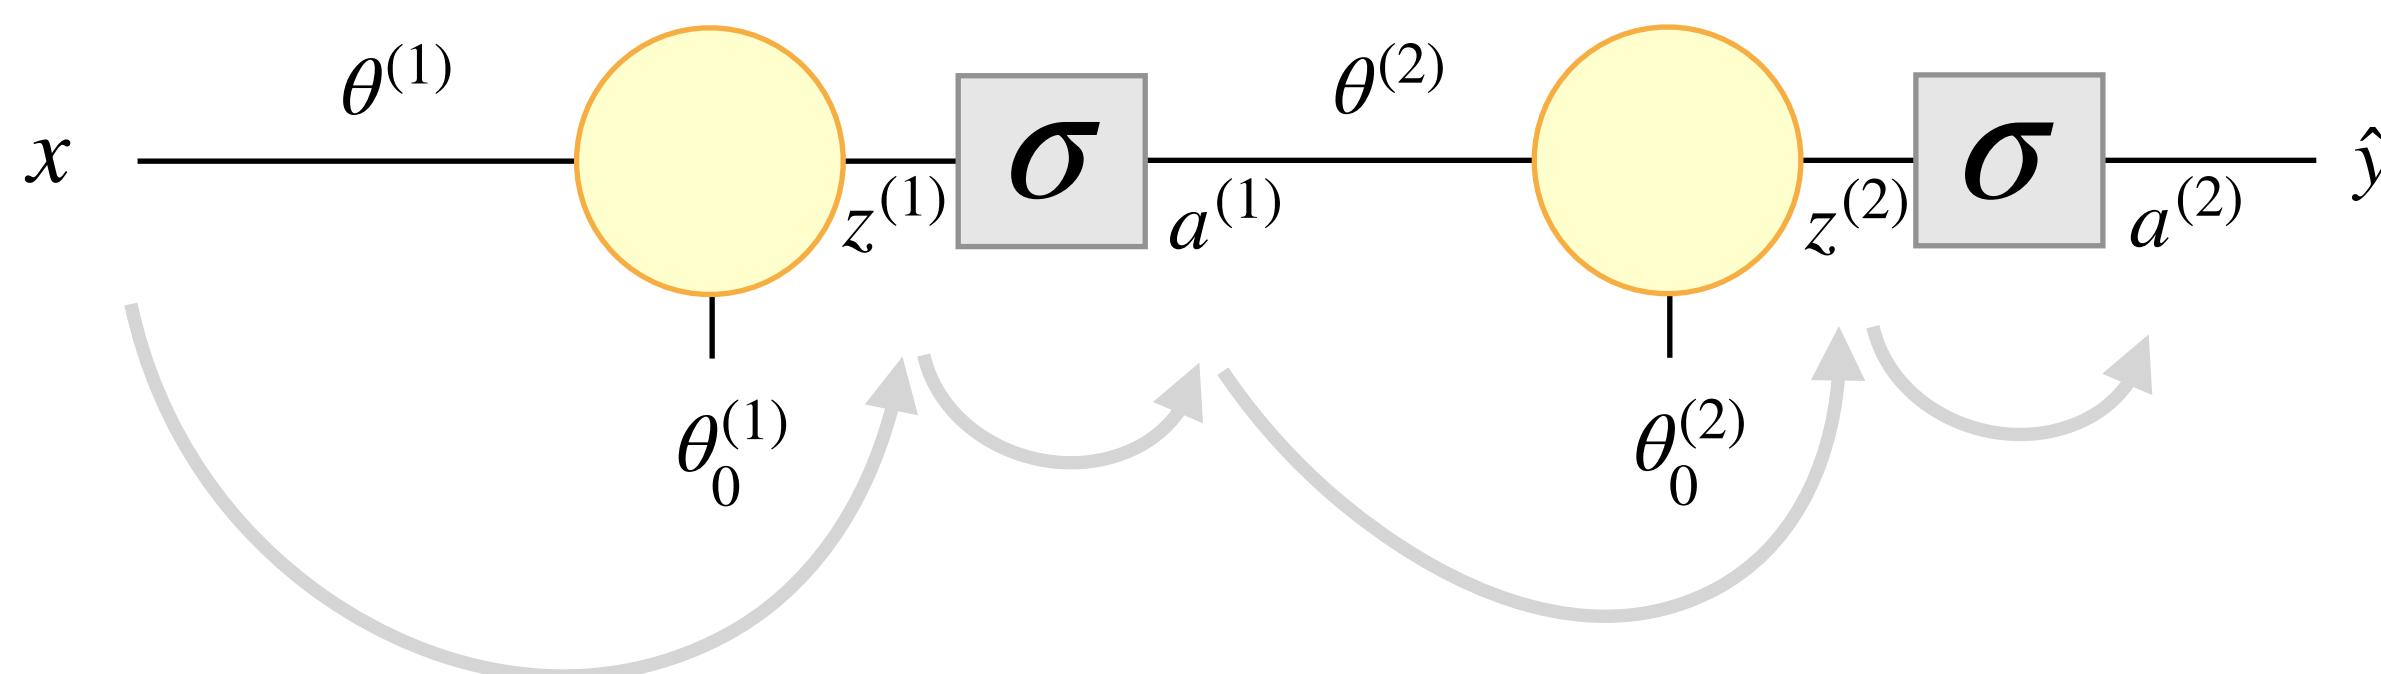
General notation for layer l :

Weighted input: $z^{(l)} = \Theta^{(l)}a^{(l-1)}$

Activation: $a^{(l)} = f^{(l)}(z^{(l)})$

The Idea of Backpropagation

for 1D Neural Network



1) Feedforward pass

$$z^{(1)} = \theta^{(1)}x + \theta_0^{(1)}$$

$$a^{(1)} = \sigma(z^{(1)})$$

$$z^{(2)} = \theta^{(2)}a^{(1)} + \theta_0^{(2)}$$

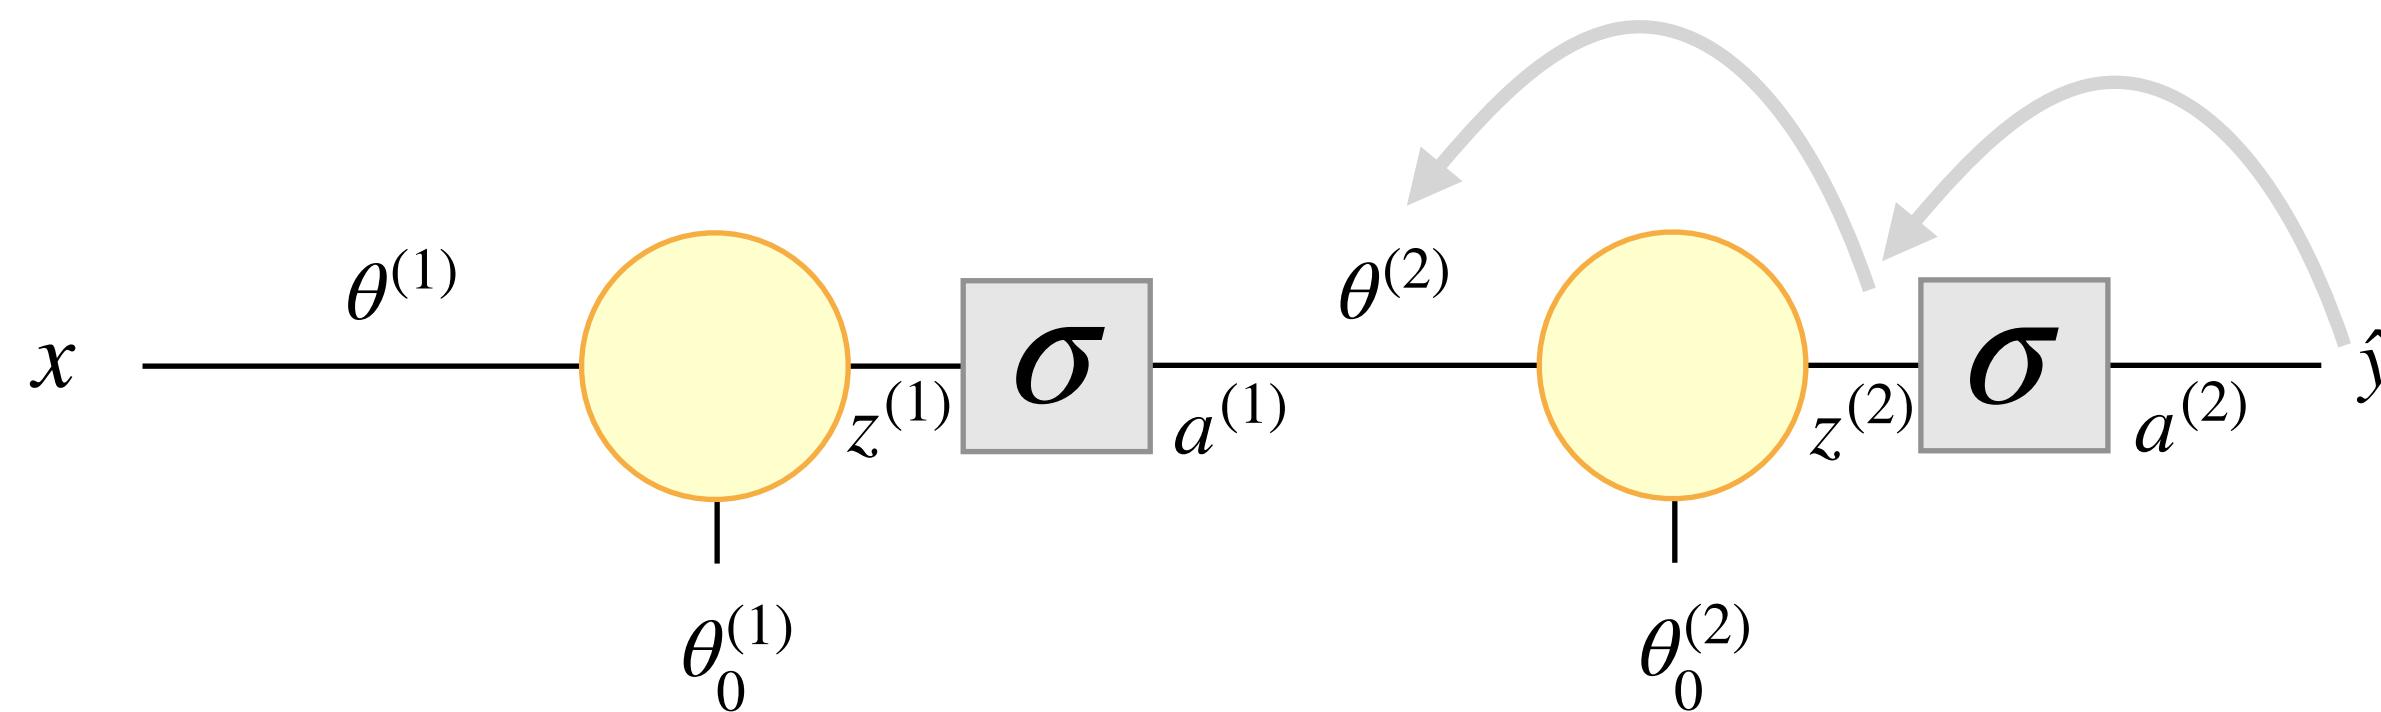
$$\hat{y} = a^{(2)} = \sigma(z^{(2)})$$

2) Compute loss

$$\varepsilon(\hat{y}) = \frac{1}{2}(y - \hat{y})^2$$

The Idea of Backpropagation

for 1D Neural Network



3) Backpropagate loss gradient

$$\frac{\partial L}{\partial \theta^{(2)}} = ? \quad L(\theta^{(2)}) = \epsilon \circ \sigma \circ l(\theta^{(2)})$$

$$\frac{\partial L}{\partial \theta^{(2)}} = \frac{\partial \epsilon}{\partial a^{(2)}} \cdot \frac{\partial \sigma}{\partial z^{(2)}} \cdot \frac{\partial l}{\partial \theta^{(2)}}$$

1) Feedforward pass

$$z^{(1)} = \theta^{(1)}x + \theta_0^{(1)}$$

$$a^{(1)} = \sigma(z^{(1)})$$

$$z^{(2)} = \theta^{(2)}a^{(1)} + \theta_0^{(2)}$$

$$\hat{y} = a^{(2)} = \sigma(z^{(2)})$$

2) Compute loss

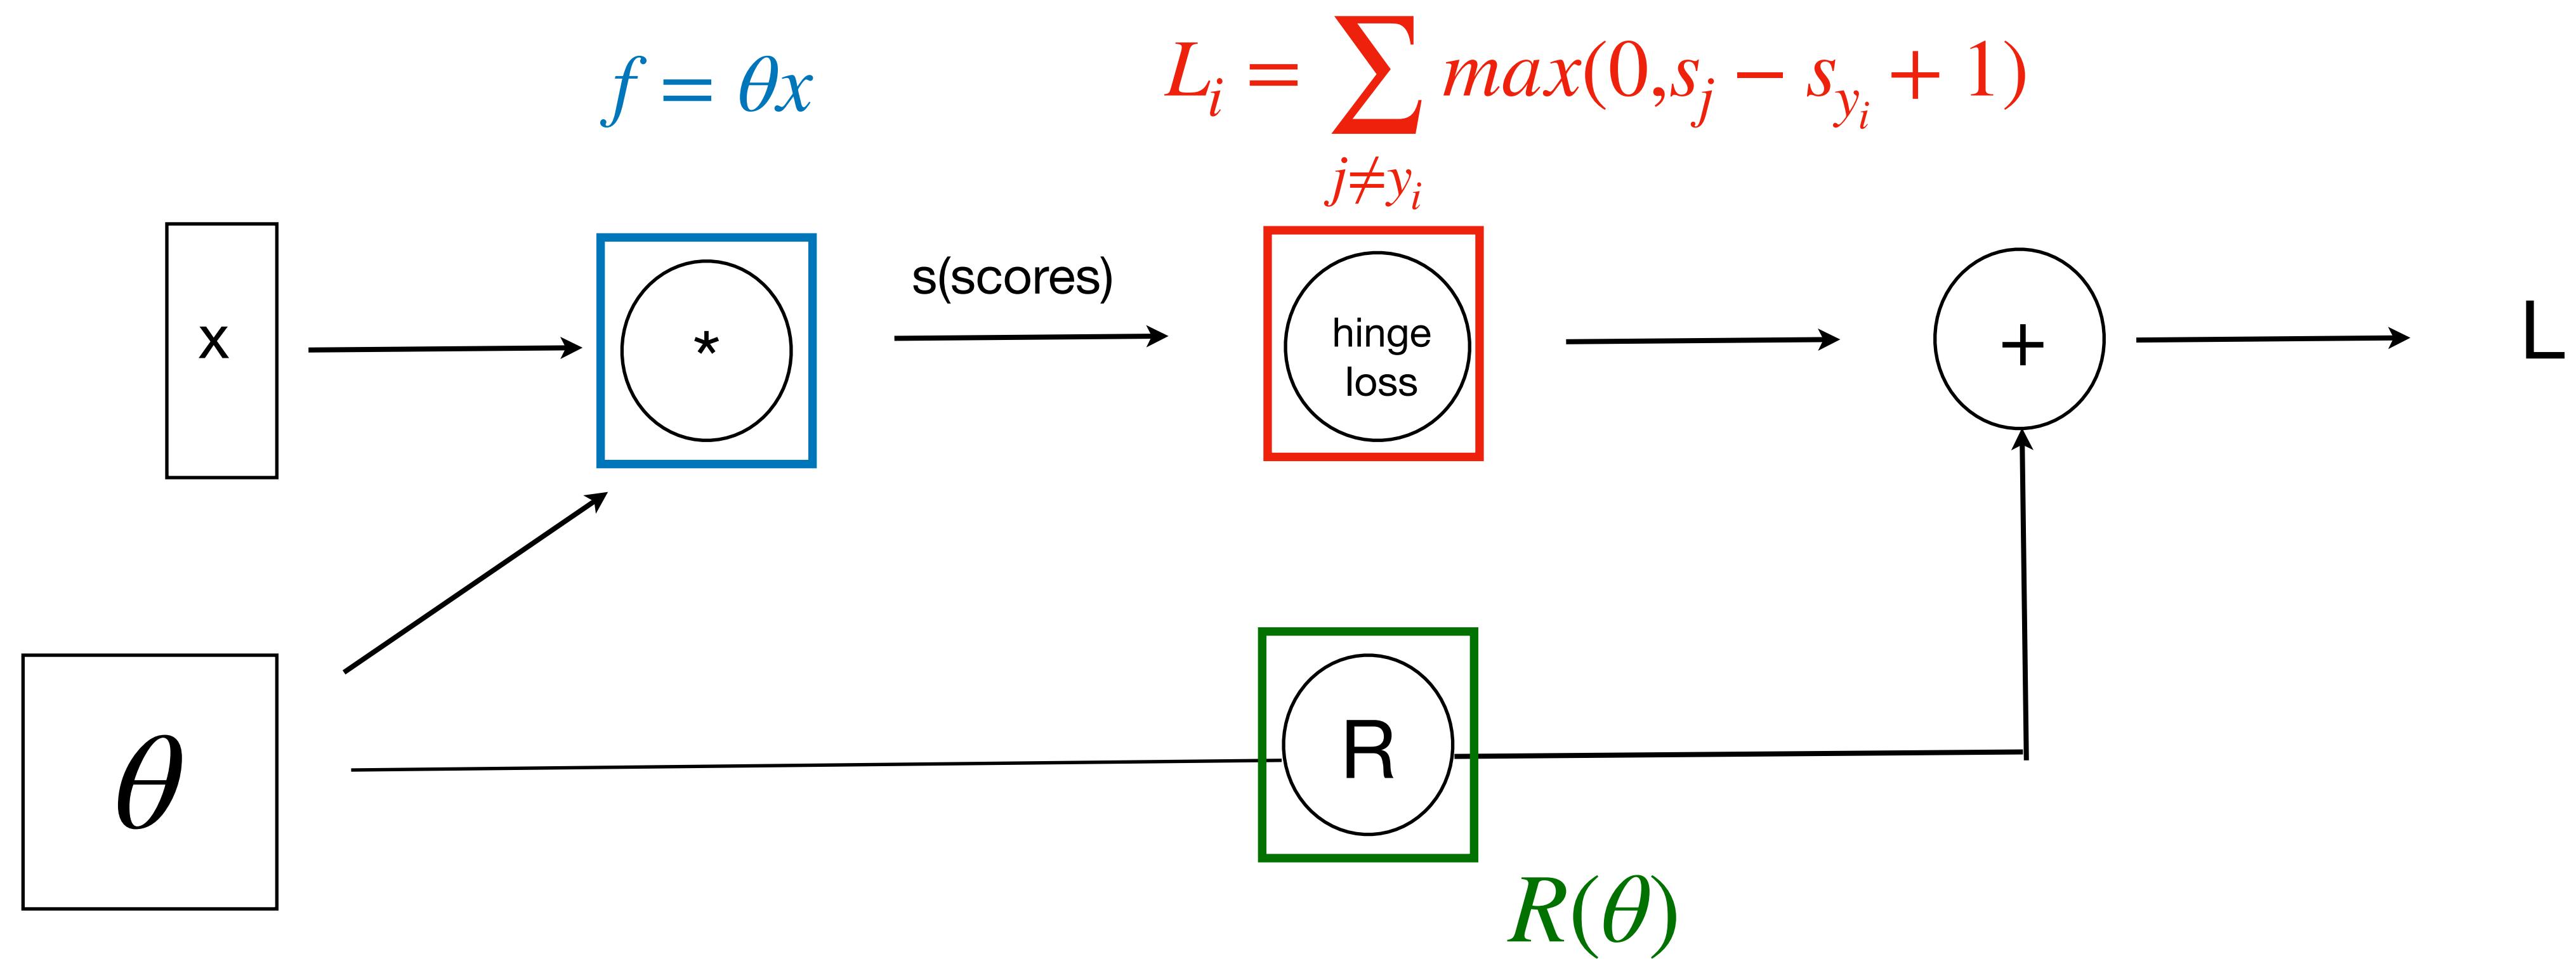
$$\epsilon(\hat{y}) = \frac{1}{2}(y - \hat{y})^2$$

Goal:
Gradient Descent

$$\theta^{t+1} = \theta^t - \lambda \frac{\partial \epsilon(\theta)}{\partial \theta}$$

Backpropagation

Computational graphs



Convolutional network

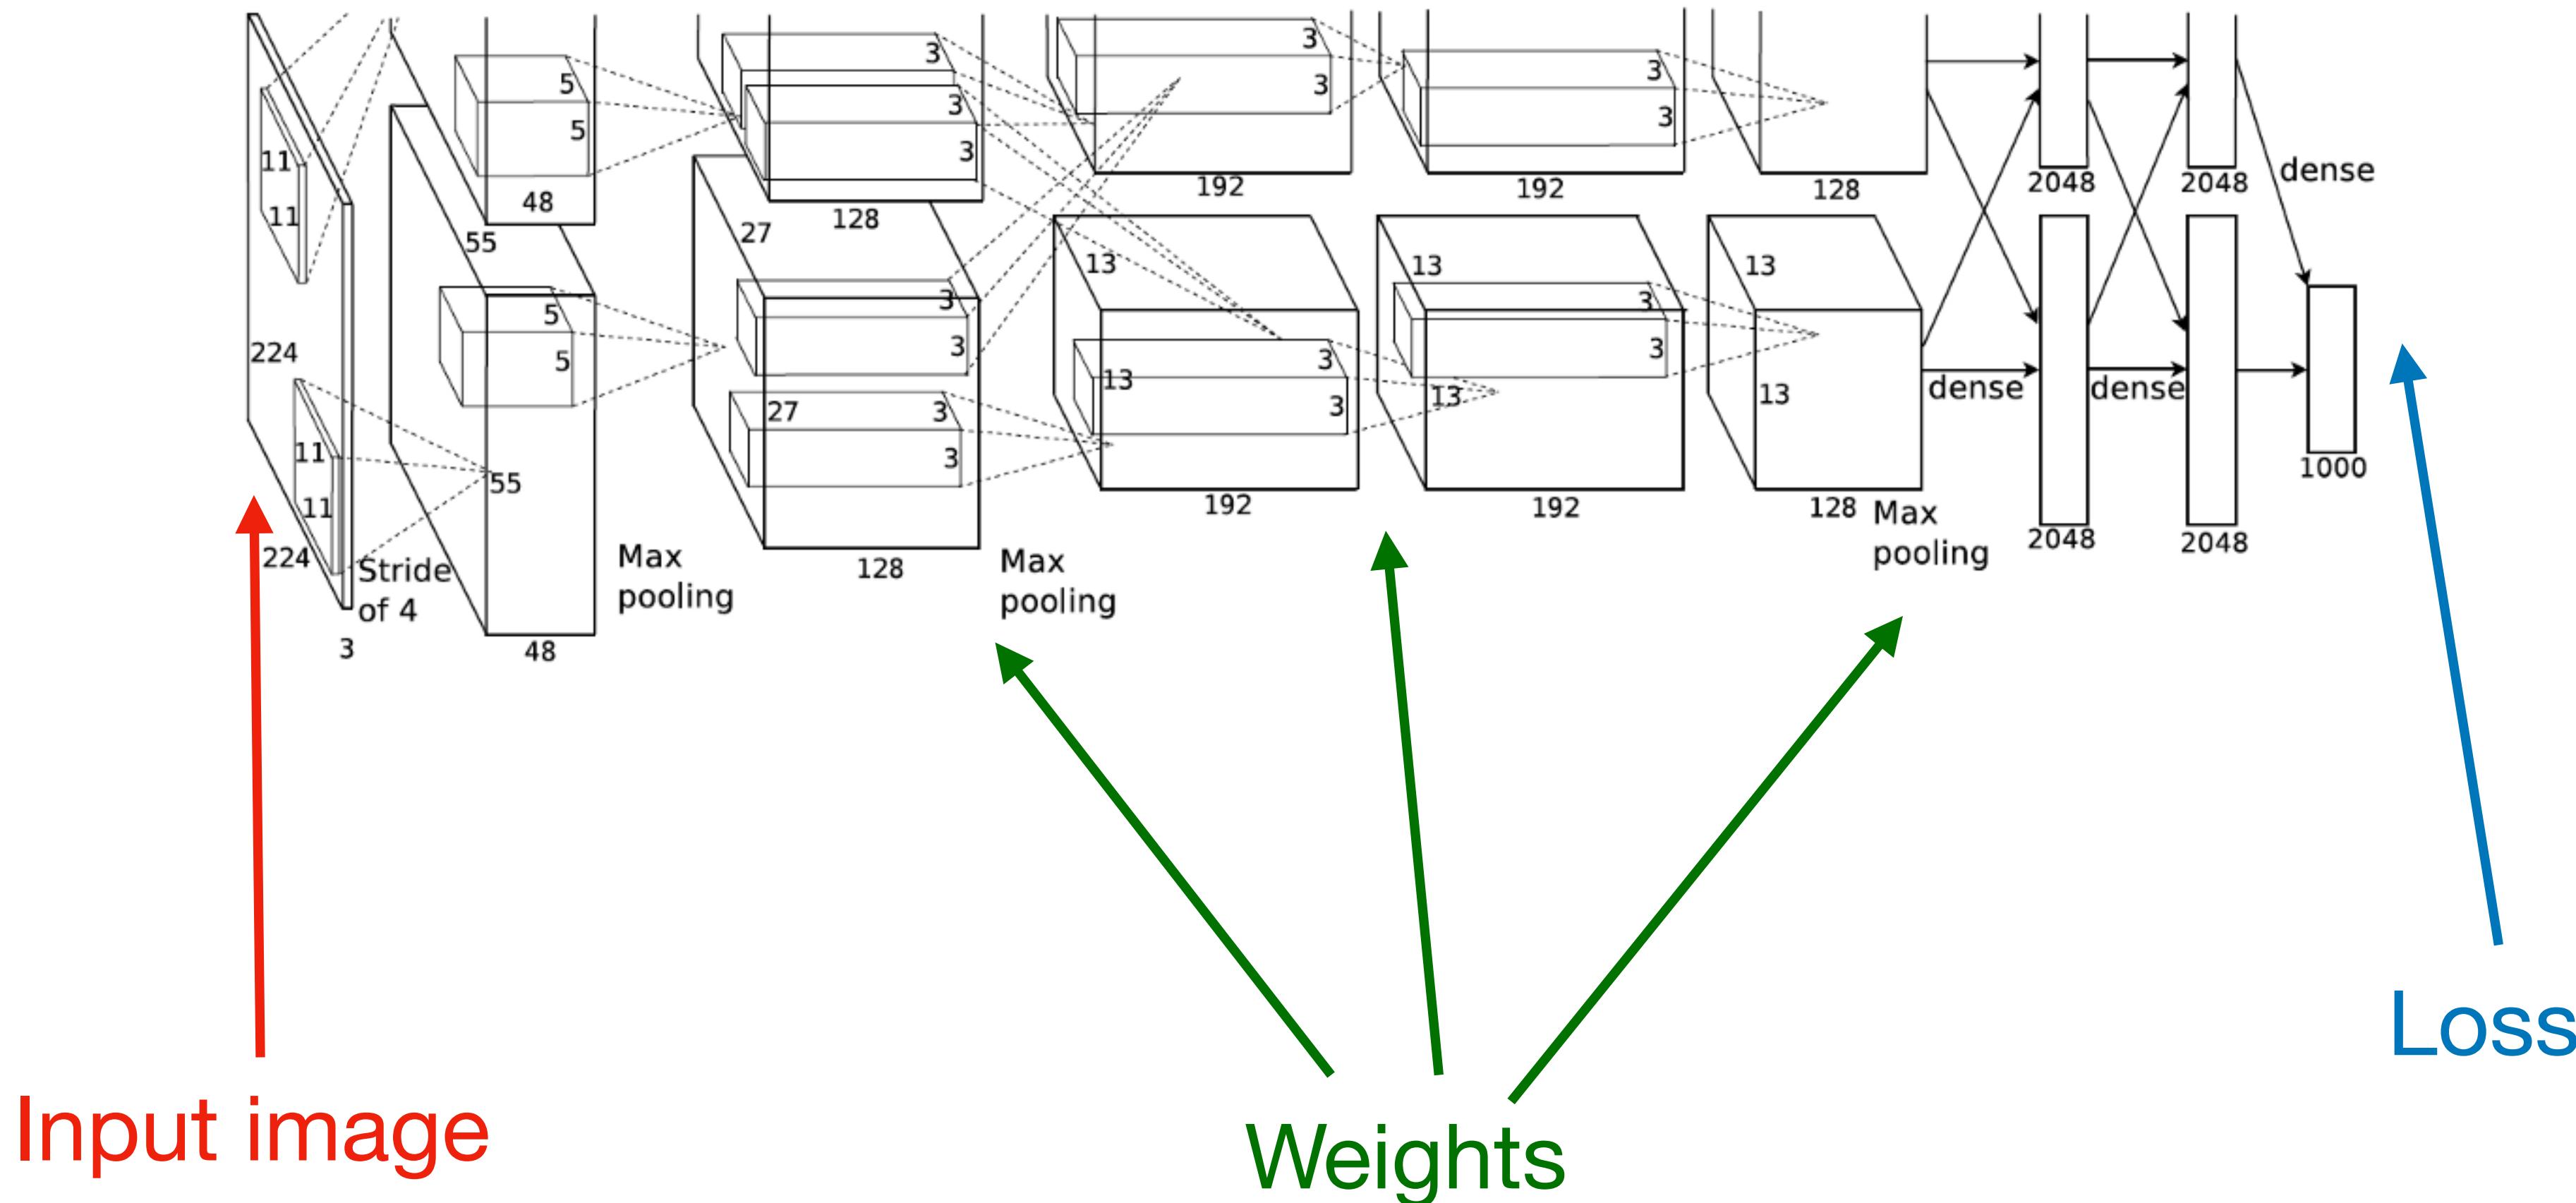
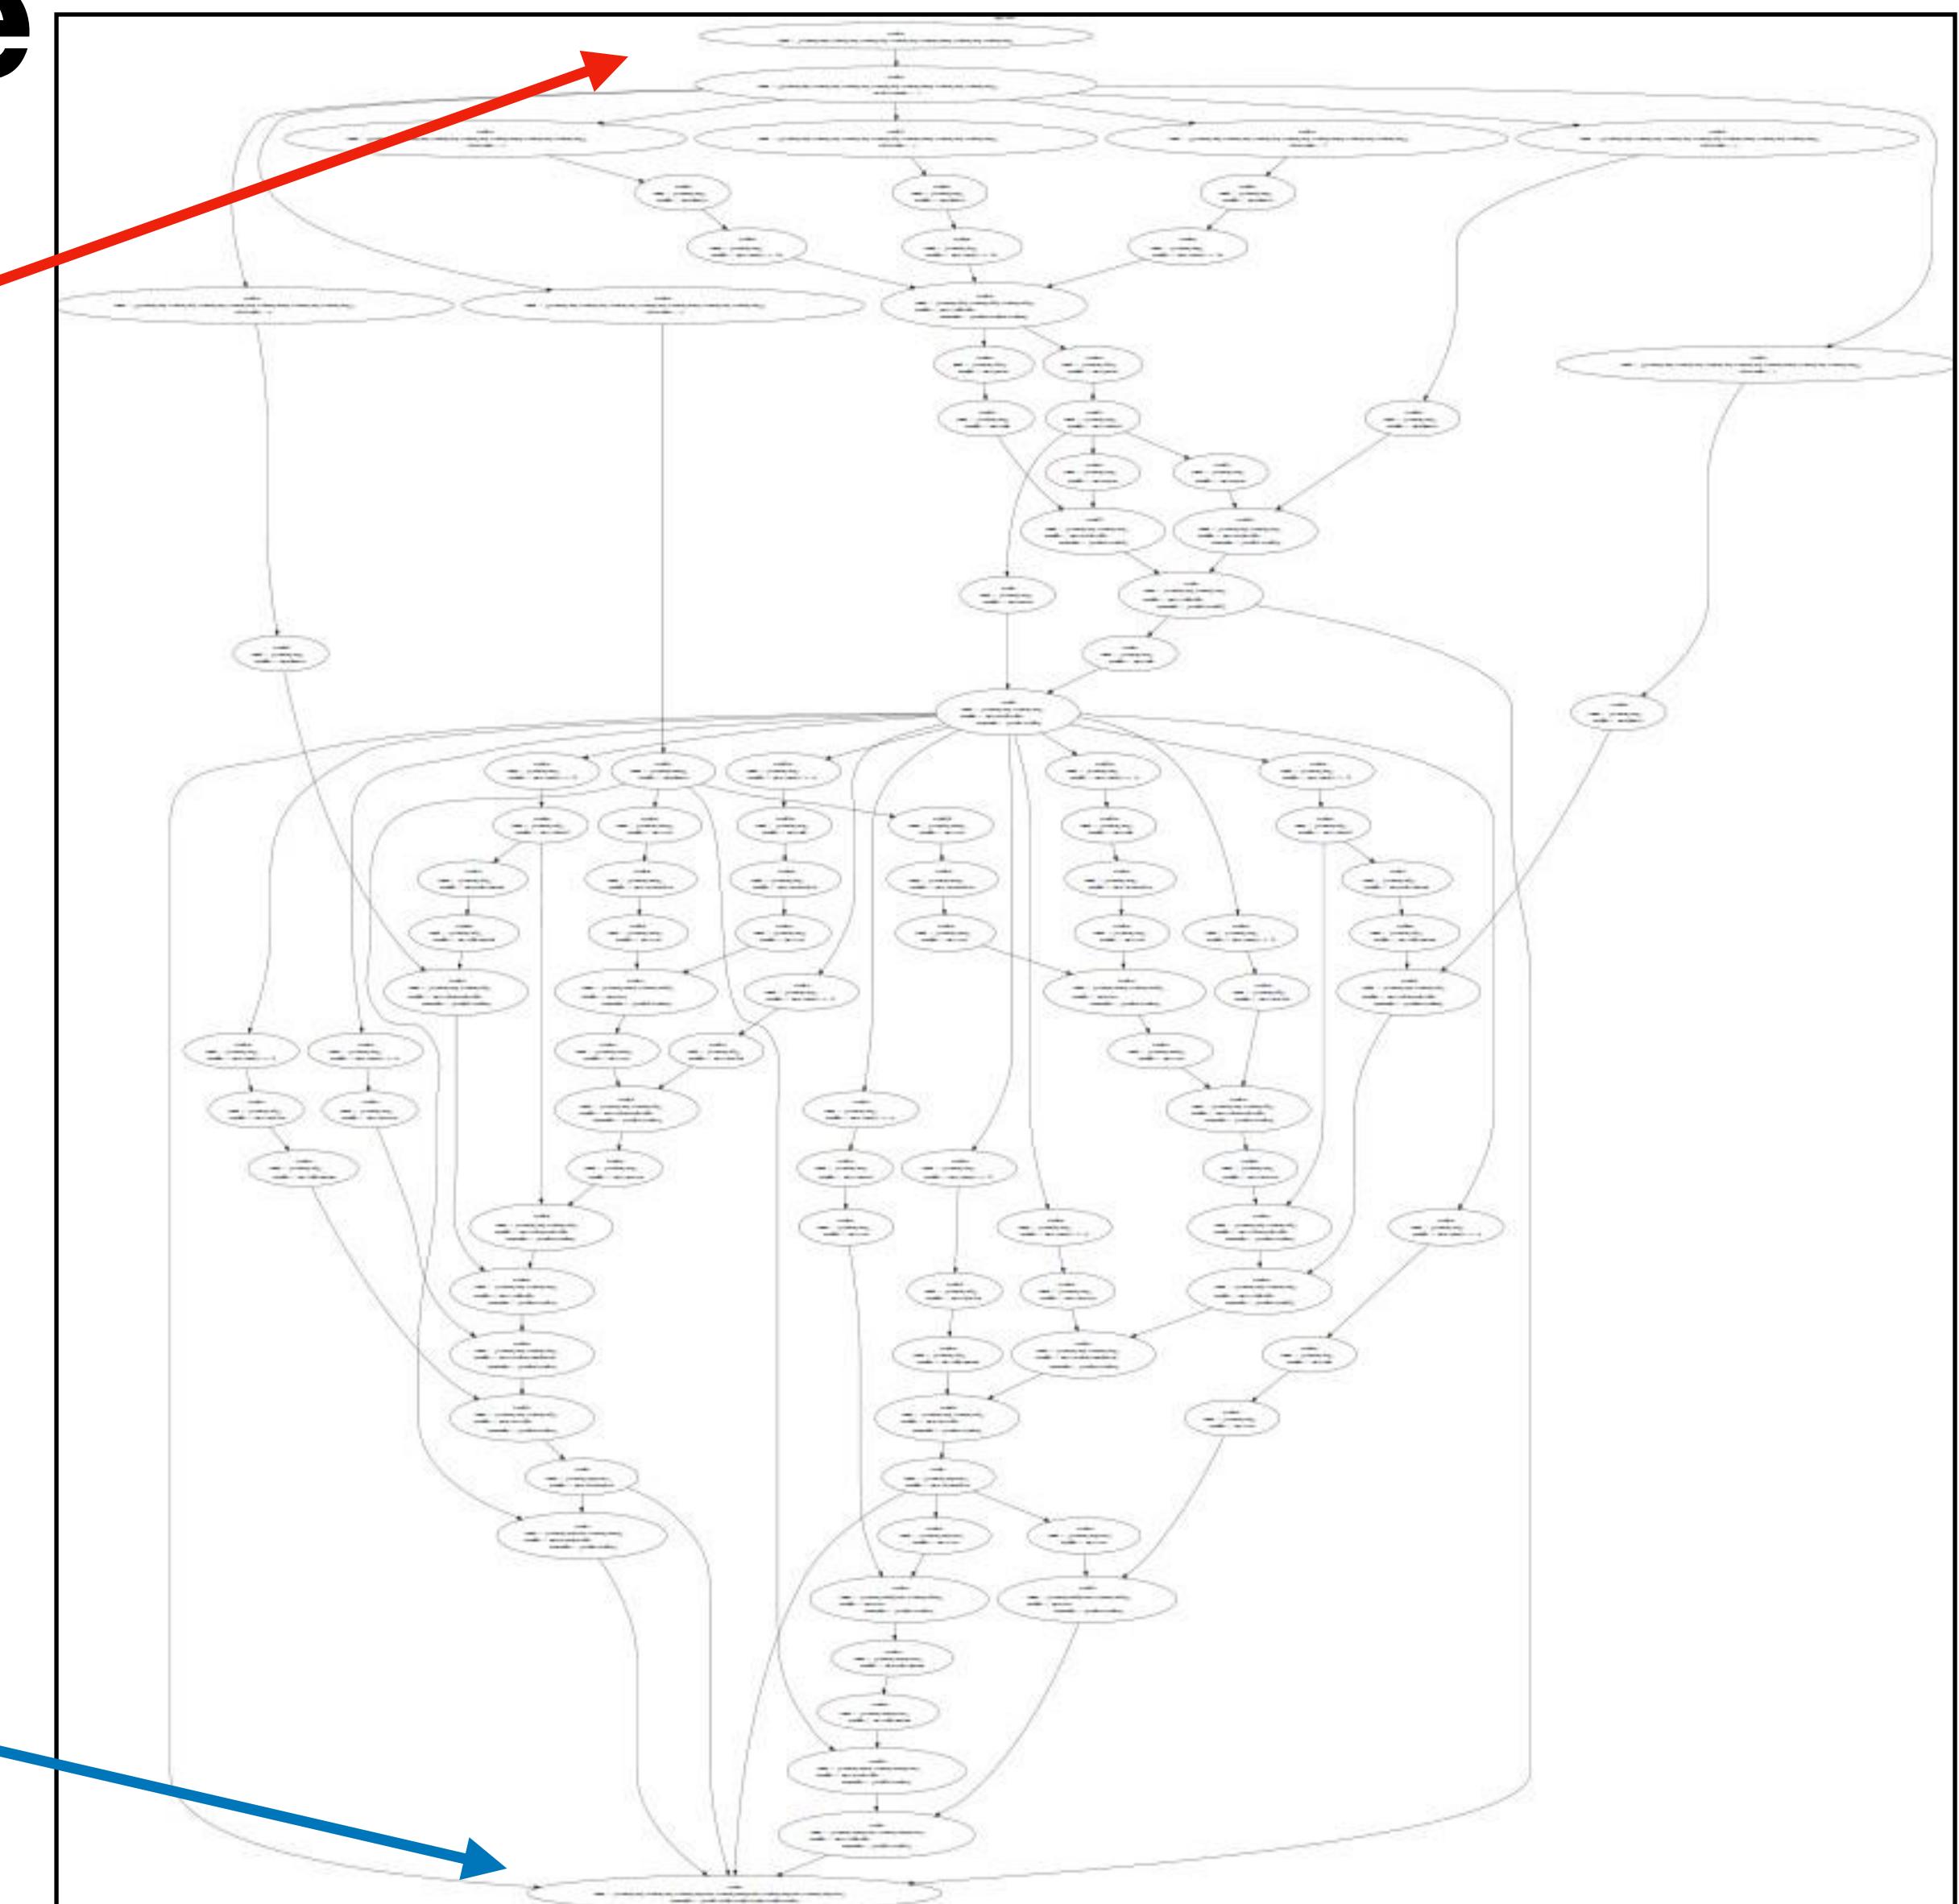


Figure copyright Alex Krizhevsky, Ilya Sutskever and Geoffrey Hinton, 2012

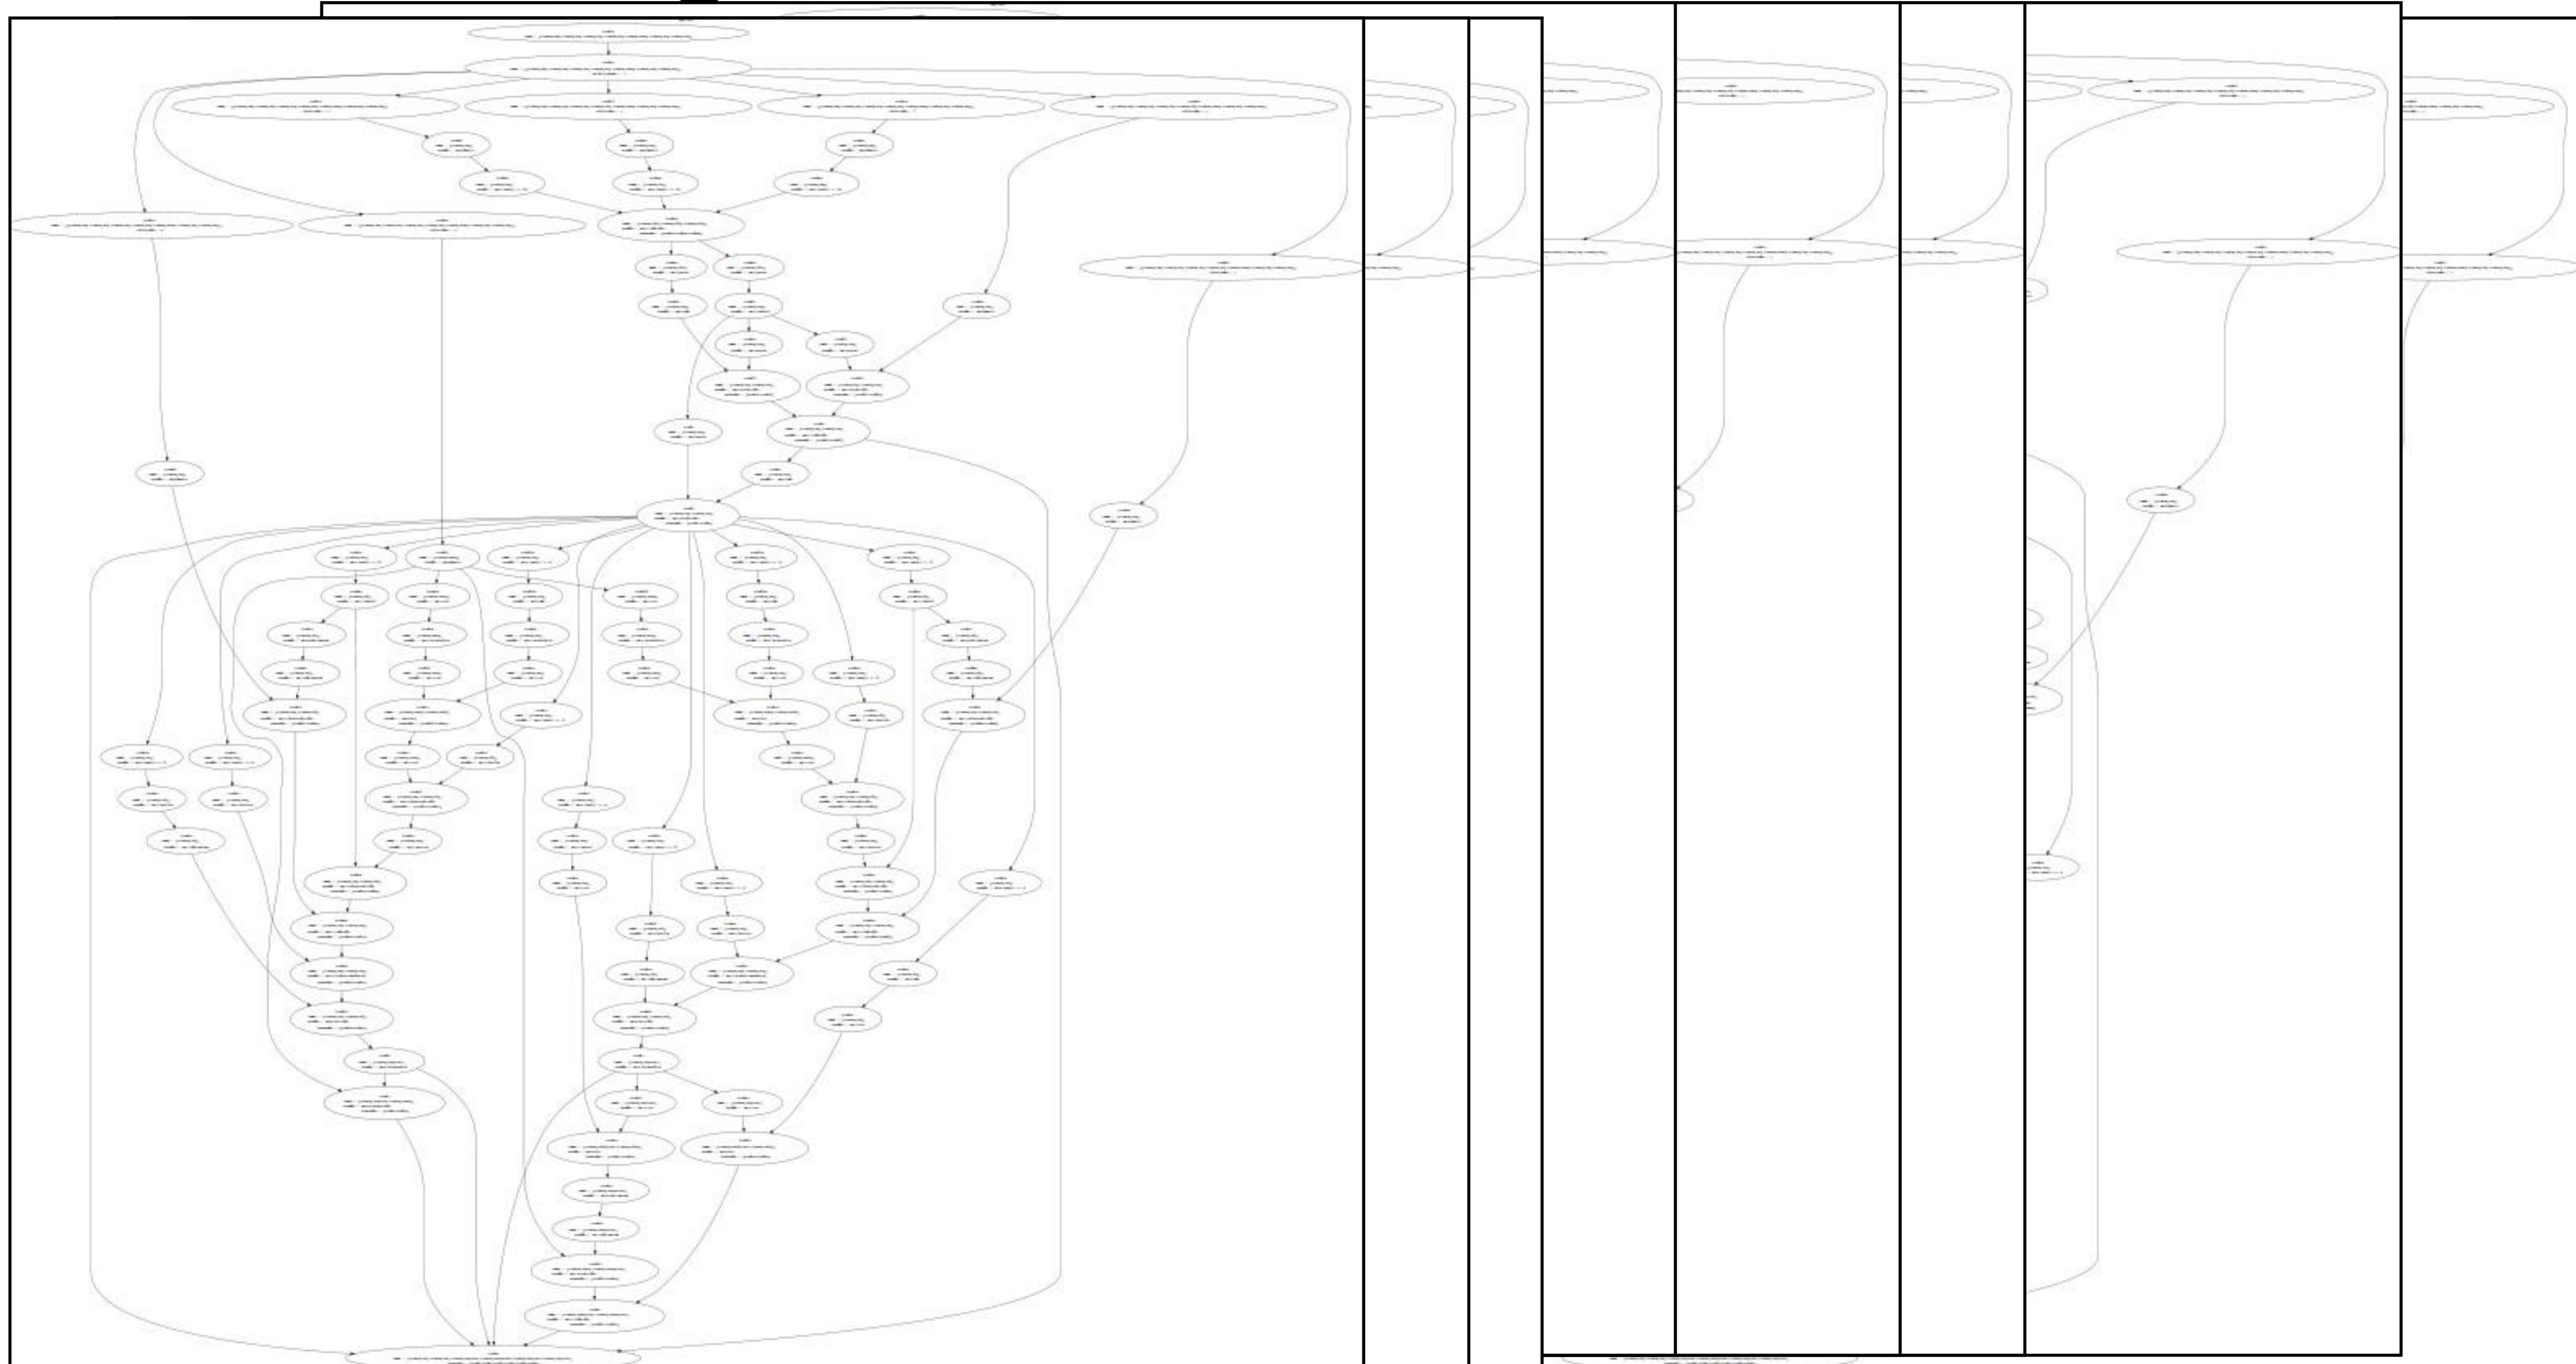
Neural Turing machine

Input image

Loss



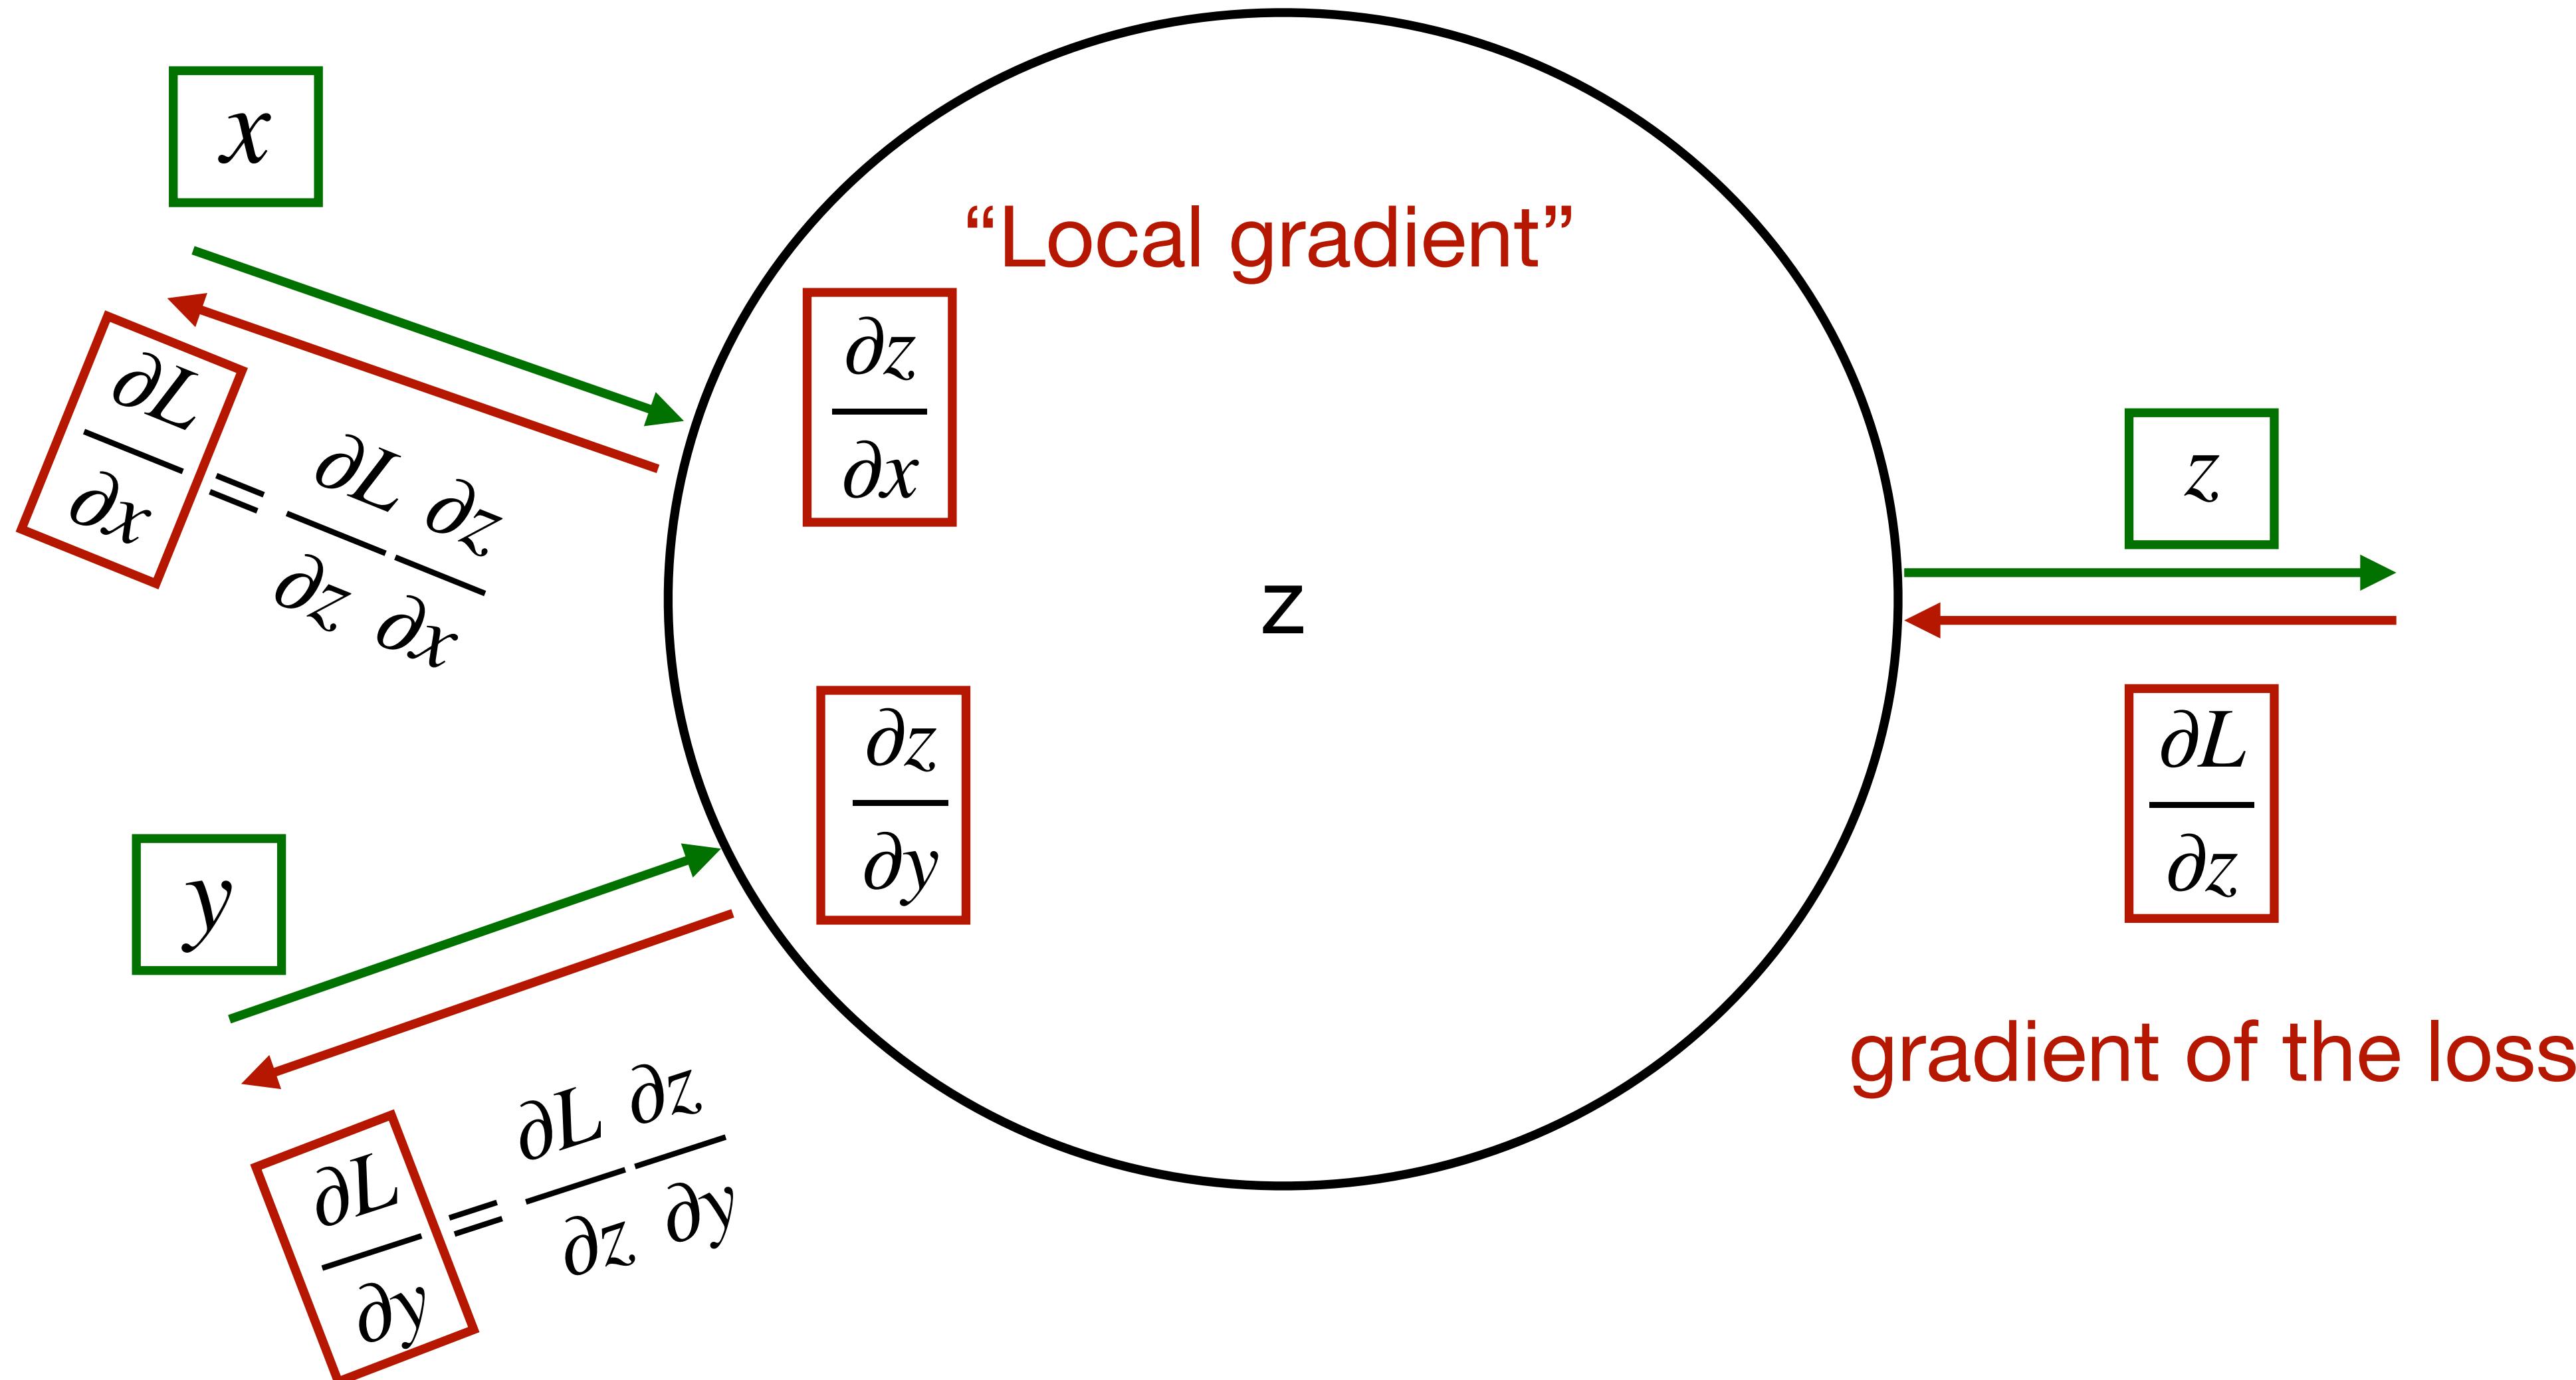
Neural Turing machine



Backpropagation

Chain rule

$$\frac{\partial L(z(x))}{\partial x} = \frac{\partial L}{\partial z} \cdot \frac{\partial z}{\partial x}$$



Backpropagation

A simple example

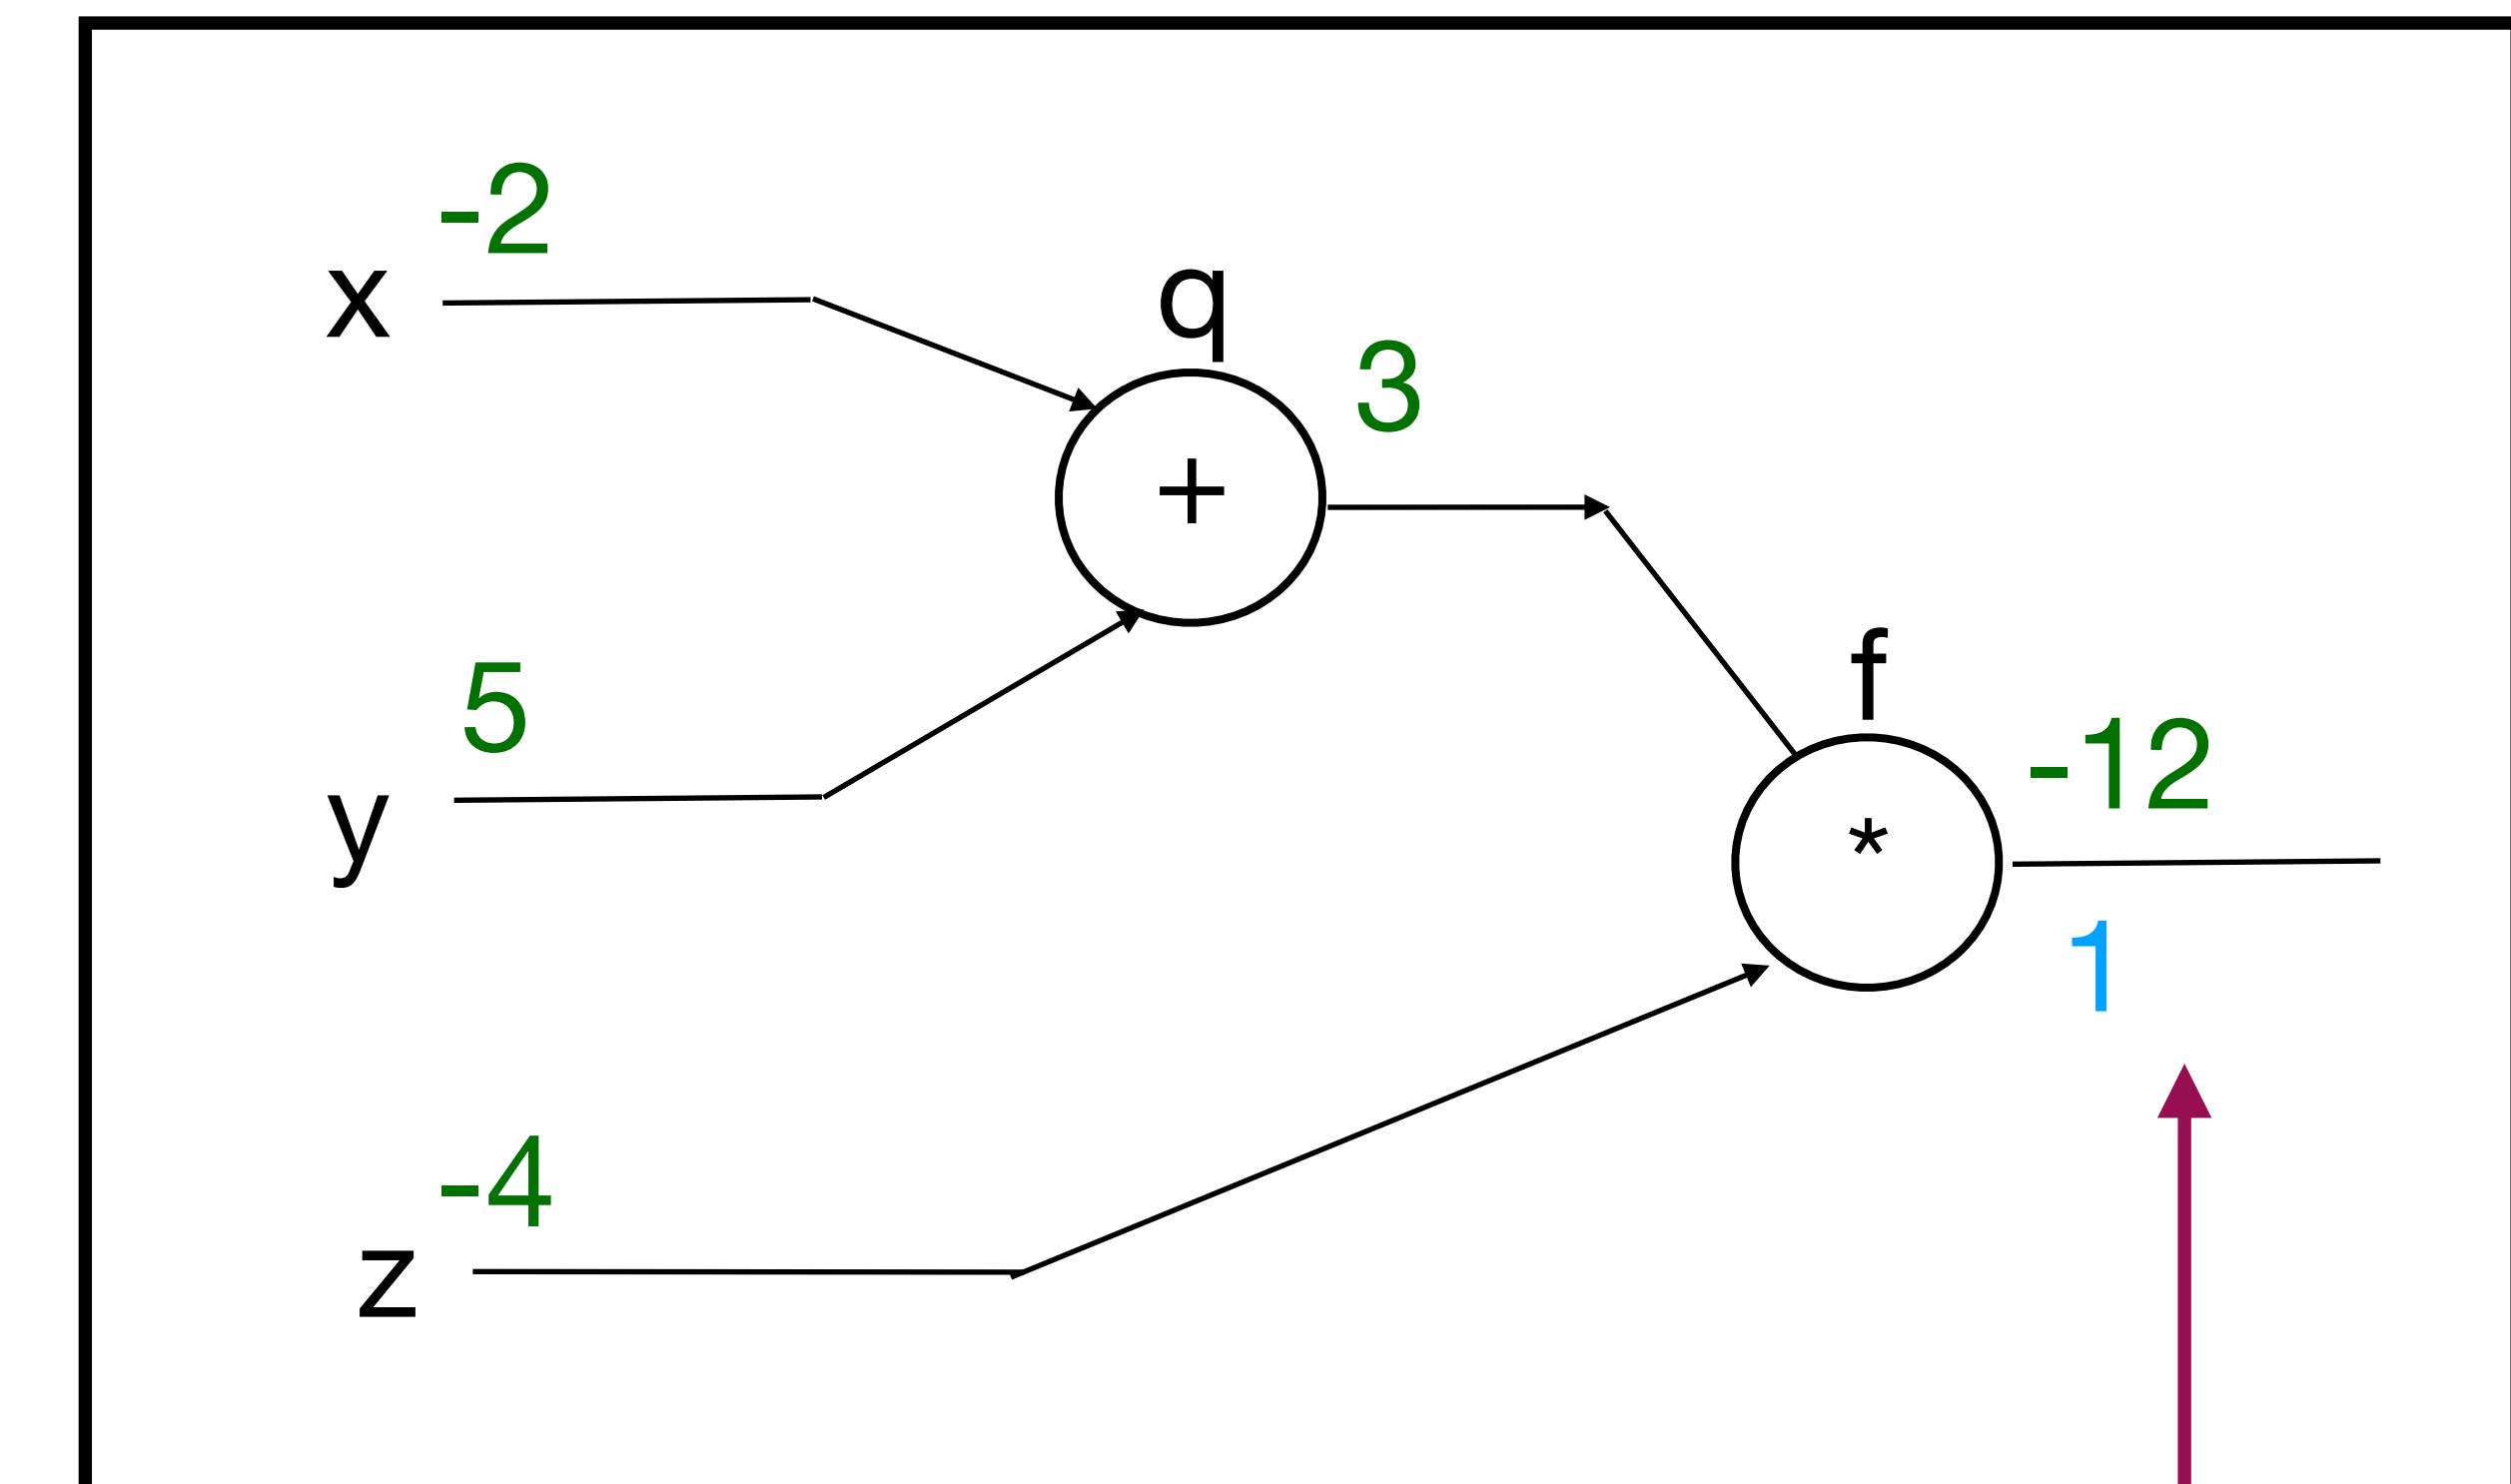
$$f(x, y, z) = (x + y)z$$

Let $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial f}$$

Backpropagation

A simple example

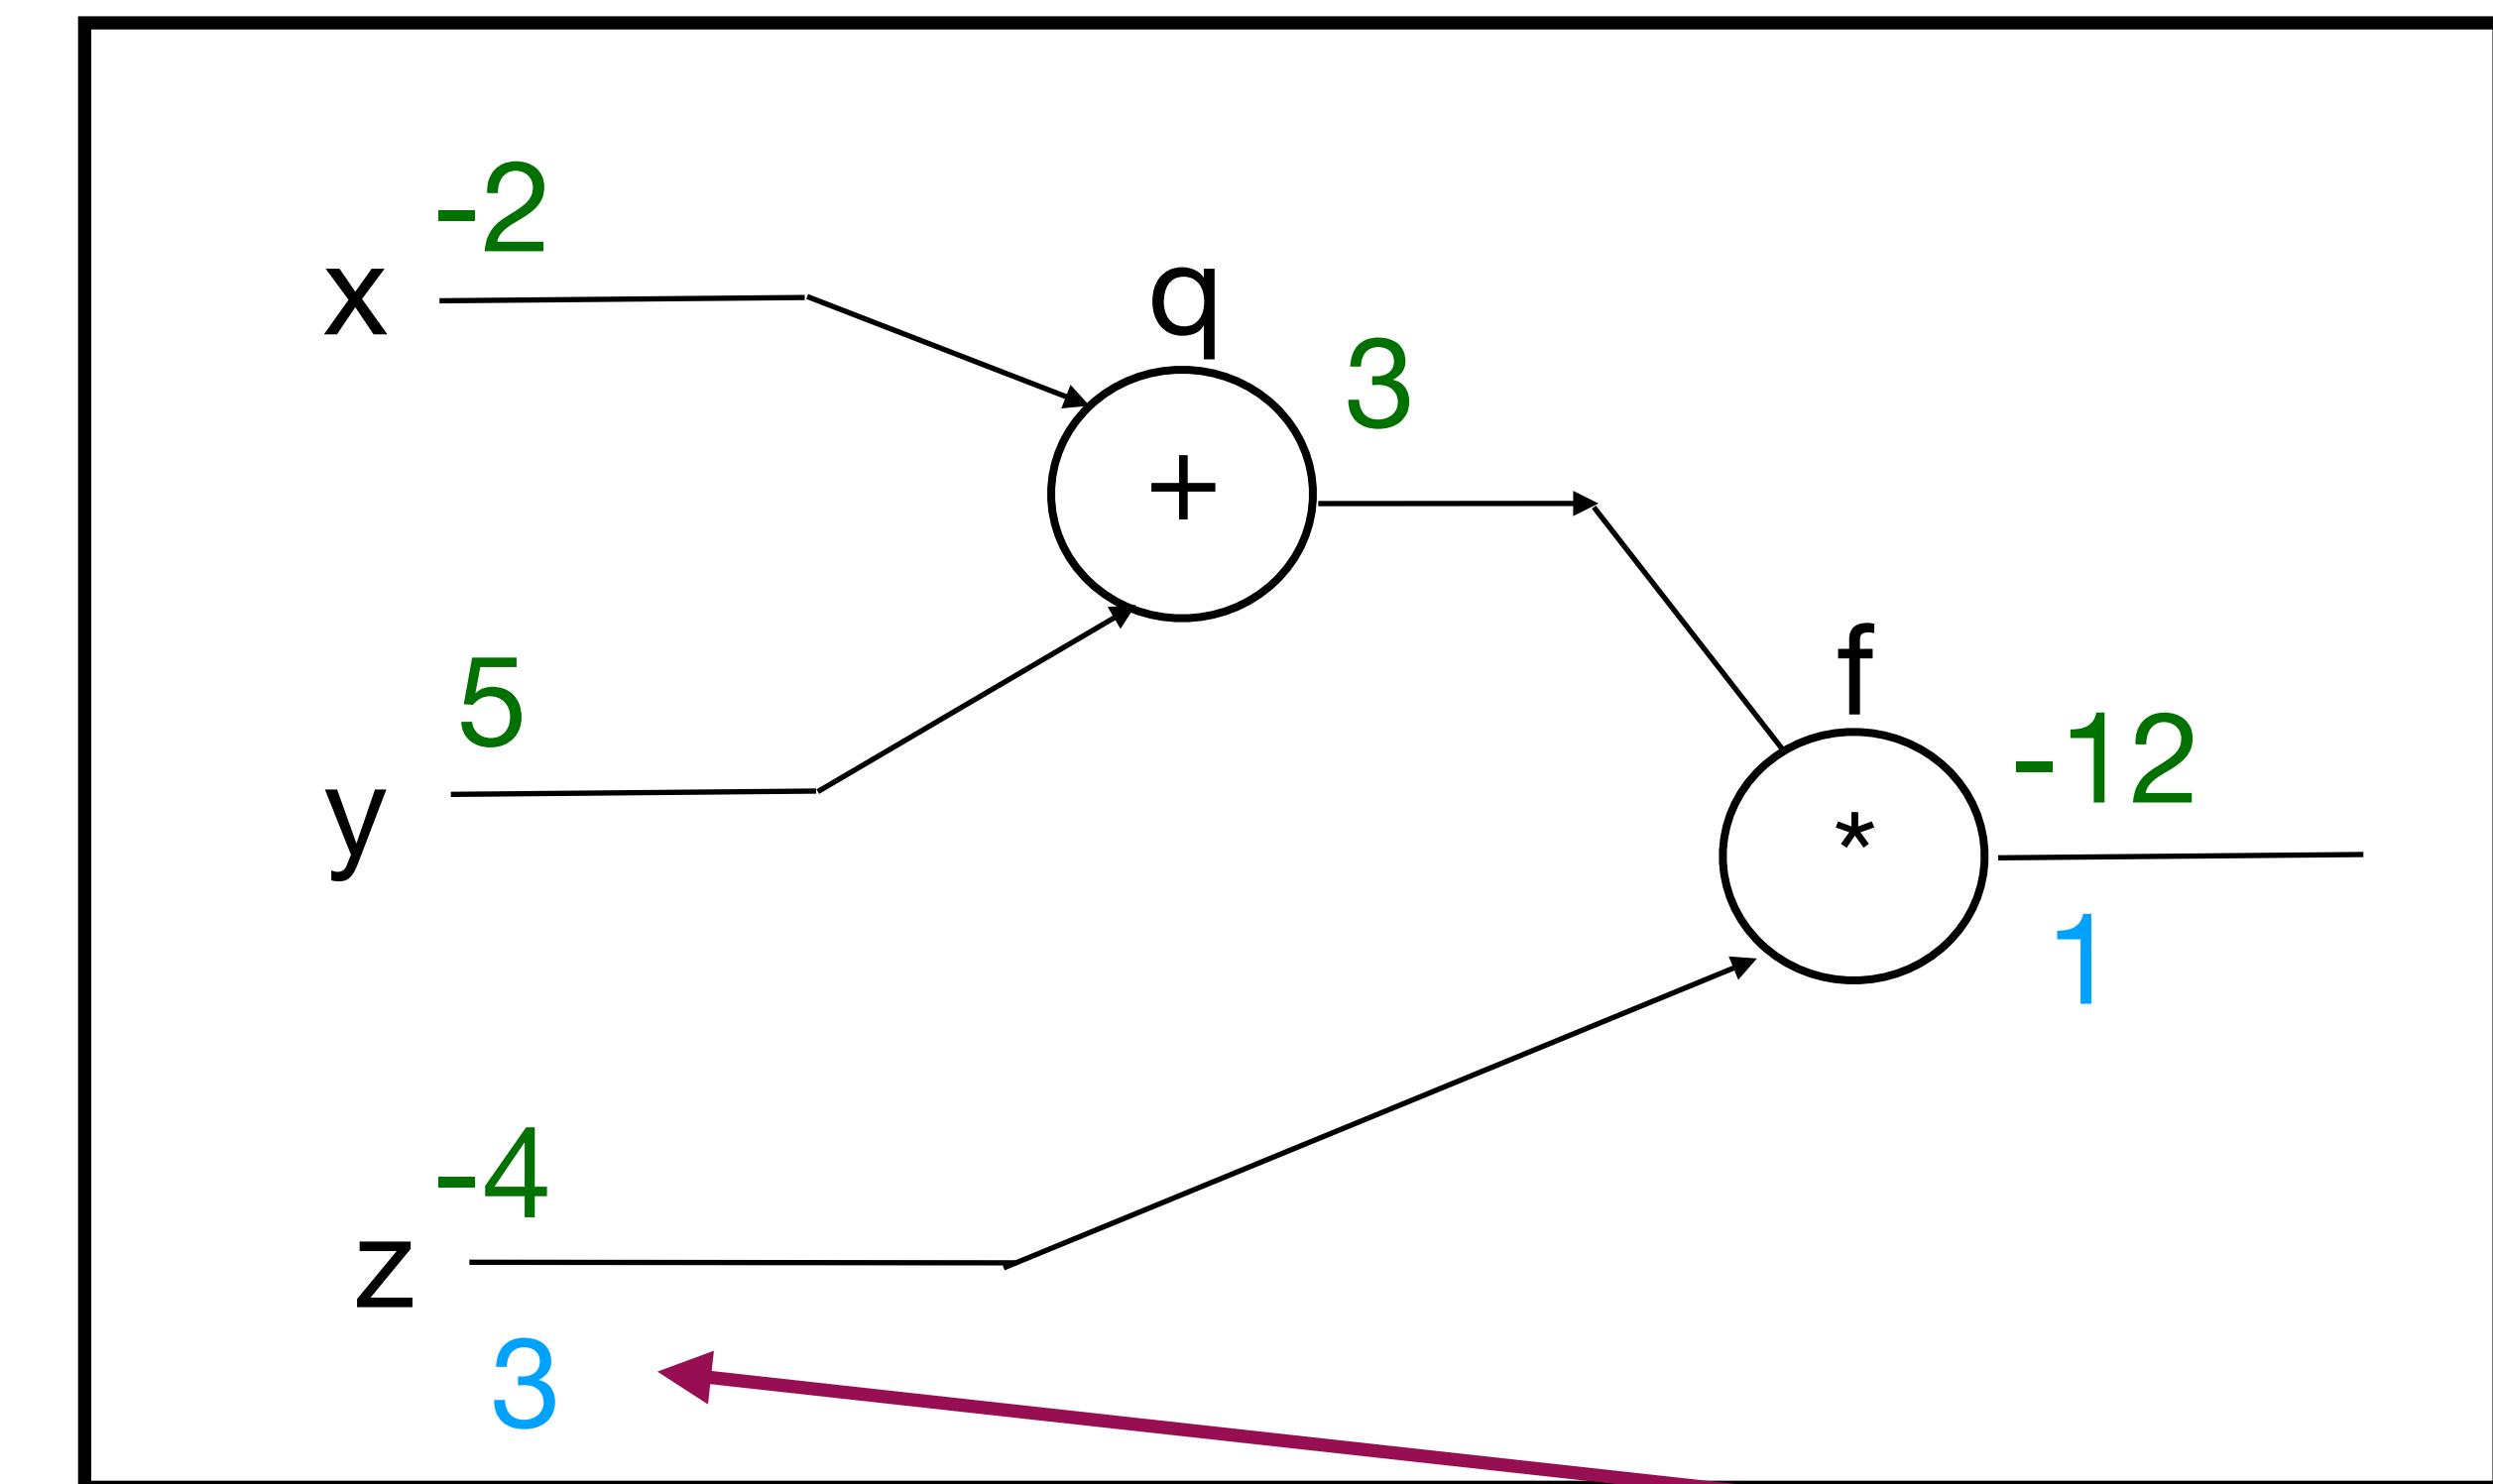
$$f(x, y, z) = (x + y)z$$

Let $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$

Backpropagation

A simple example

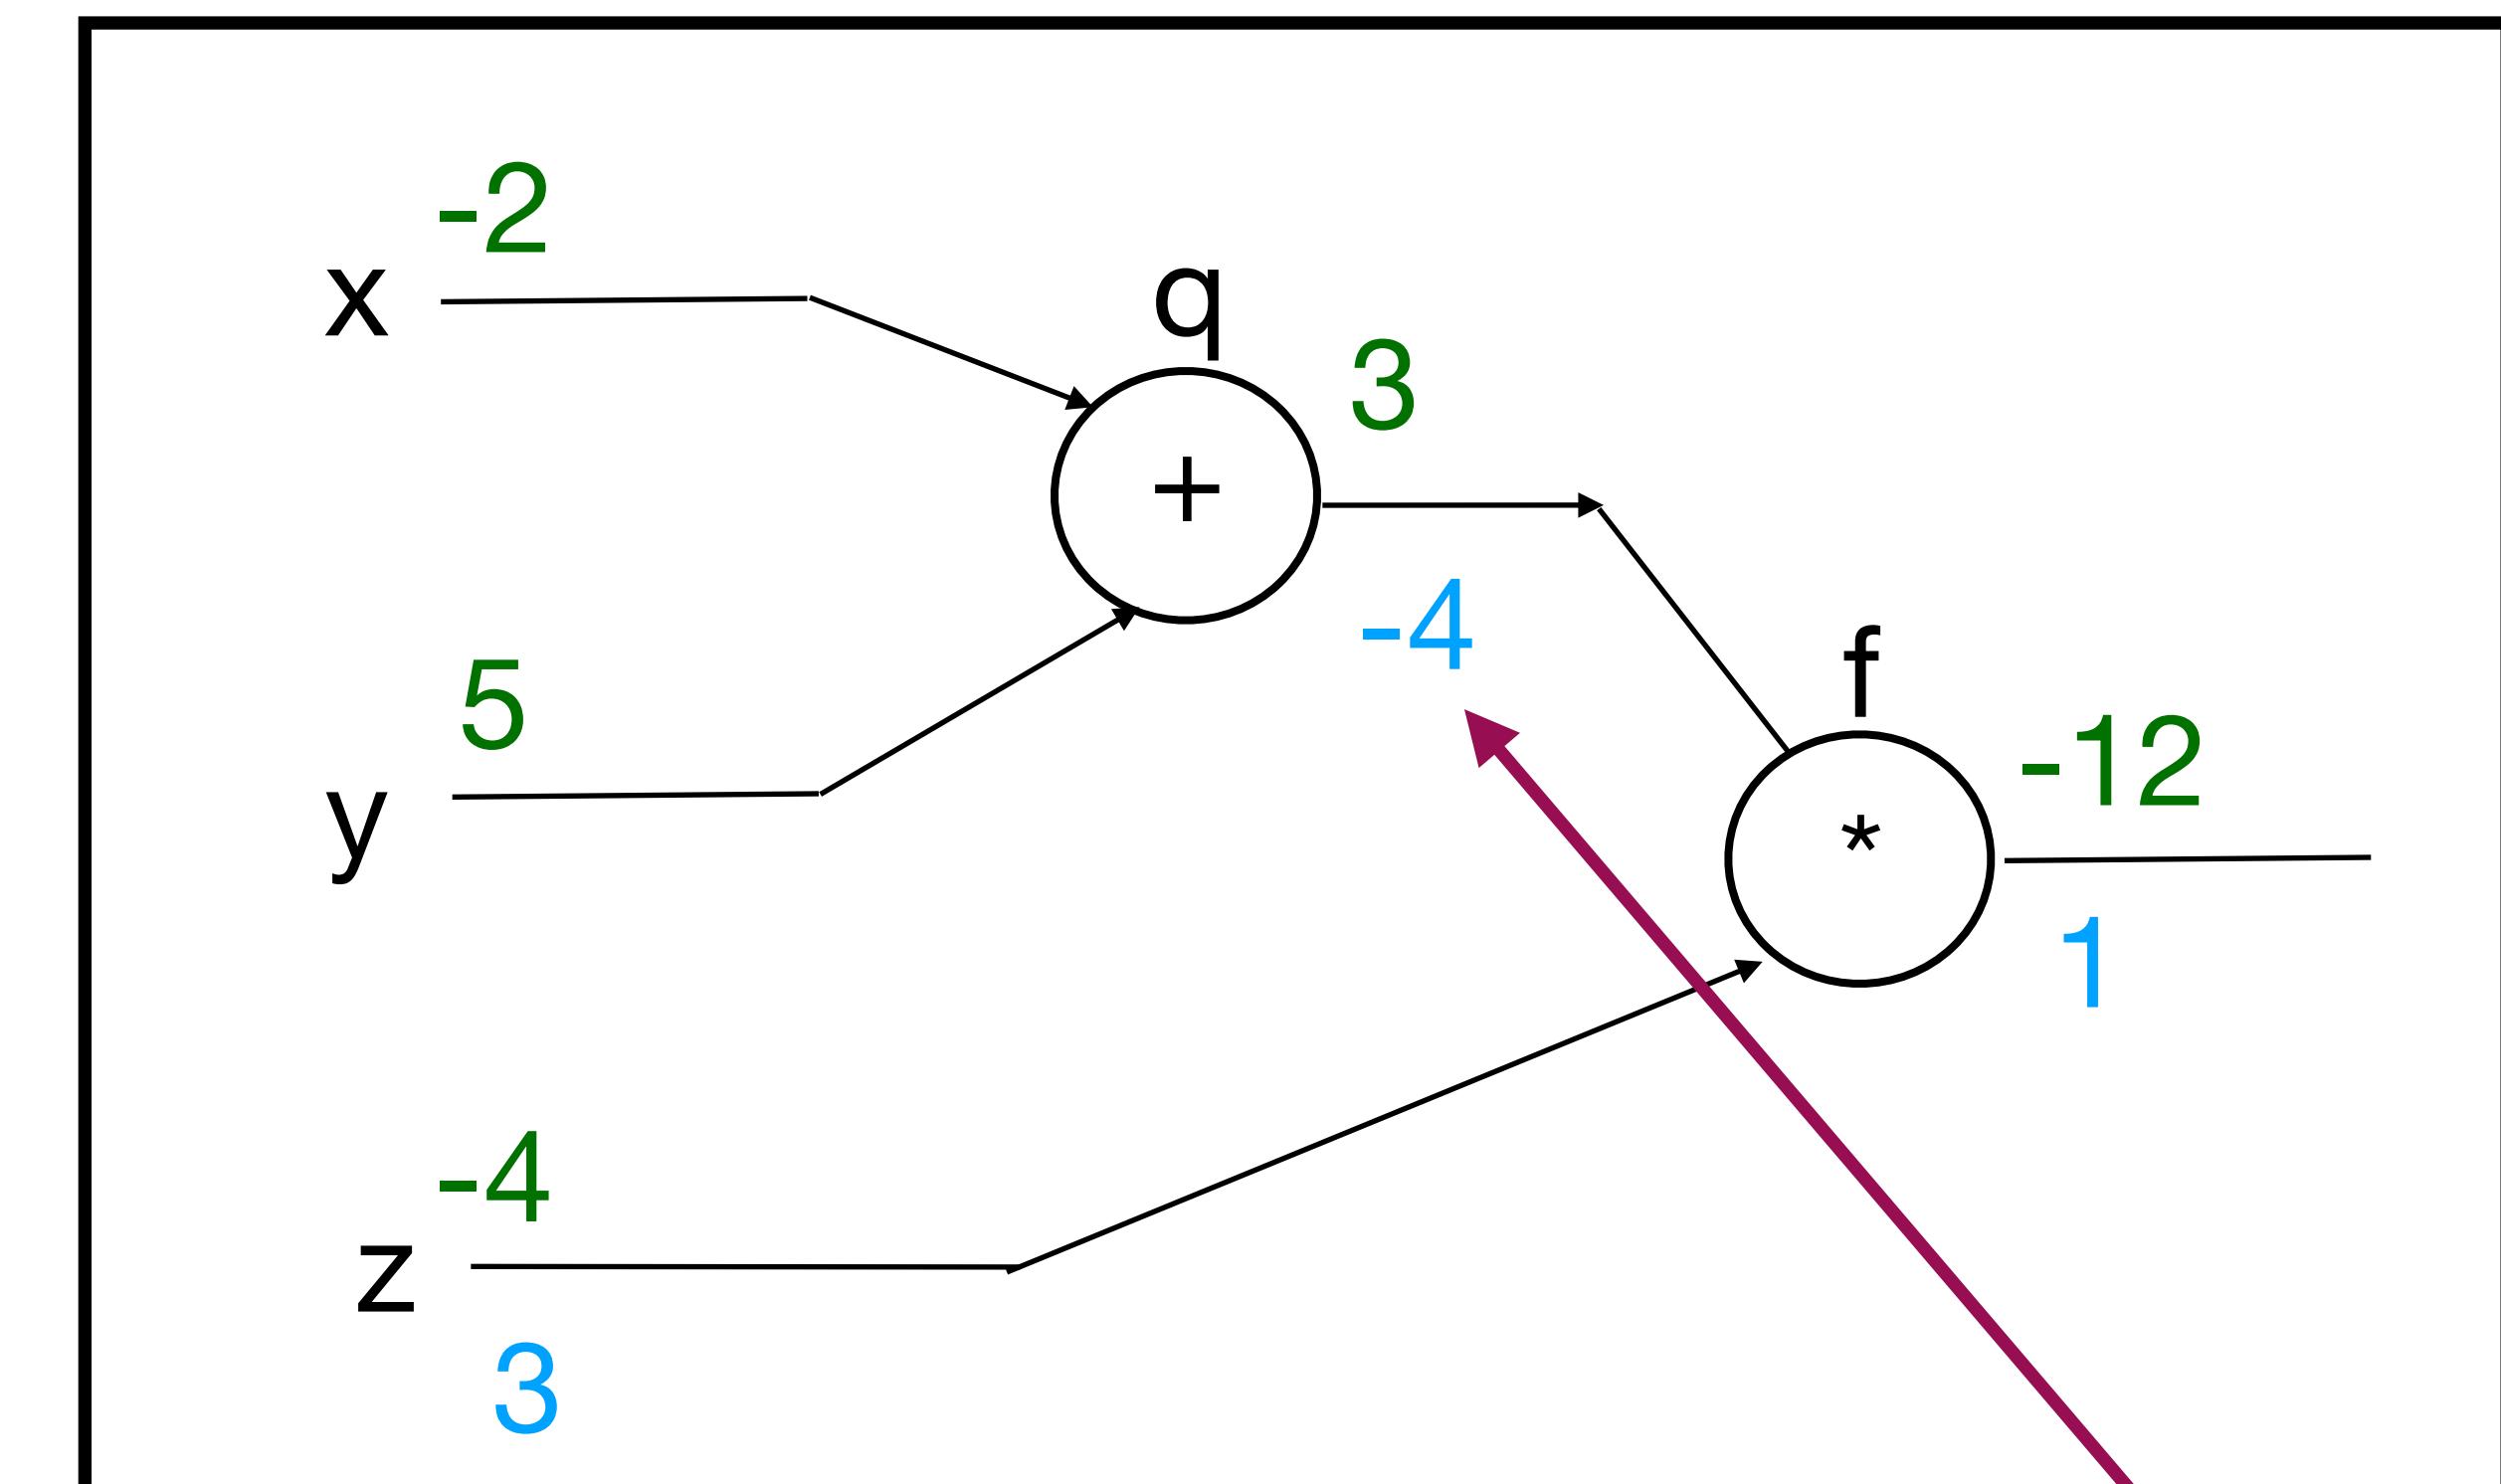
$$f(x, y, z) = (x + y)z$$

Let $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial q}$$

Backpropagation

A simple example

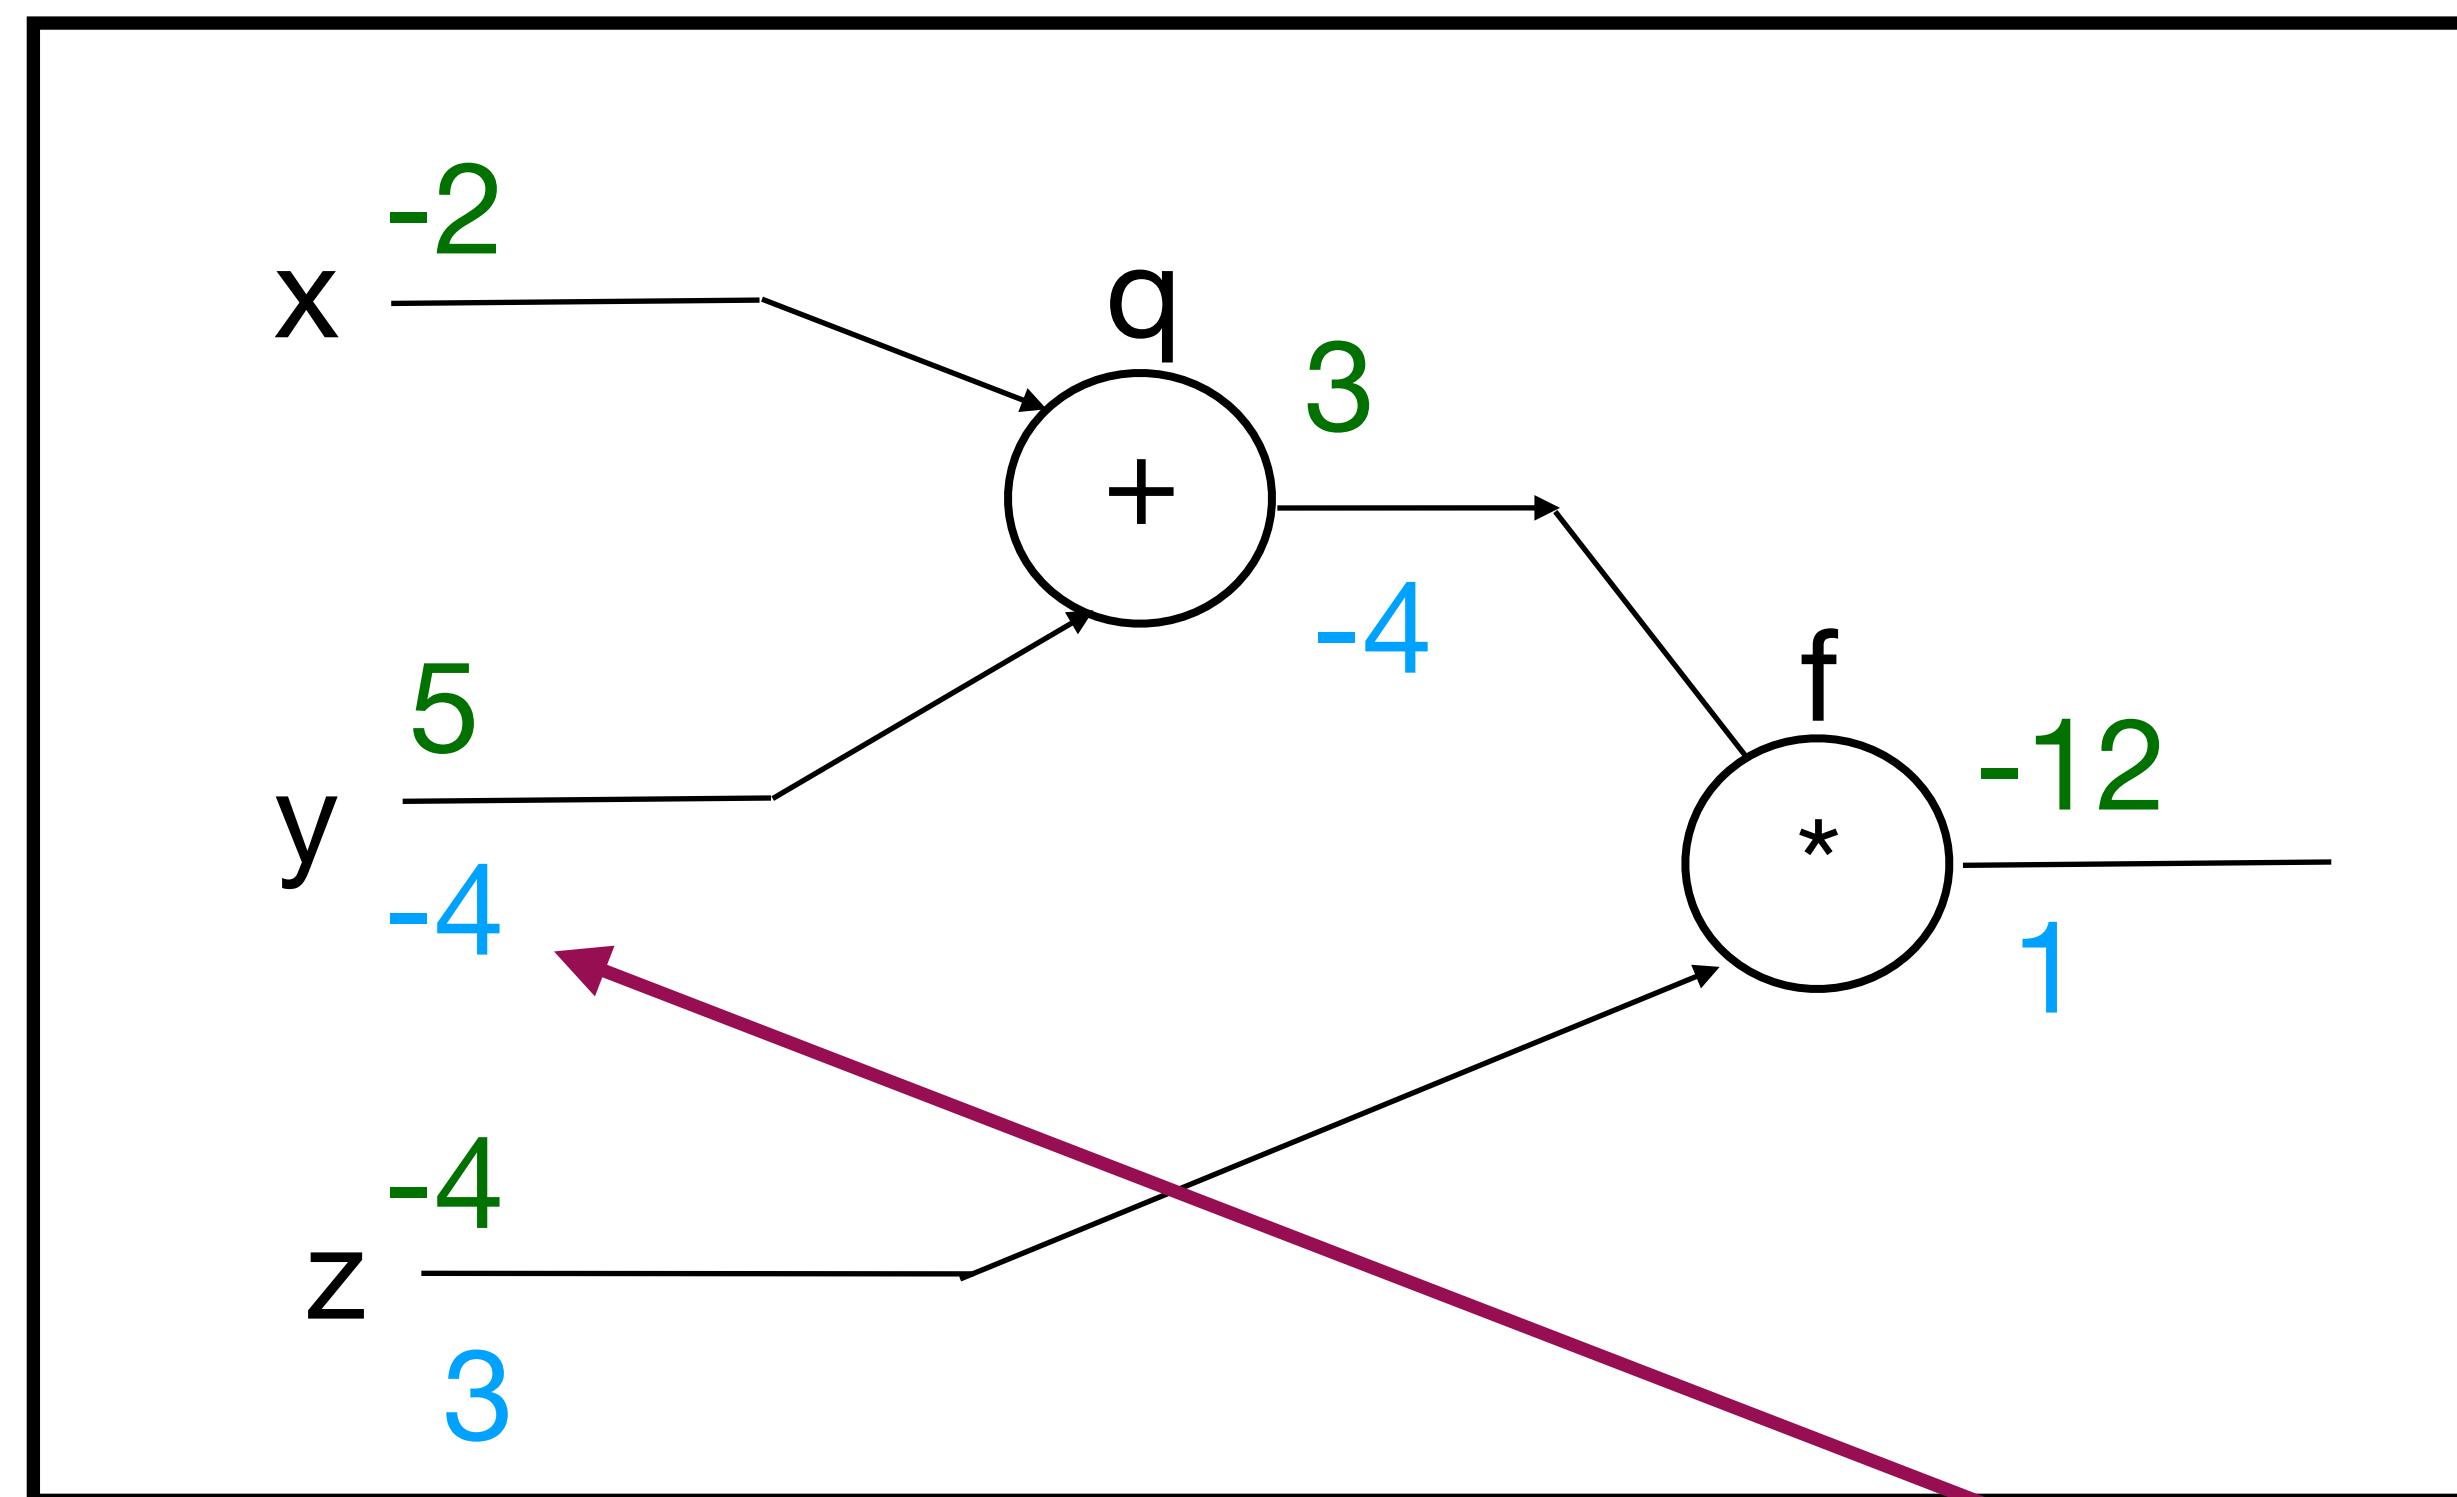
$$f(x, y, z) = (x + y)z$$

Let $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain Rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

$$\frac{\partial f}{\partial y}$$

Backpropagation

A simple example

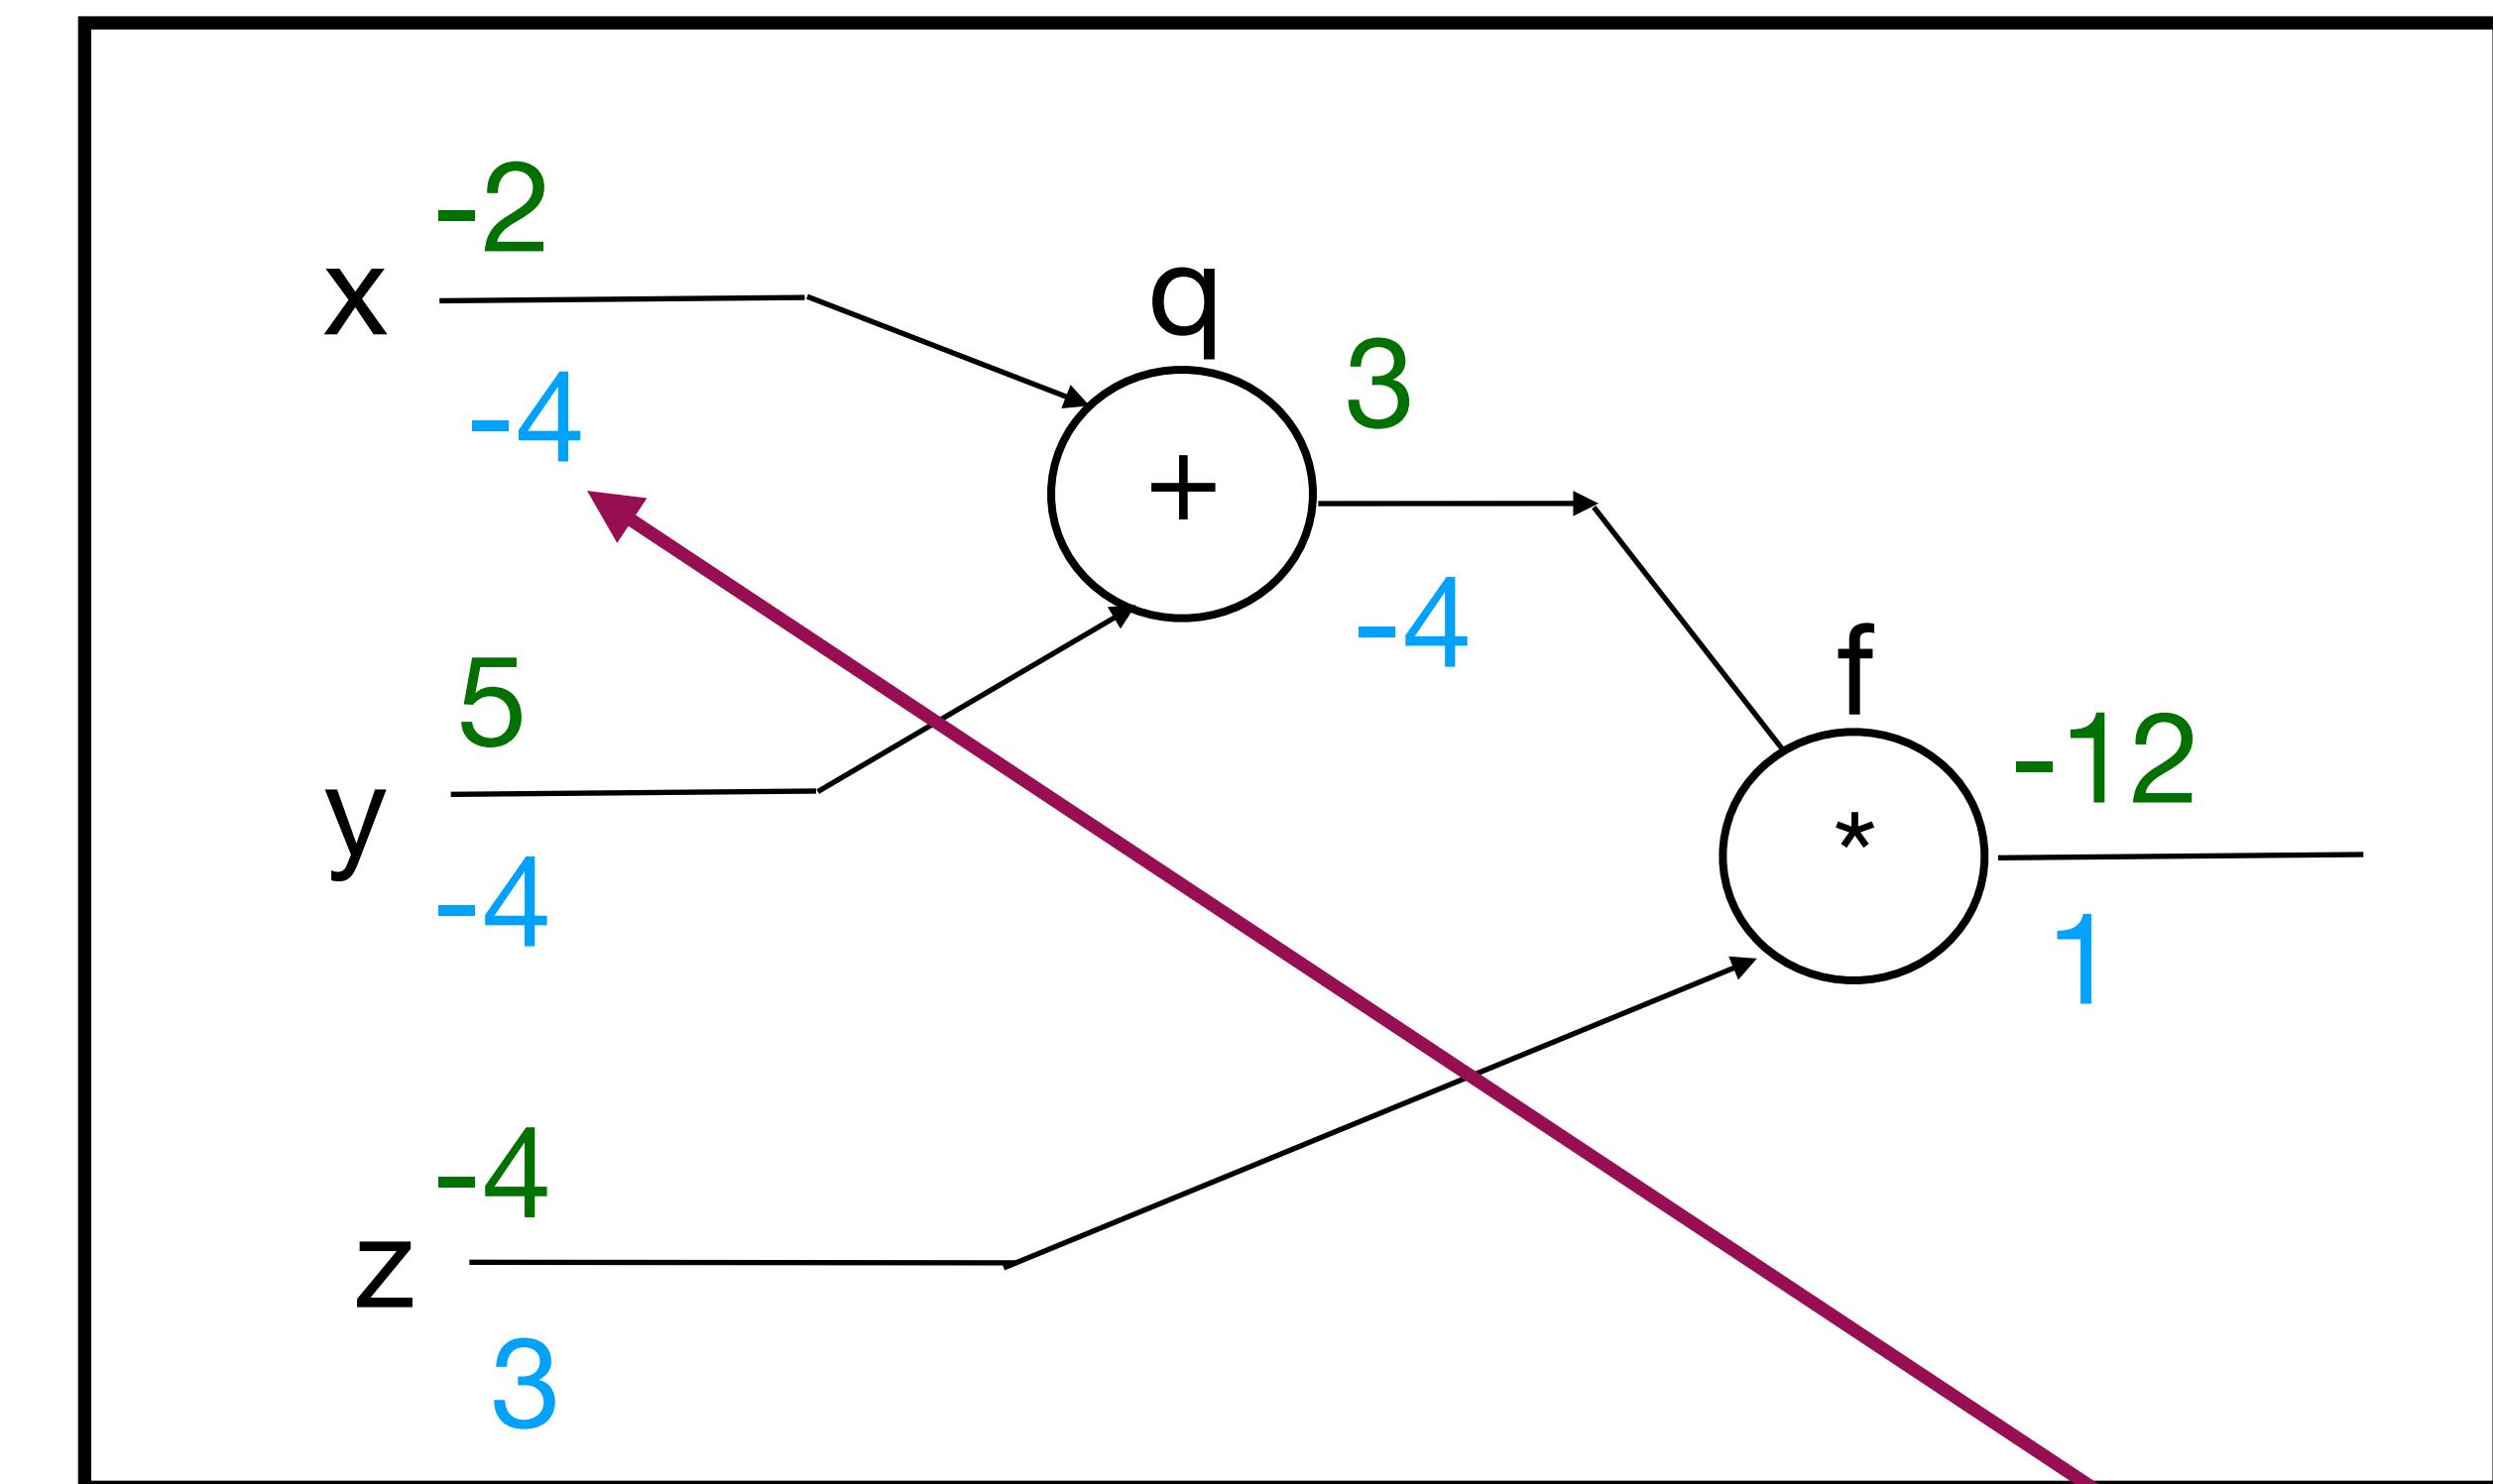
$$f(x, y, z) = (x + y)z$$

Let $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain Rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

$$\frac{\partial f}{\partial x}$$

Neural Networks

- Multi-layer Perceptrons
- Backpropagation

Neural Networks

- Multi-layer Perceptrons
- Backpropagation
- ... on Images
- Loss functions for more than two classes
- Better activation functions
- Better gradient descent

Course Outline

- Introduction to Image Analysis
- **Basics: Neural Networks**
- **Convolutional Neural Networks**
- Transformers
- Model Interpretability
- Self-supervised Learning
- Generative Models (GANs, Diffusion)